

# COGS300

# Deep Learning

Instructor: Márton Sóskuthy  
[marton.soskuthy@ubc.ca](mailto:marton.soskuthy@ubc.ca)

TAs: Daichi Furukawa • Victoria Lim • Amy  
Wang  
[cogs.300@ubc.ca](mailto:cogs.300@ubc.ca)



## Participants needed for a user-study

**Explainable AI in Intelligent Tutoring Systems.**



Max upto 3 Hours

**Typically 2 to 2.5 hours**

ICICS/CS Building

**Get \$40**

**Looking for participants who are majoring in Computer Science or related fields and meet the following requirements:**

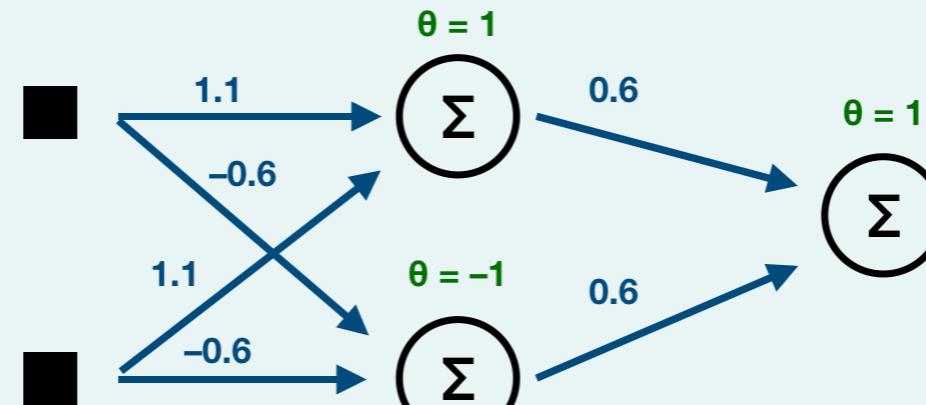
- HAVE taken CPSC 110.
- HAVE NOT taken or ARE CURRENTLY taking CPSC 322.
- do not have color vision deficiency.

**Sign Up:**

[https://hai.cs.ubc.ca/xai\\_study](https://hai.cs.ubc.ca/xai_study)



implements



implements

Connectionist system implemented by symbol system, and implementing symbol system

0	0	0
0	1	1
1	0	1
1	1	0

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$

Is this a symbol system or a connectionist system?

# Marr's levels of analysis

implementational level

physical implementation?

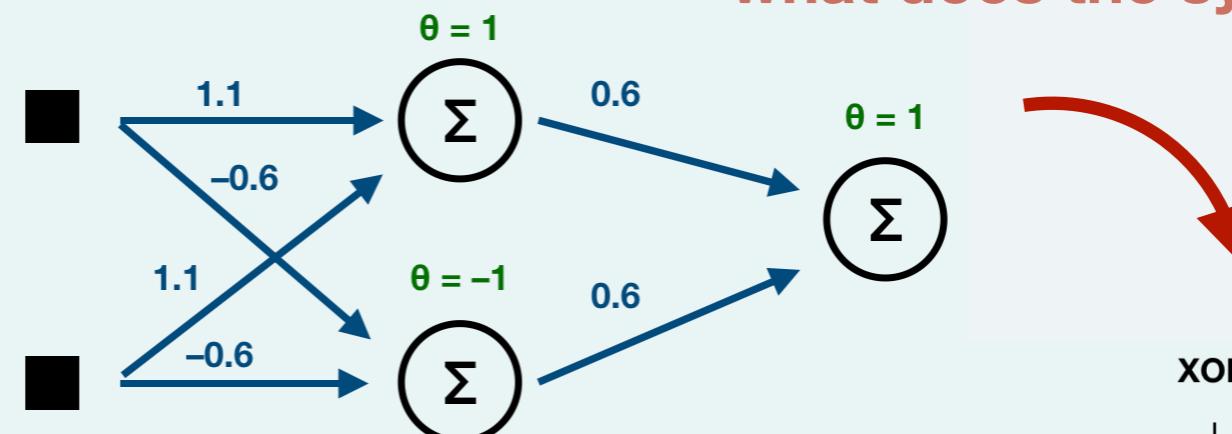
algorithmic/representational level

representations / processes?



computational level

what does the system do? goals?



XOR

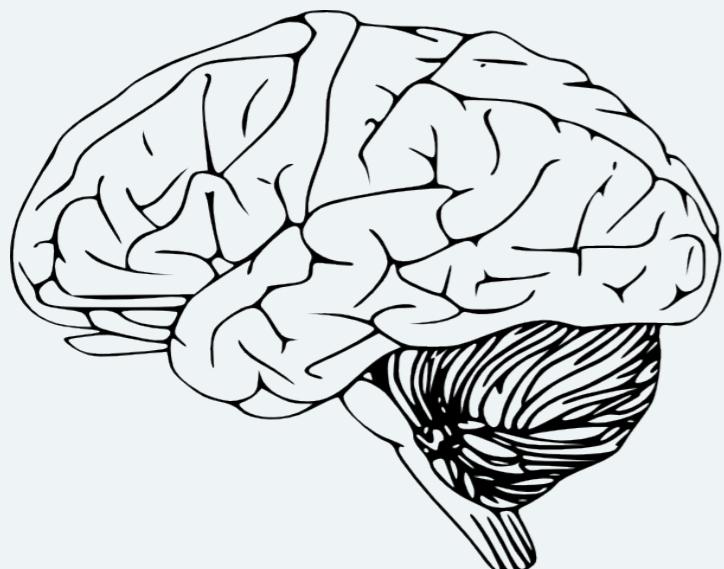
0	0	0
0	1	1
1	0	1
1	1	0

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$

# Marr's levels of analysis

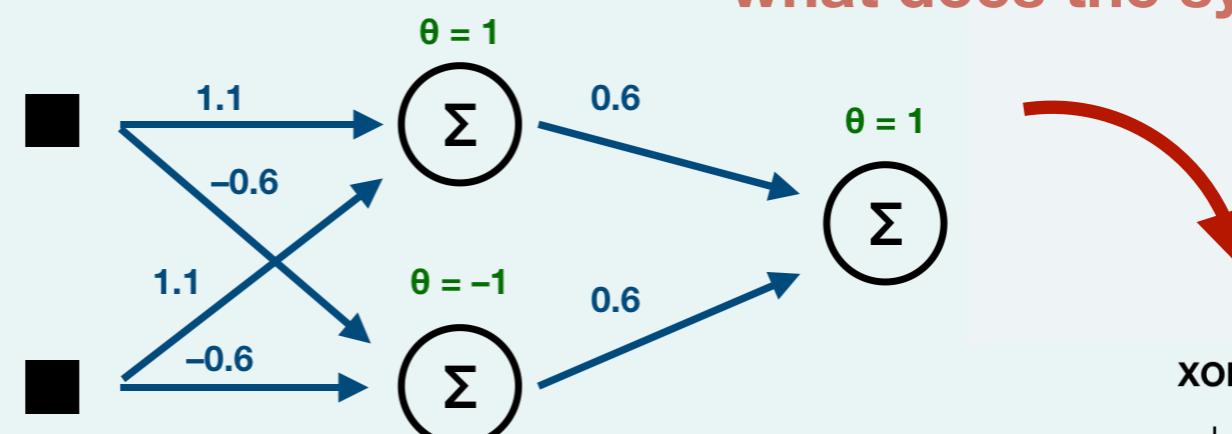
implementational level

physical implementation?



algorithmic/representational level

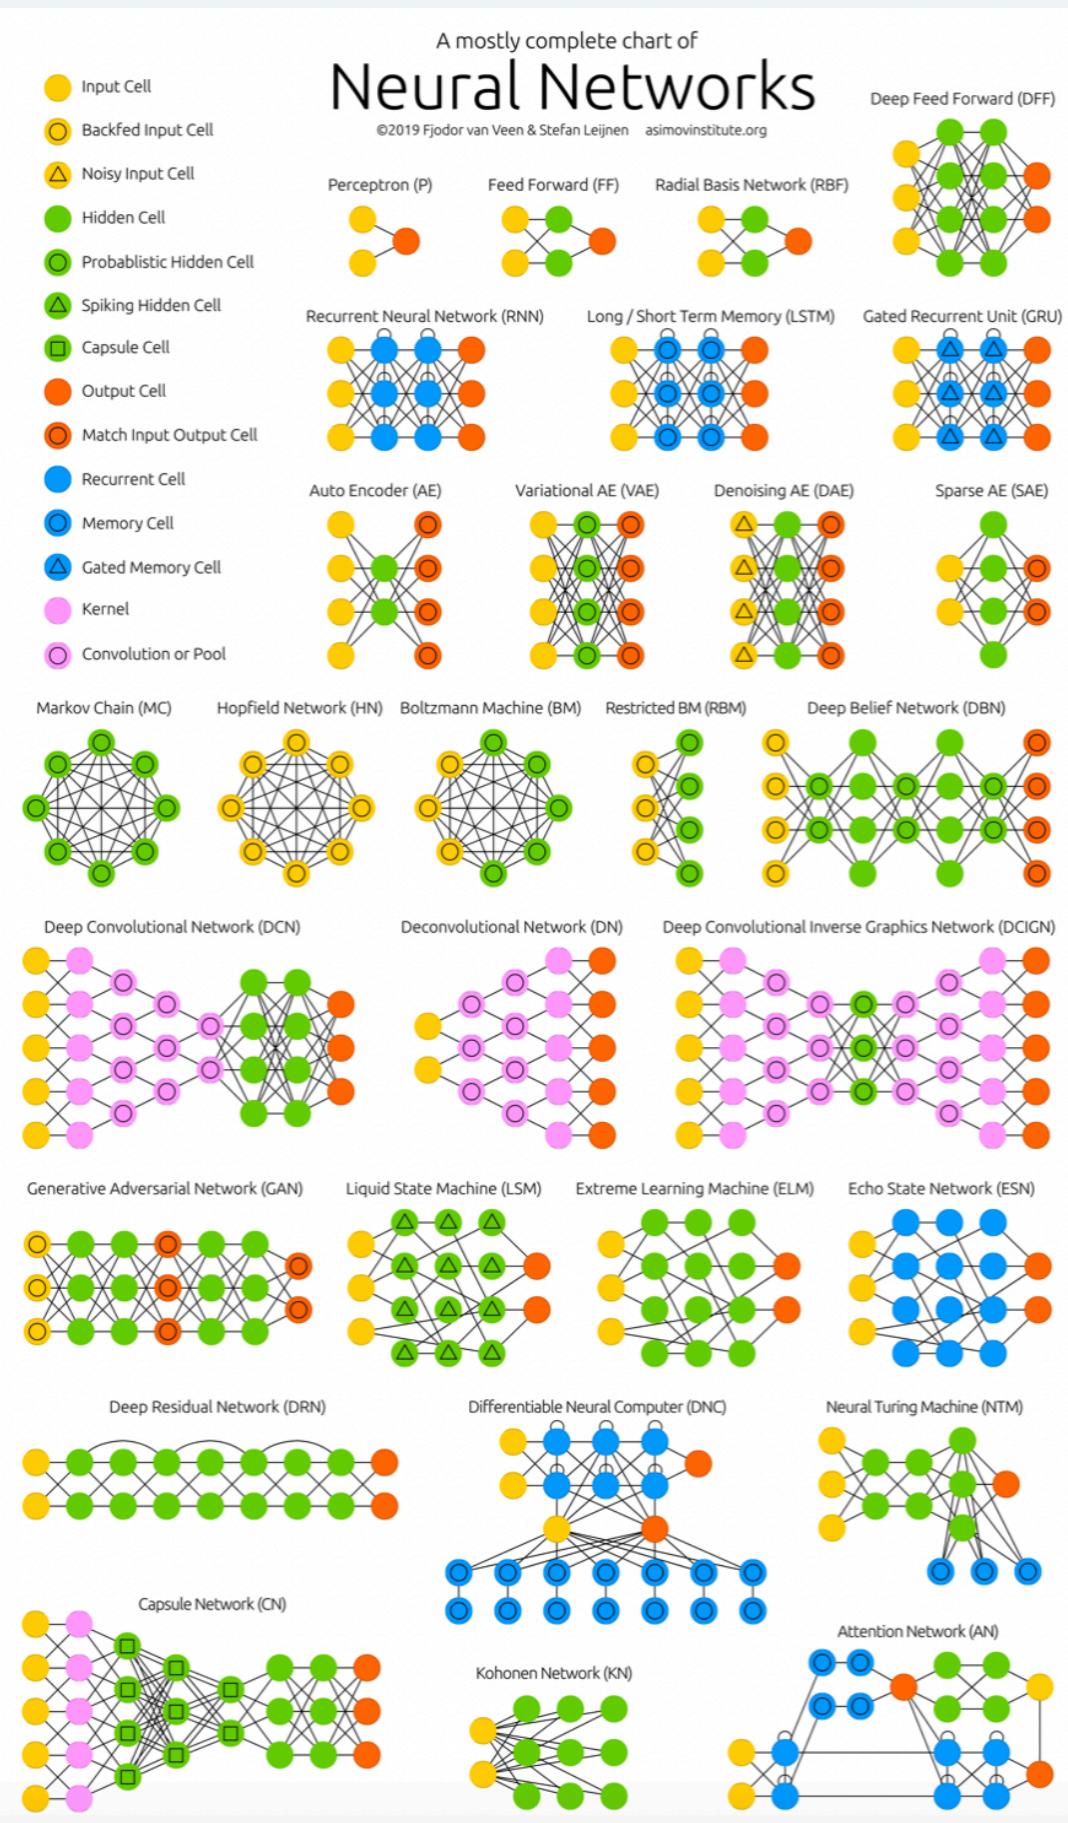
representations / processes?



computational level  
what does the system do? goals?

XOR	
0	0
0	1
1	0
1	1

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$

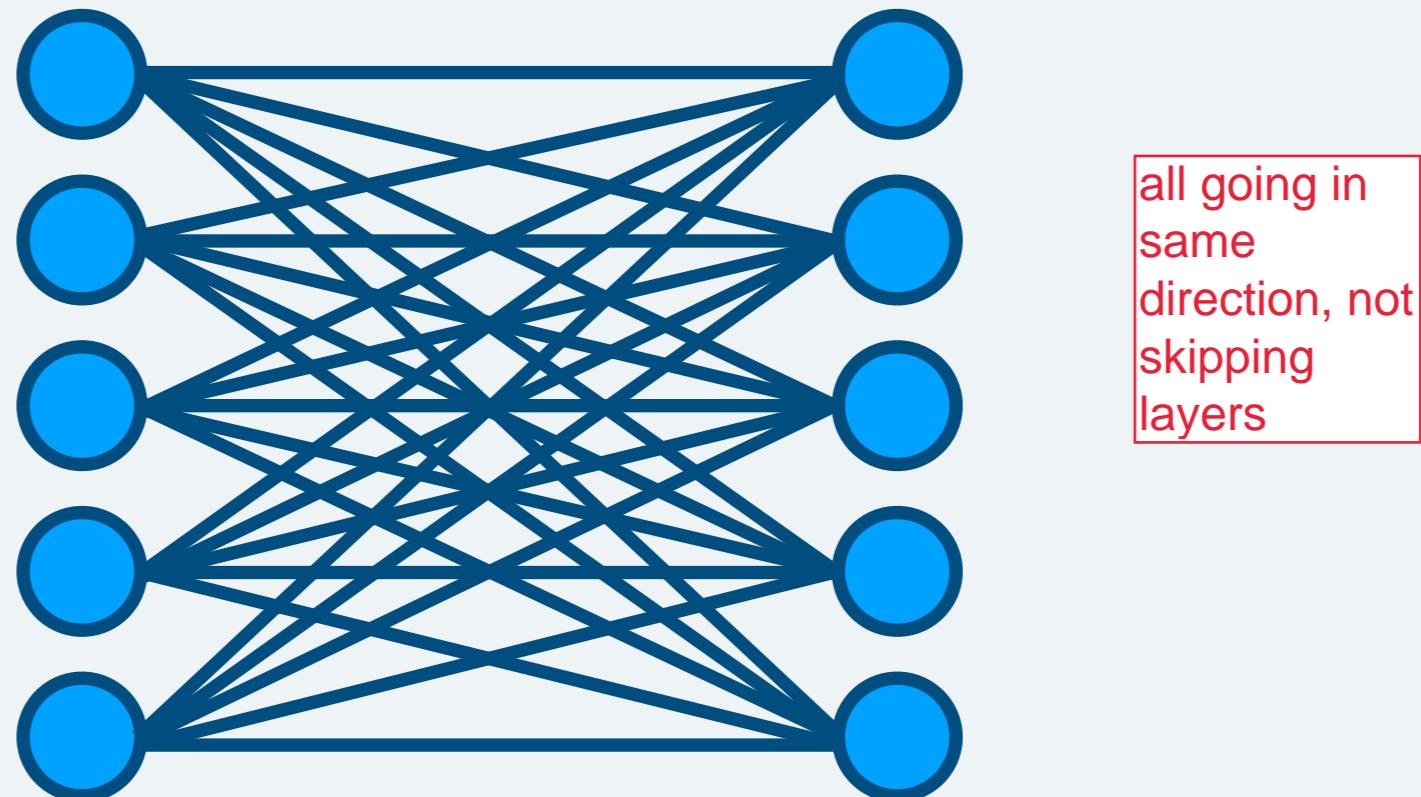


# Connectionist networks (2019)

[https://www.asimovinstitute.org/  
neural-network-zoo/](https://www.asimovinstitute.org/neural-network-zoo/)

# Connectionist networks

## Fully Connected (Feed Forward)



# Connectionist networks

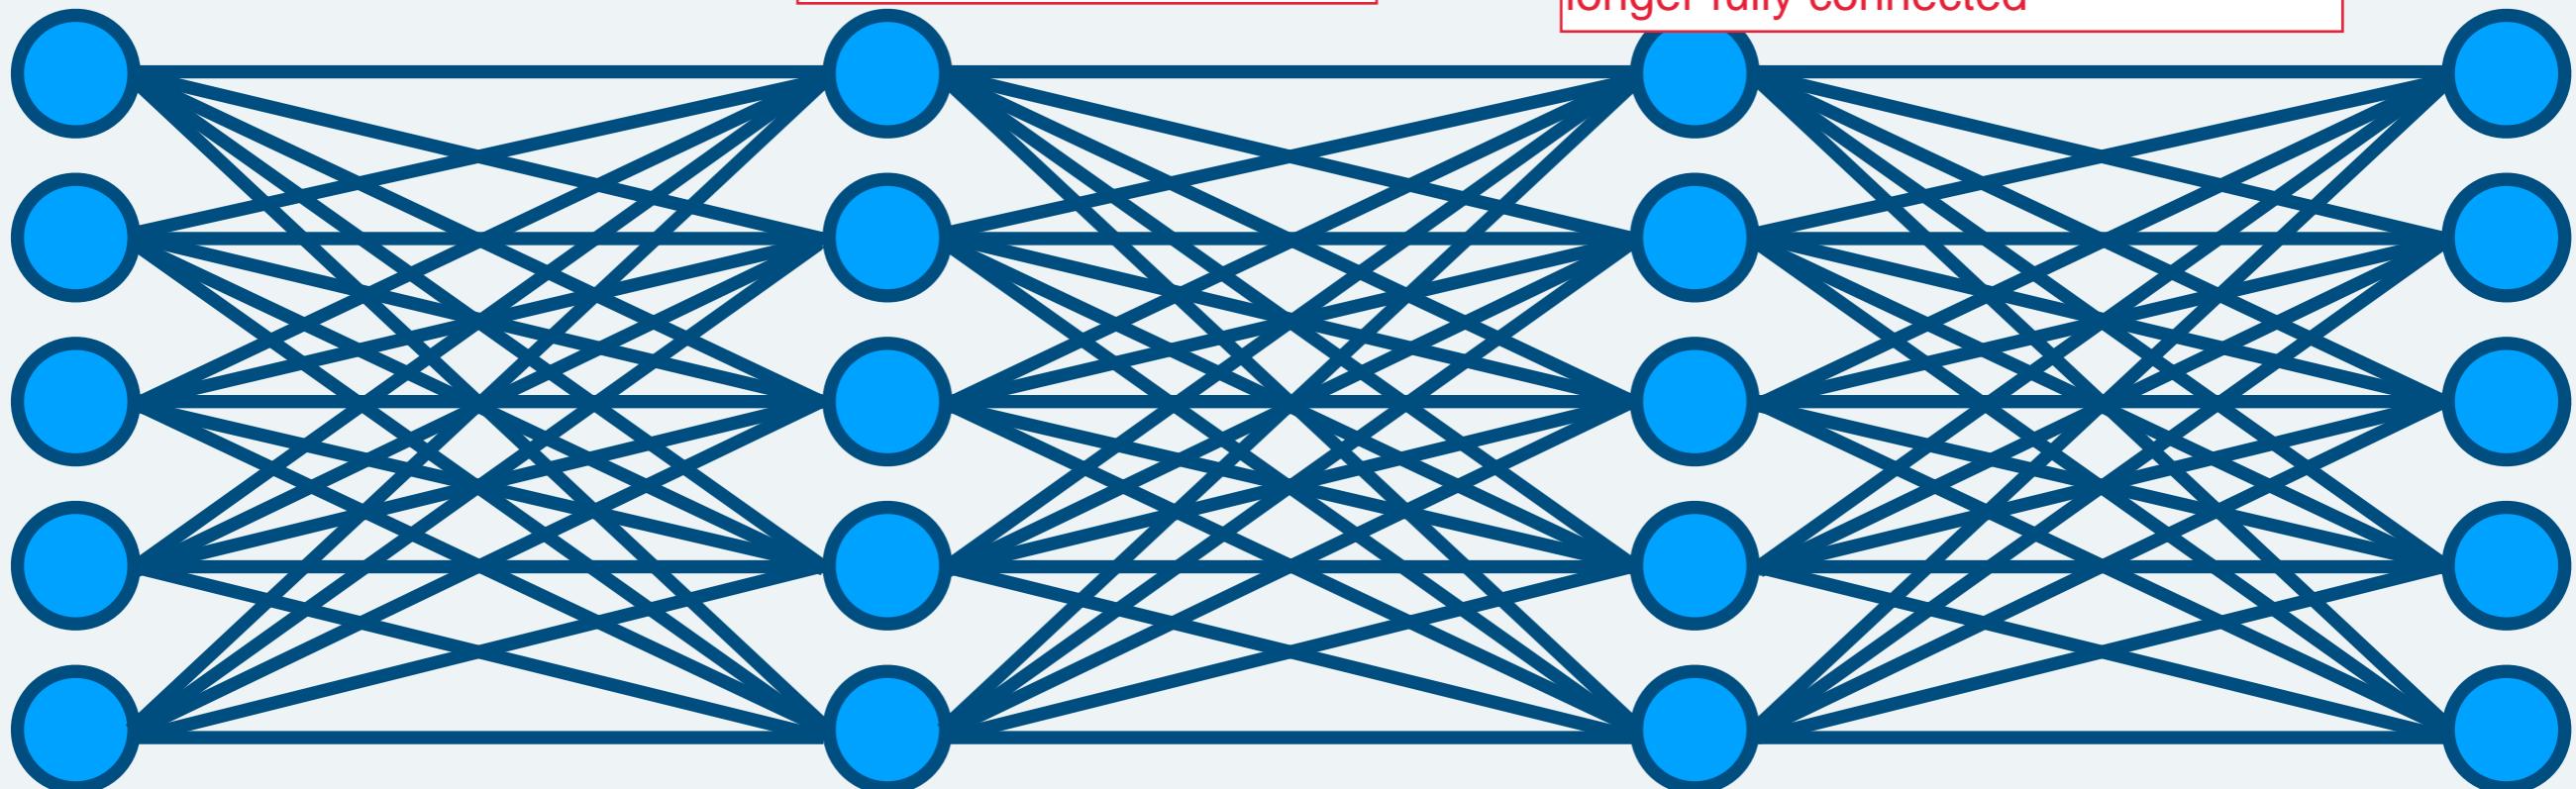
## Deep Fully Connected (Feed Forward)

Sequential

Loops =  
recurrent NN

Typically deep networks  
have at least 3 layers

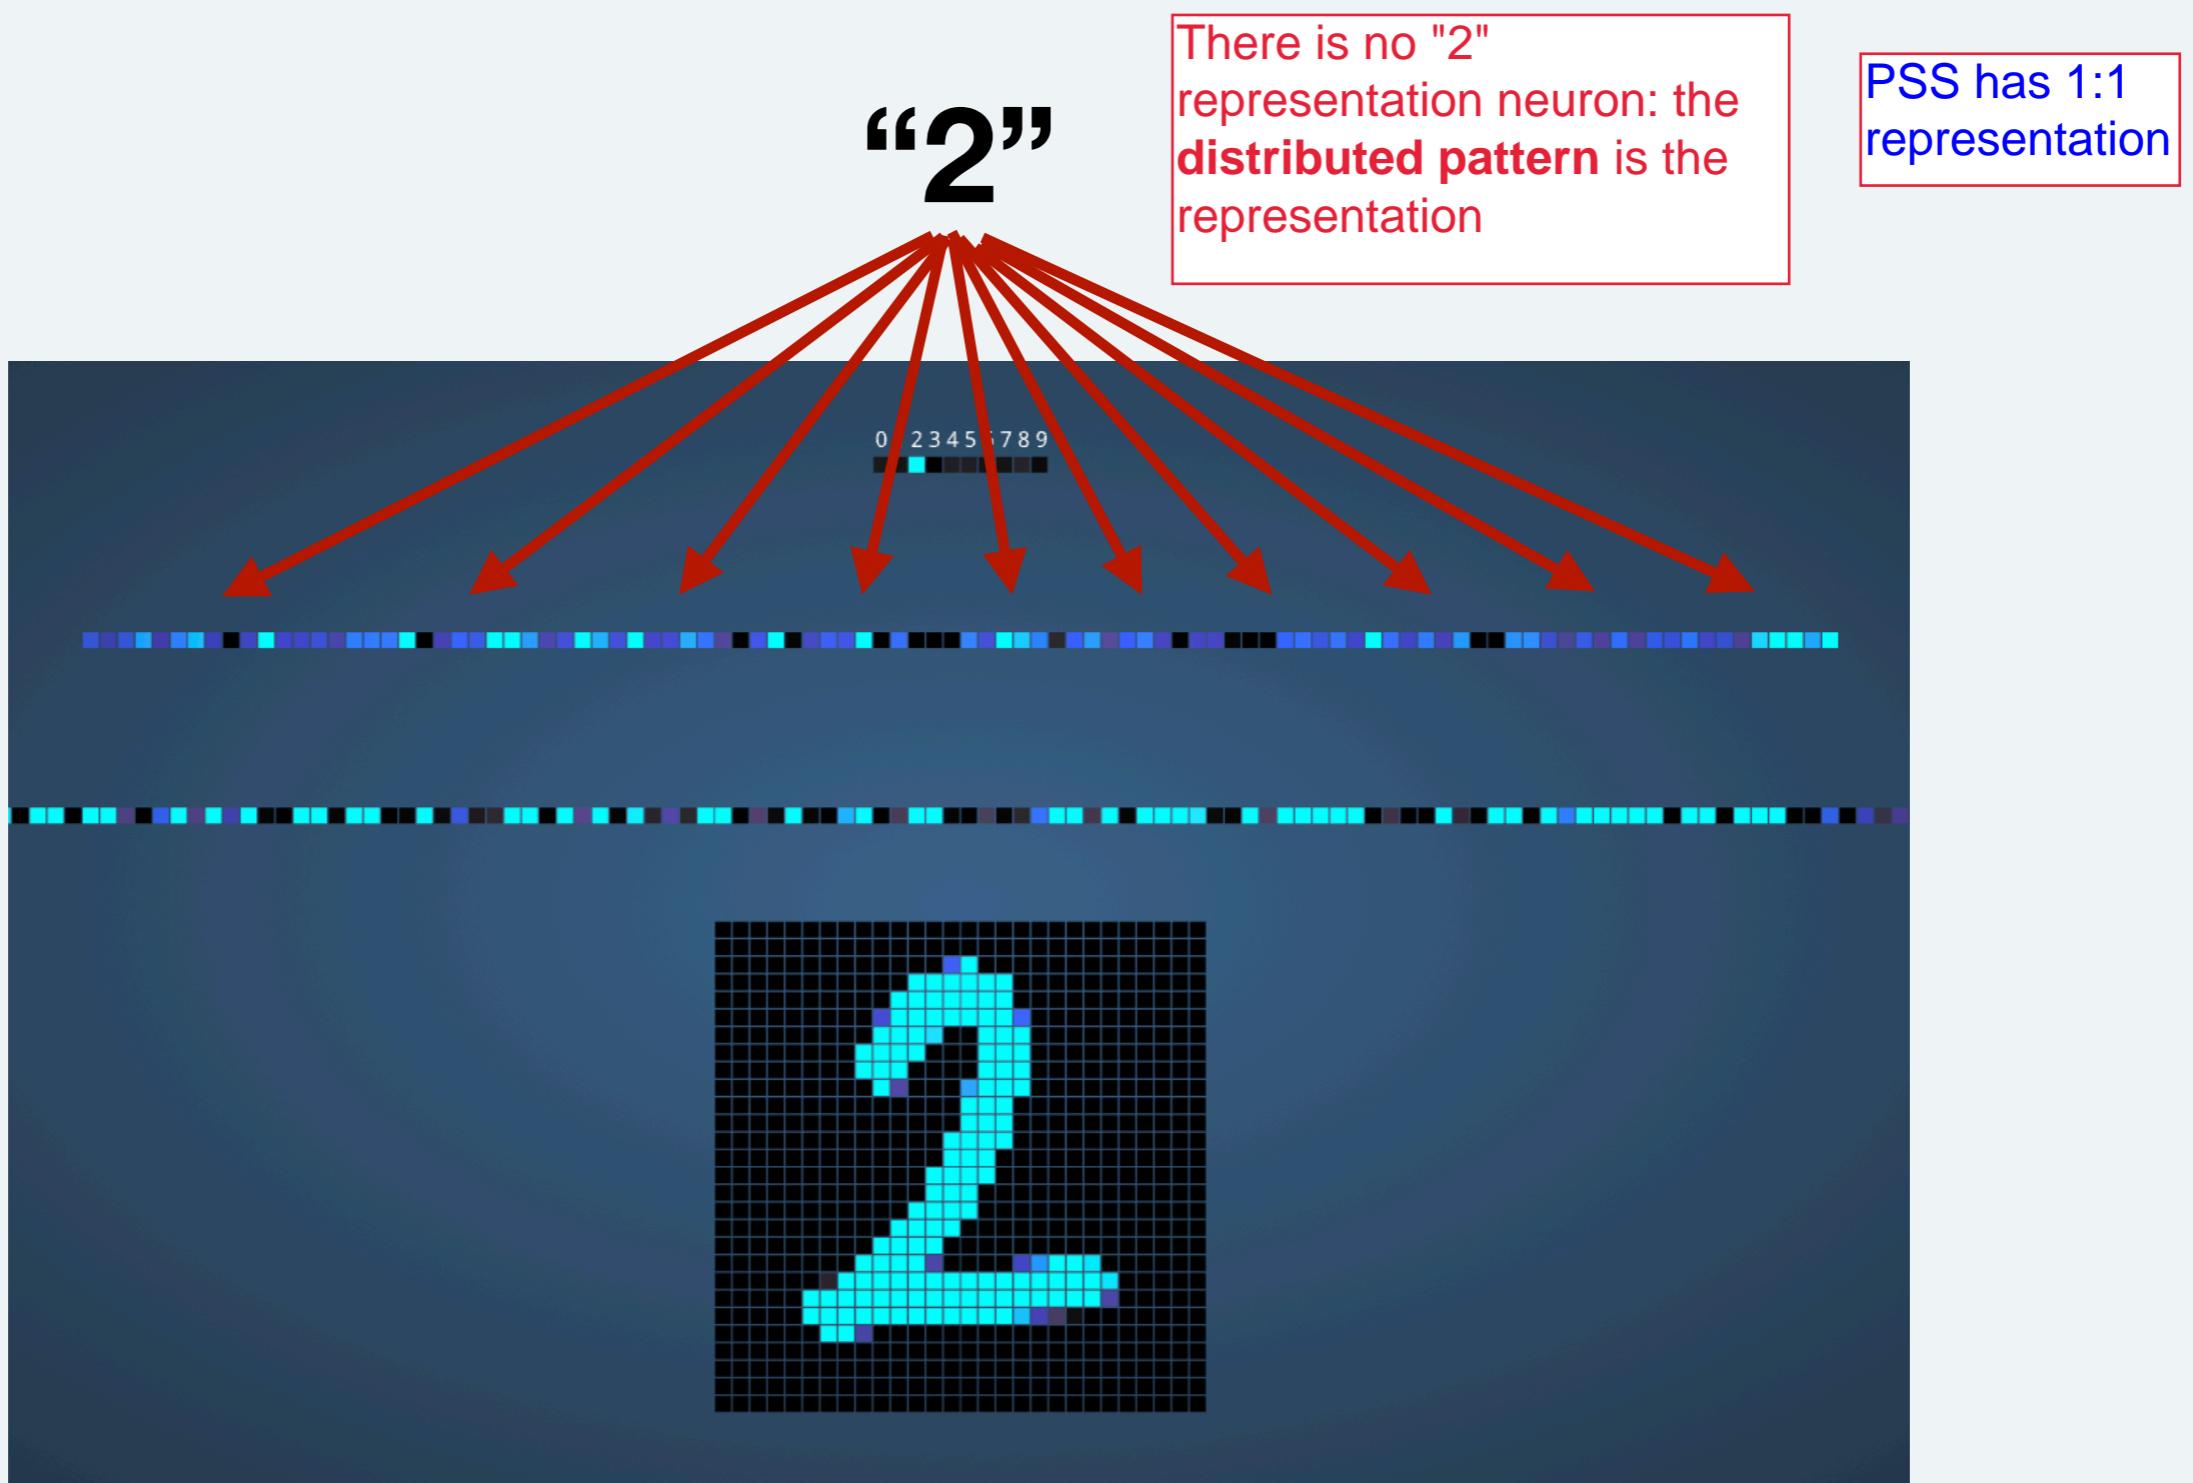
Often NNs aren't fully connected  
(just dropping 1 connection --> no  
longer fully connected)



# Connectionist networks

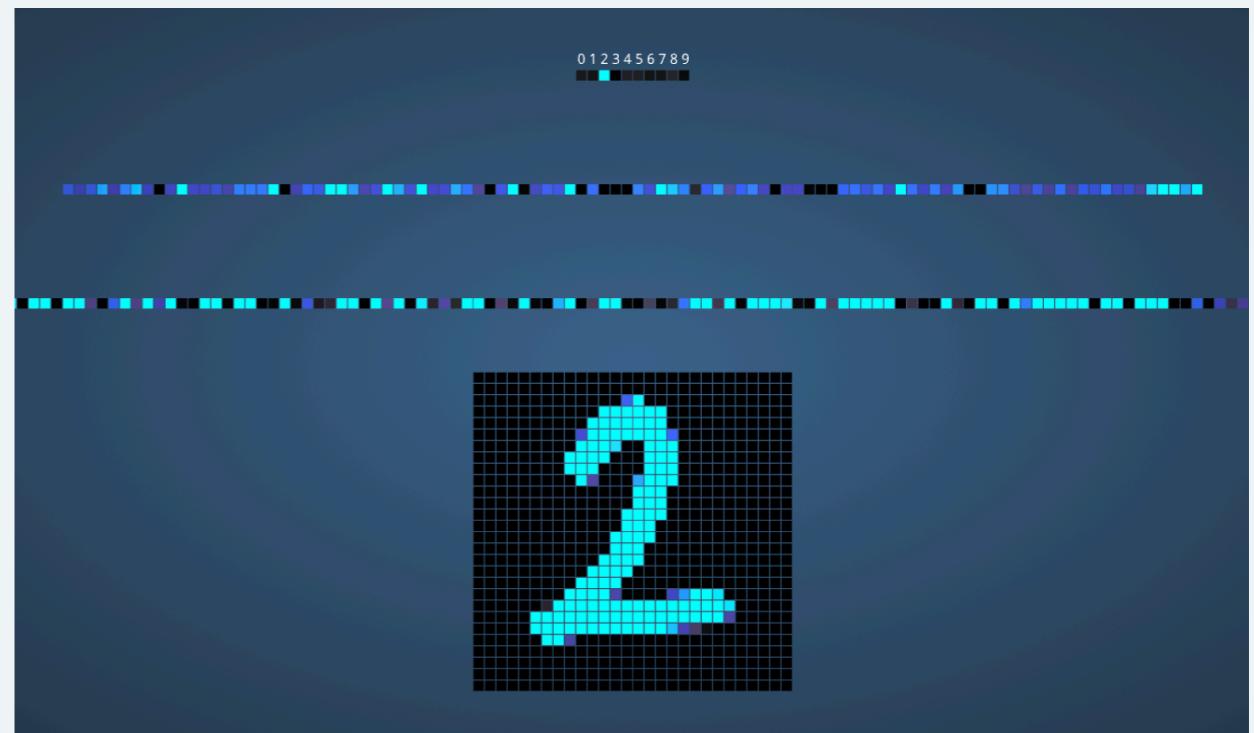
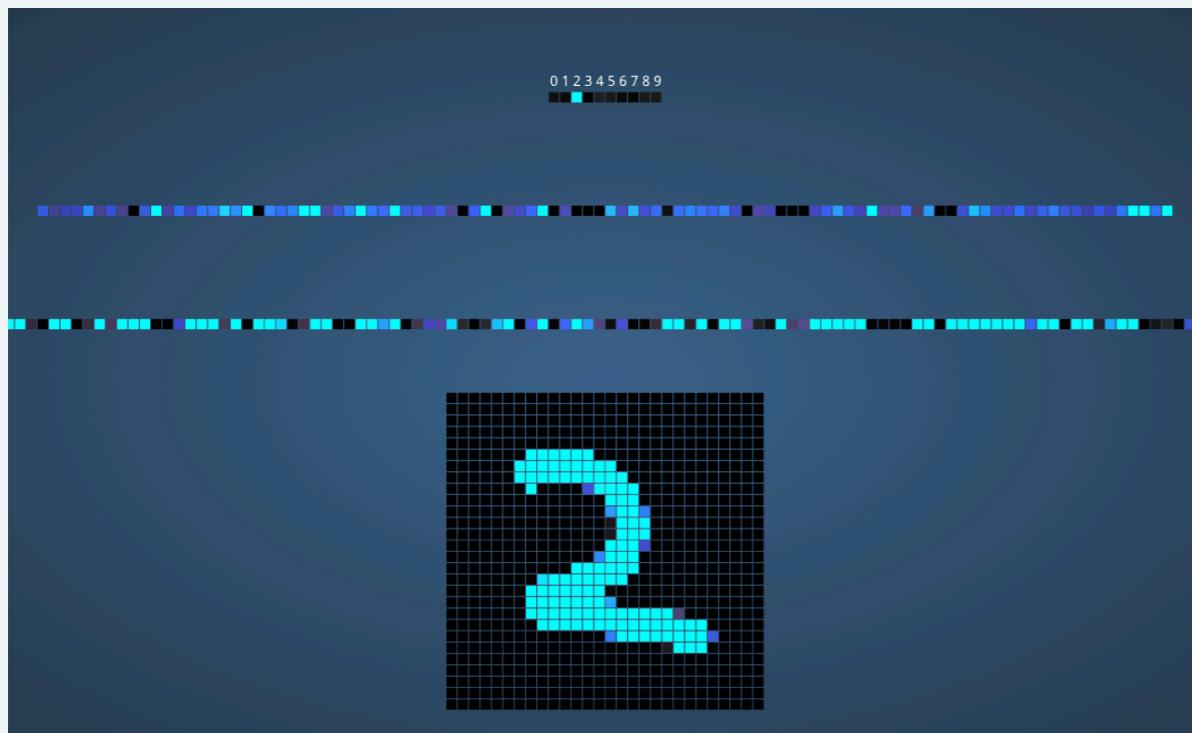
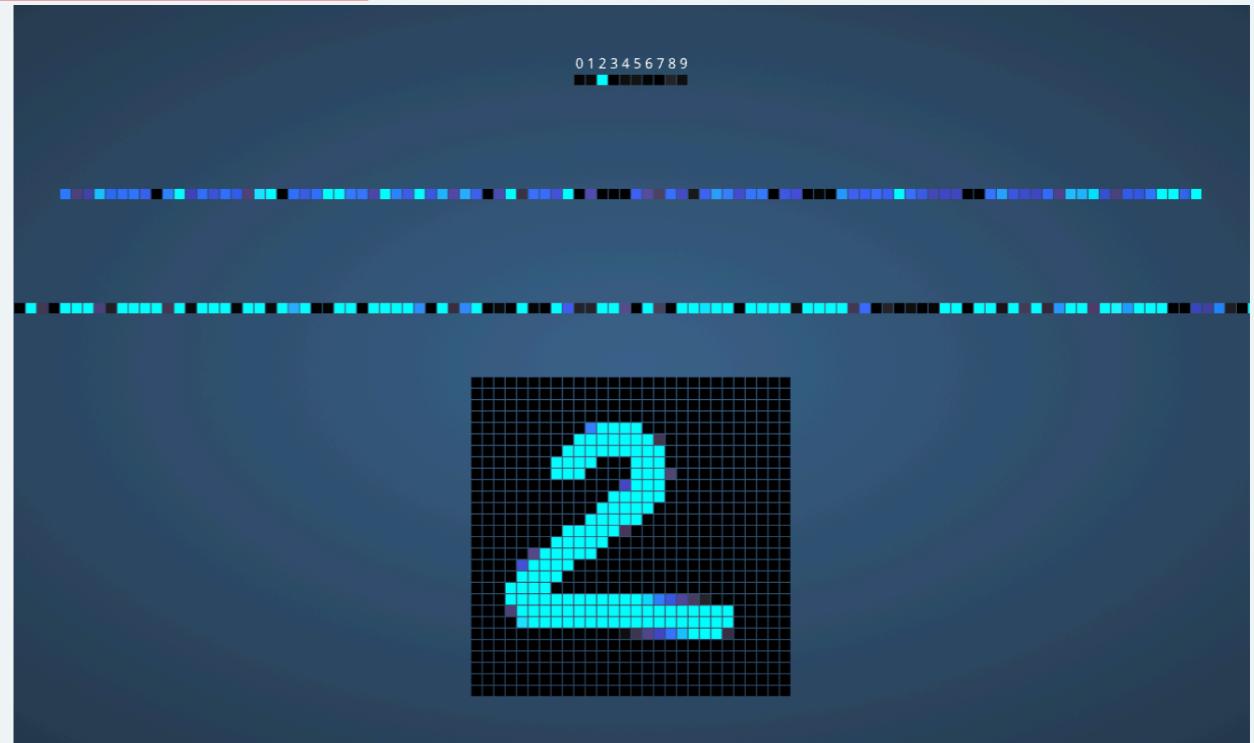
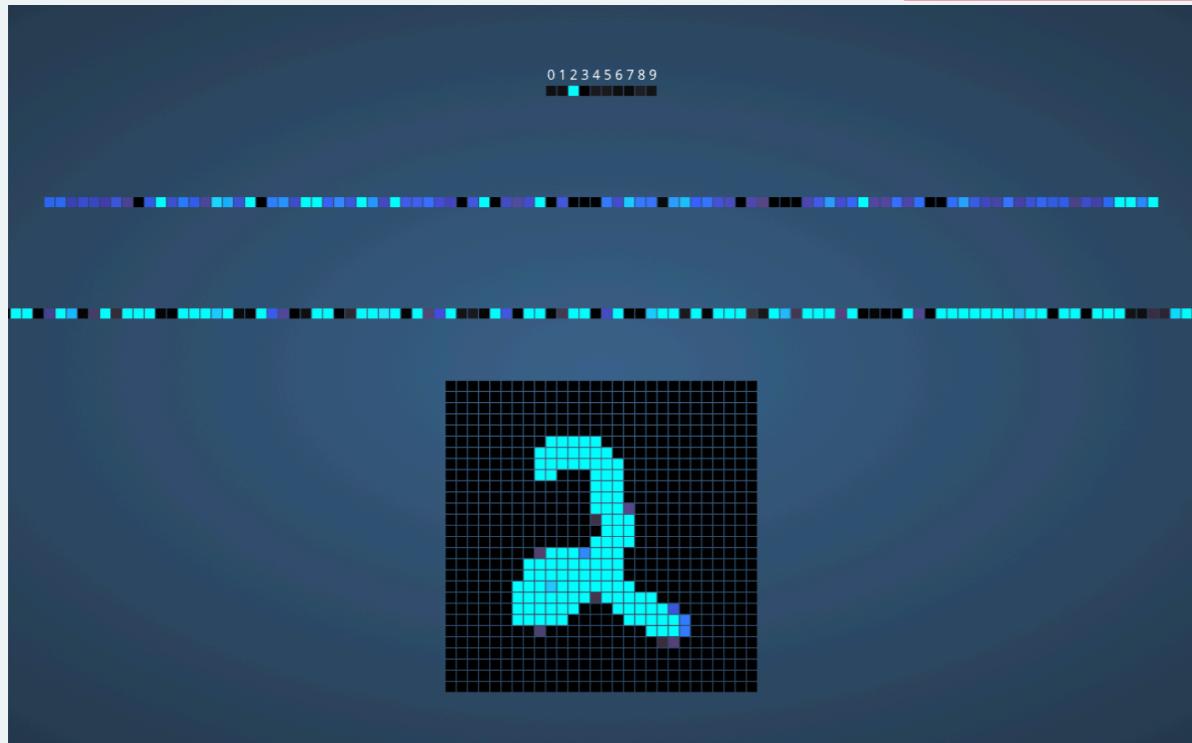
[https://adamharley.com/nn\\_vis/](https://adamharley.com/nn_vis/)

# Distributed representations

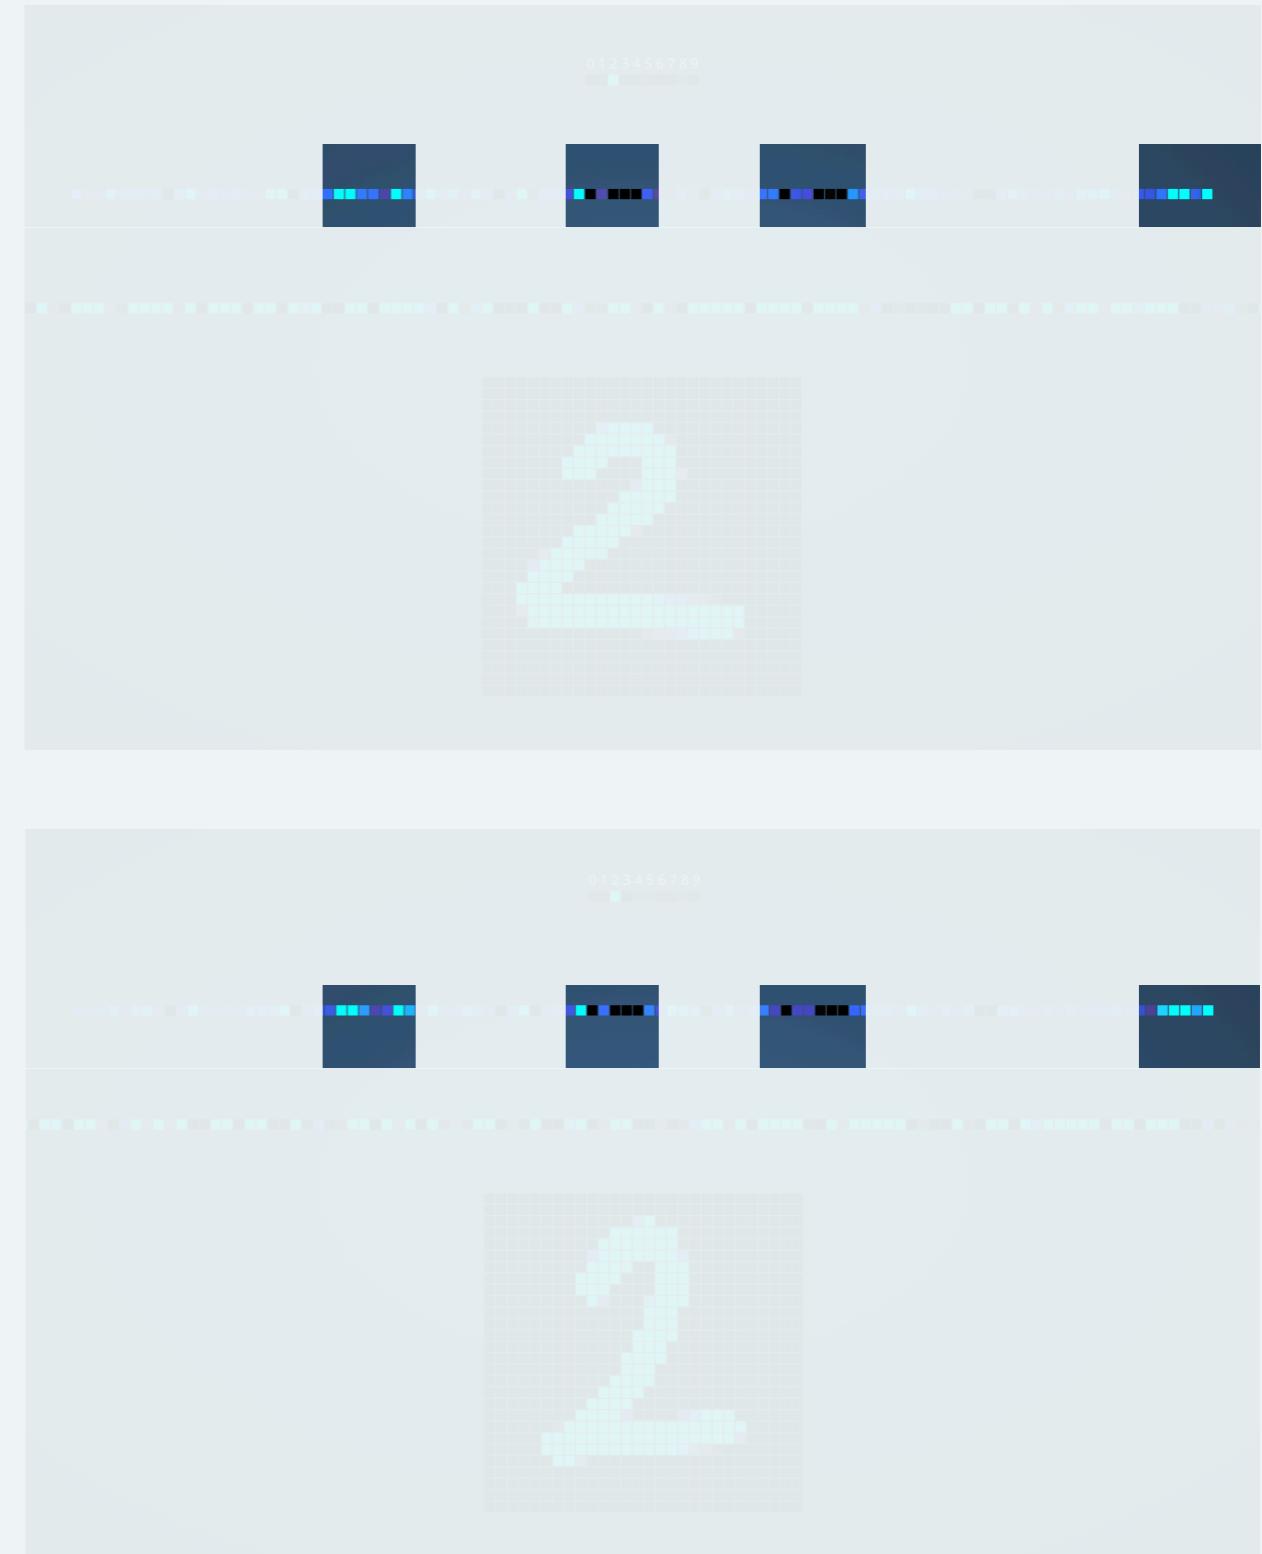
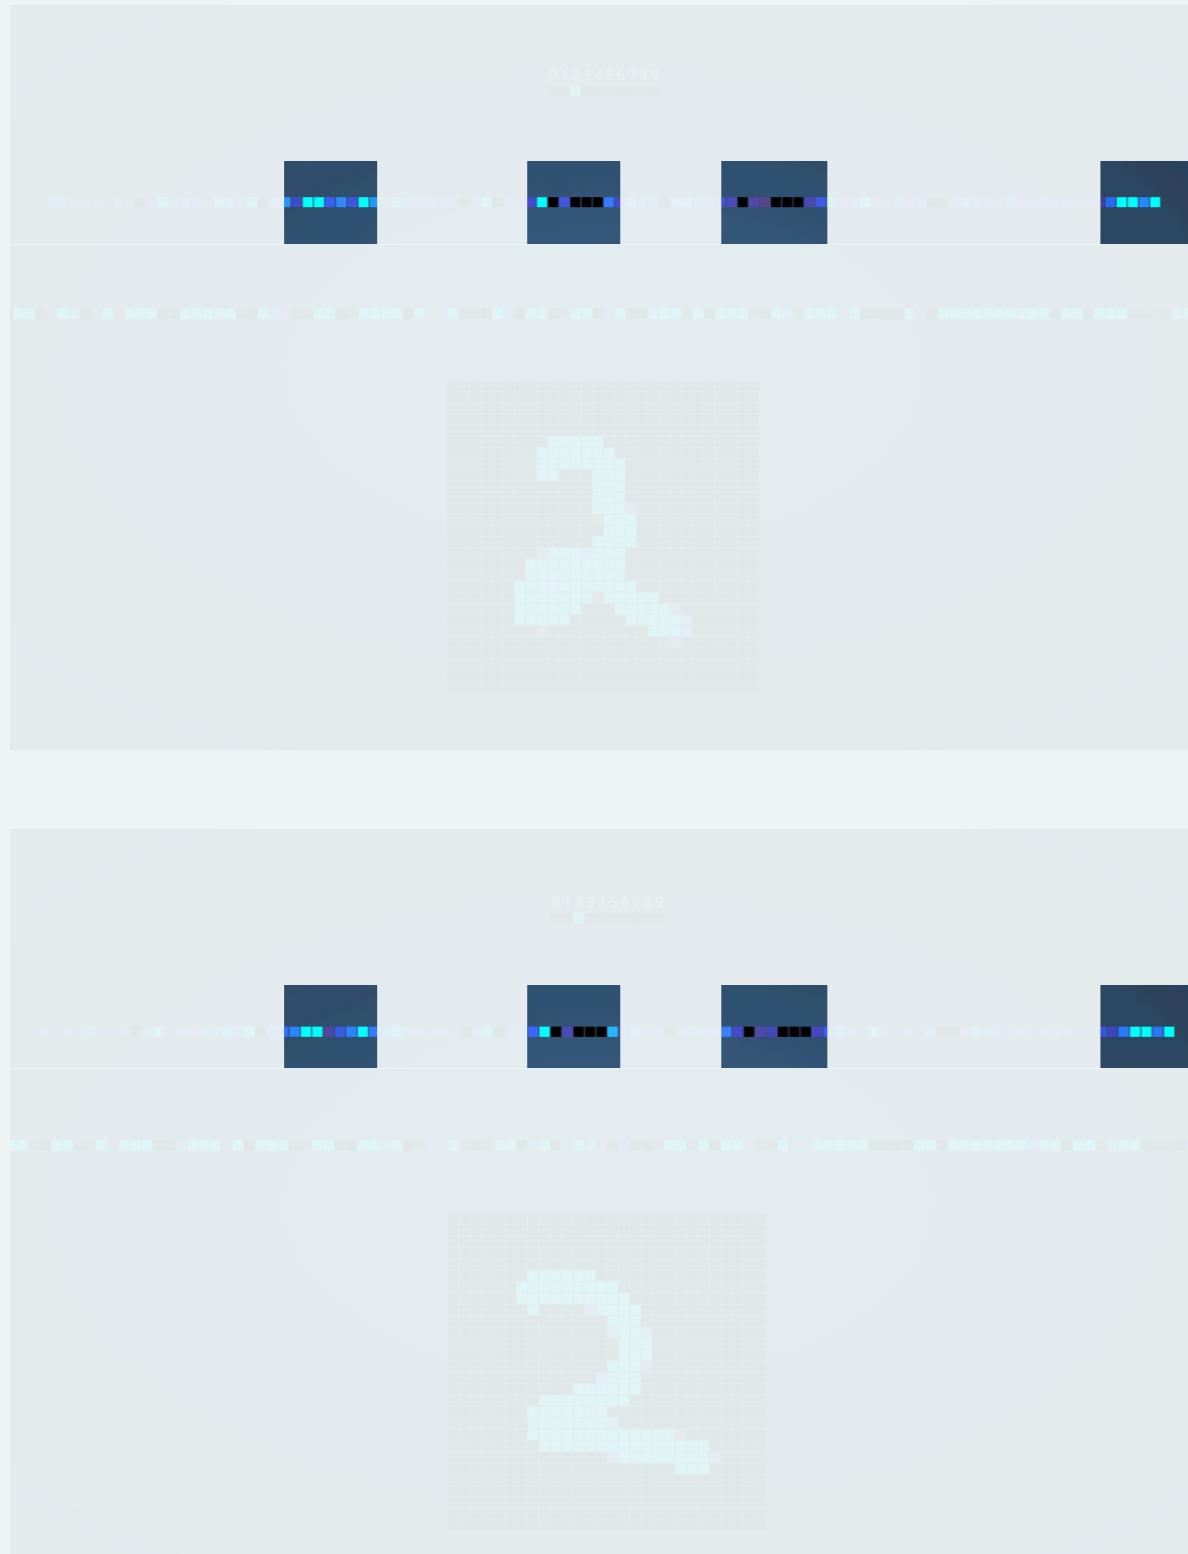


# Superpositional representations

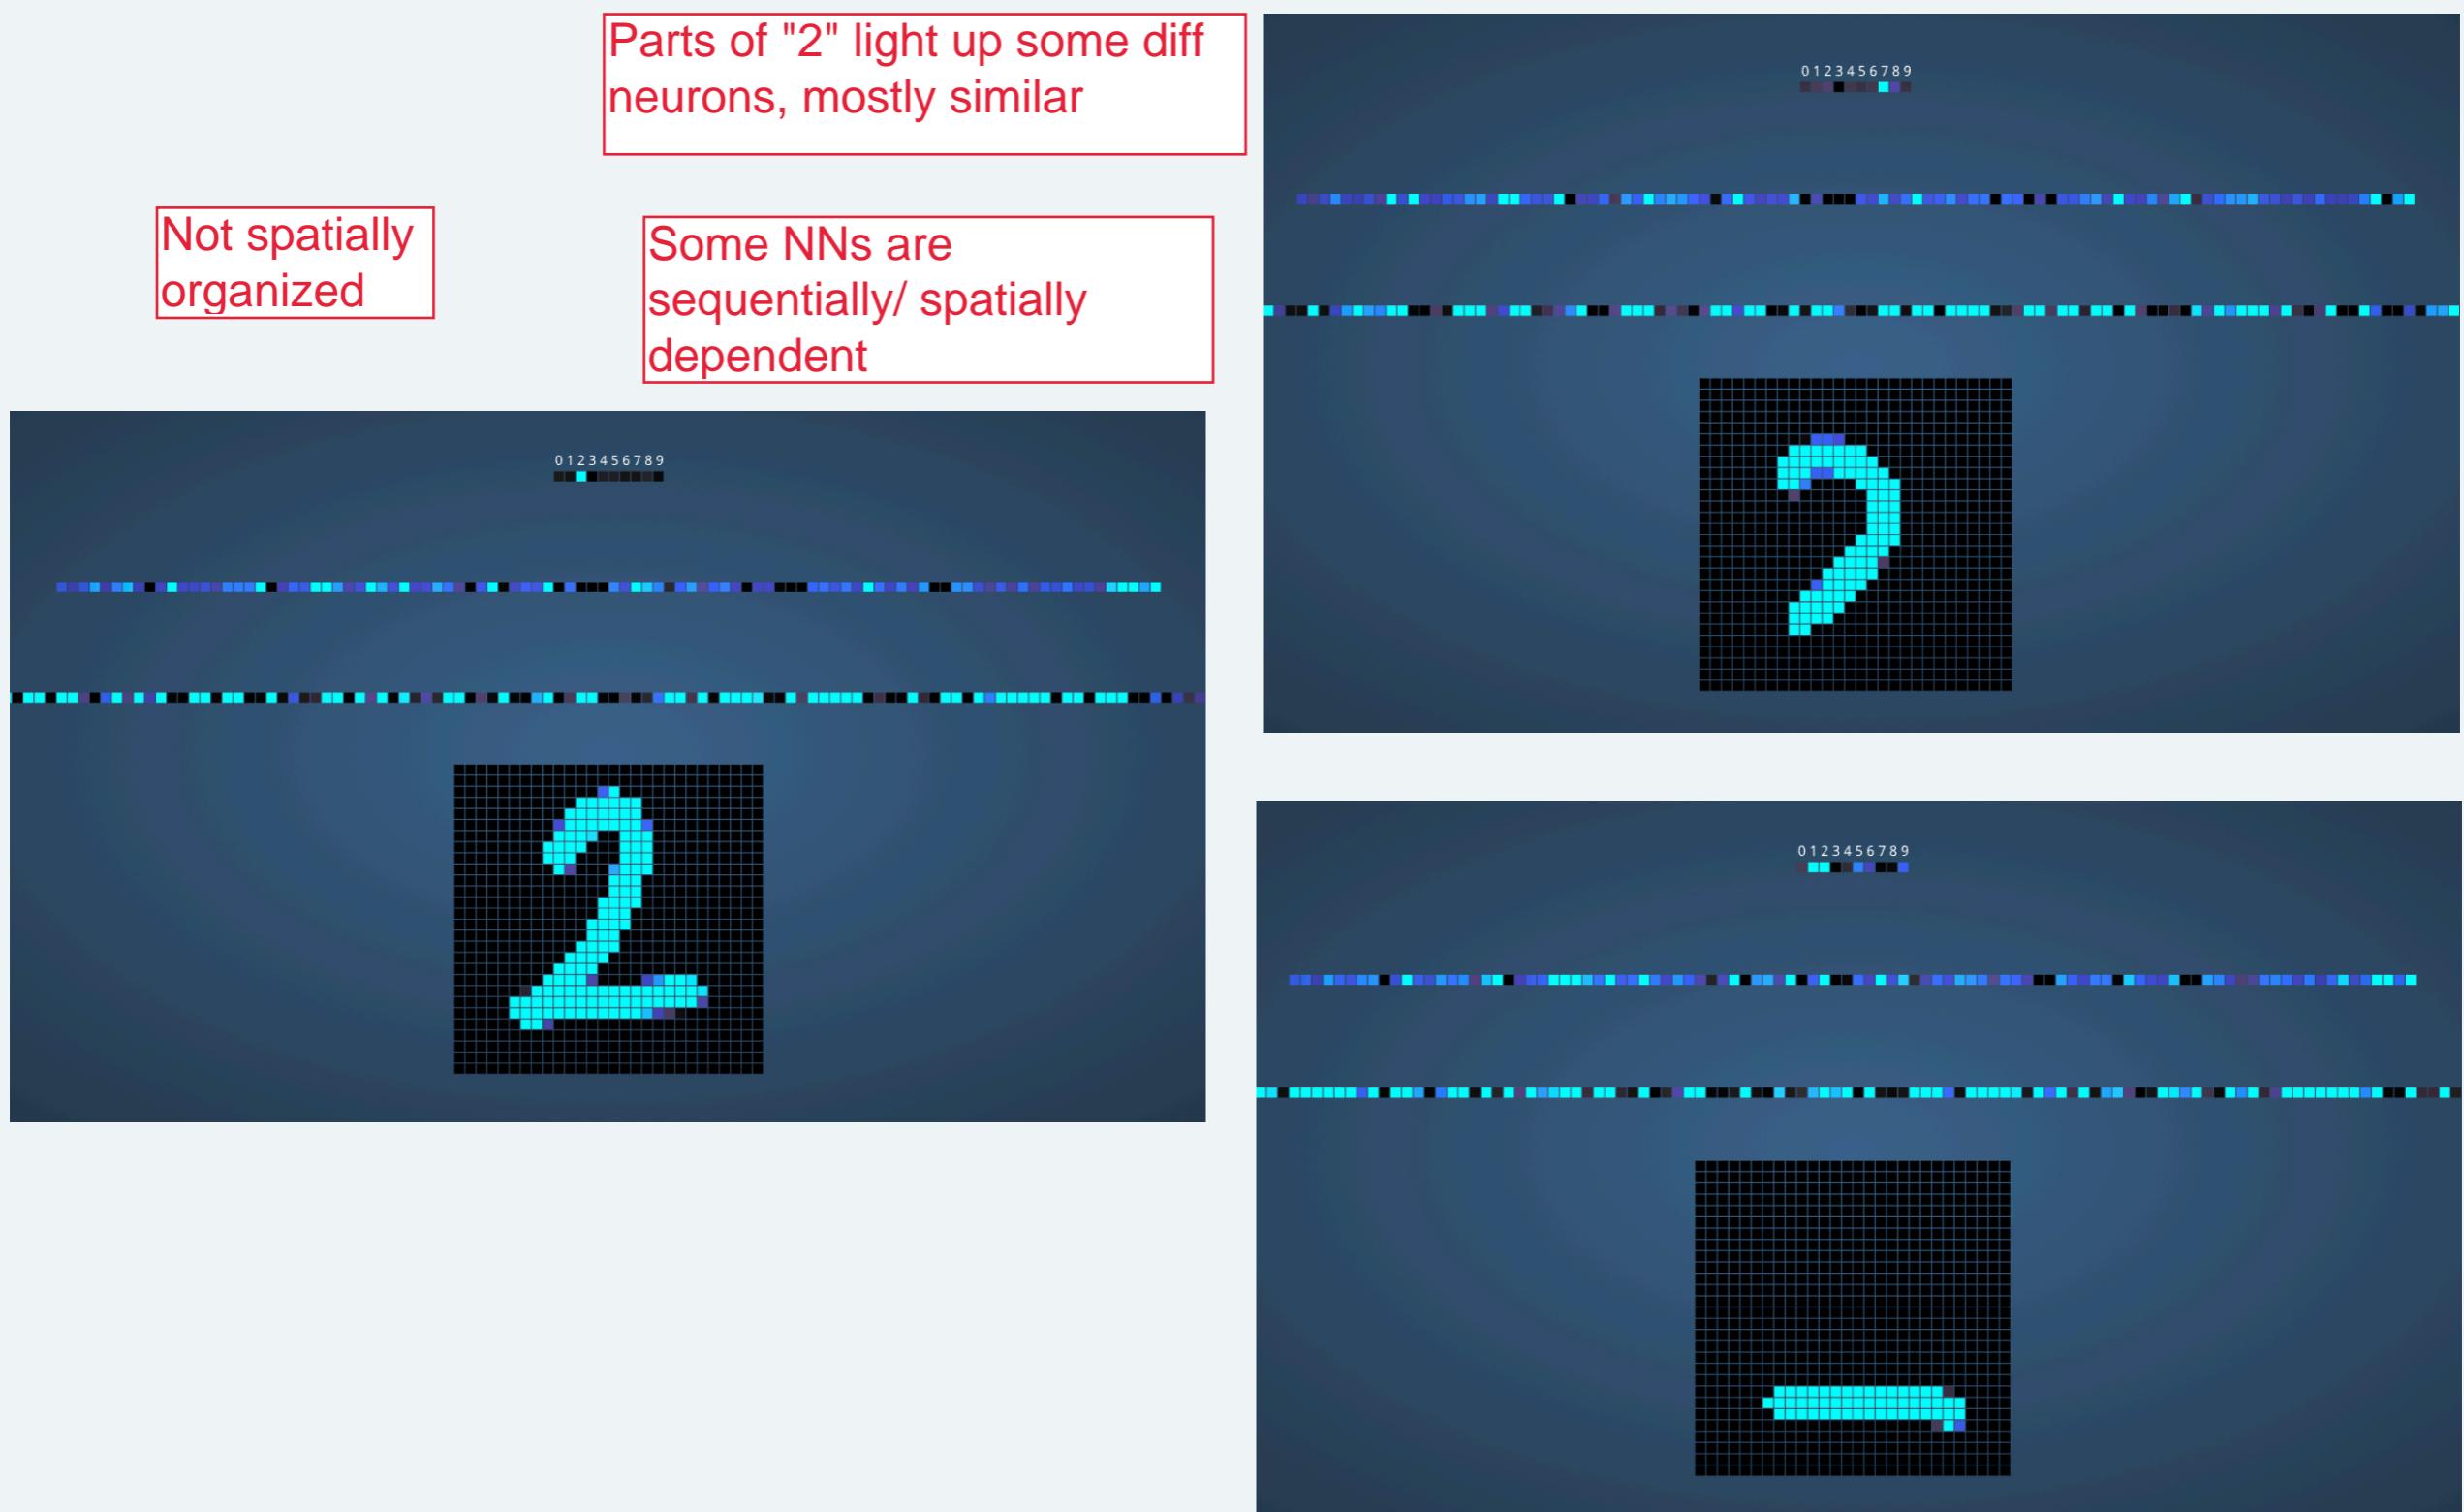
all similar activation patterns



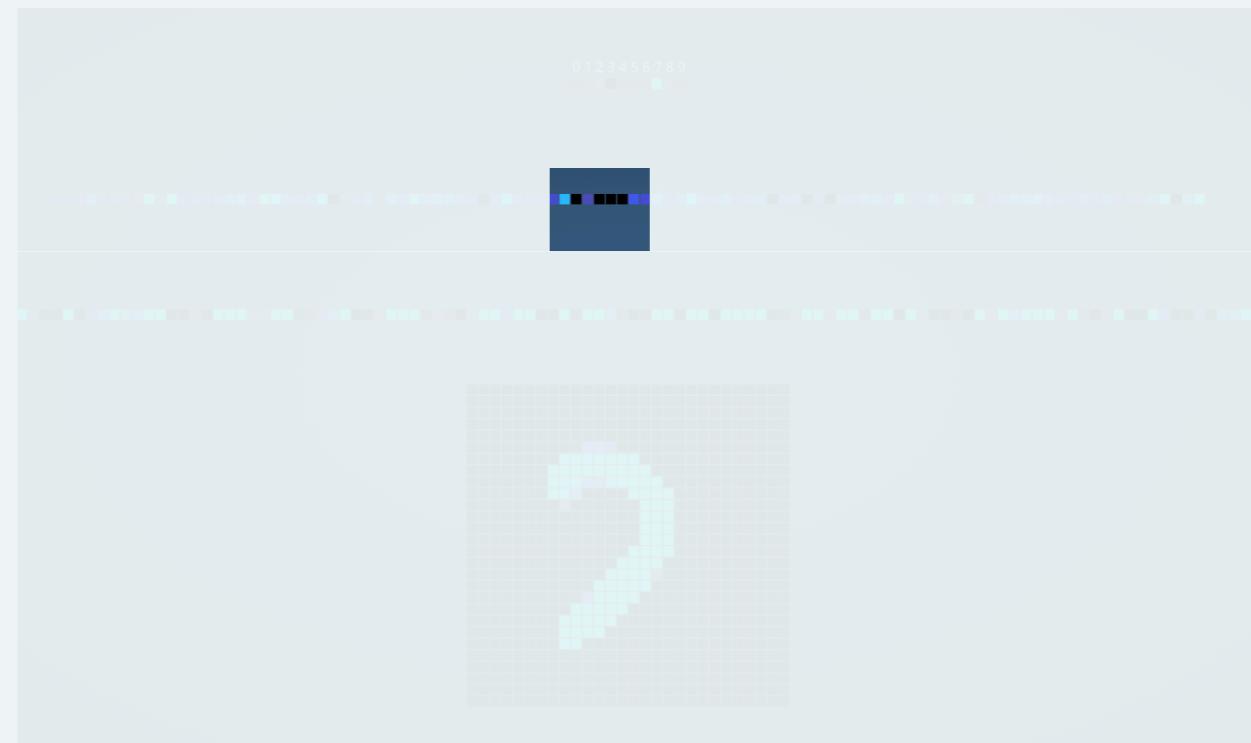
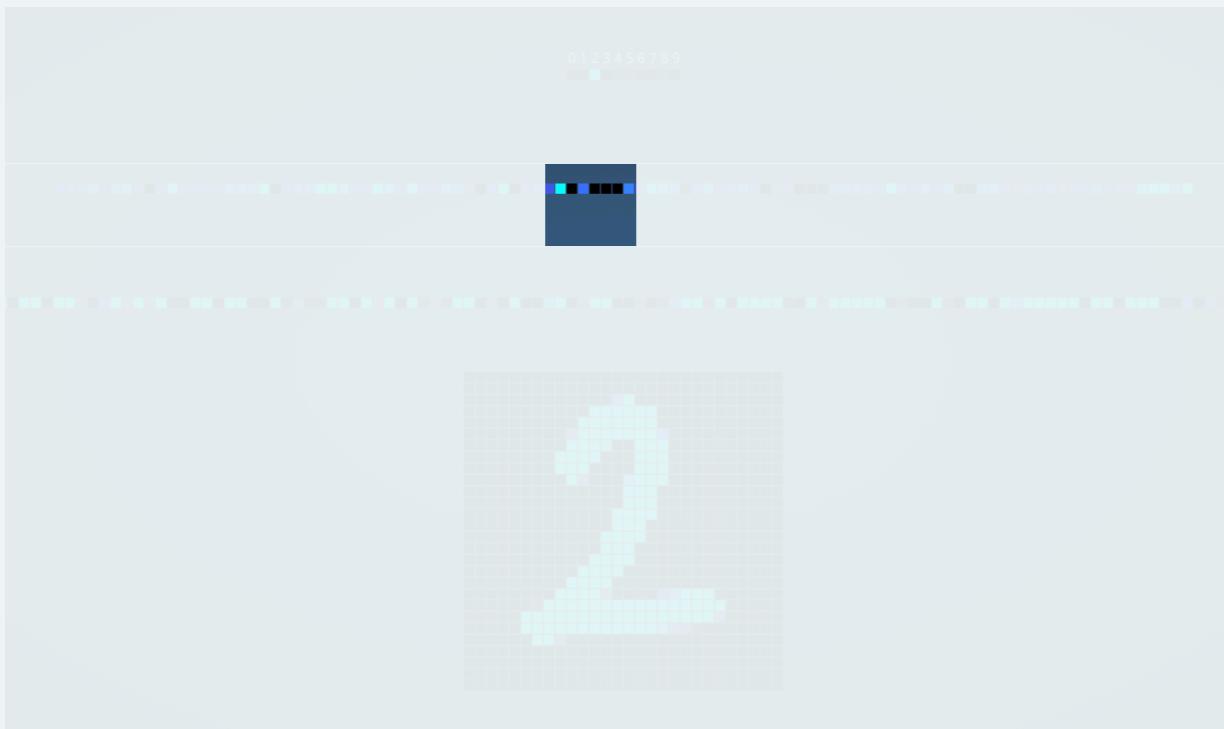
# Superpositional representations



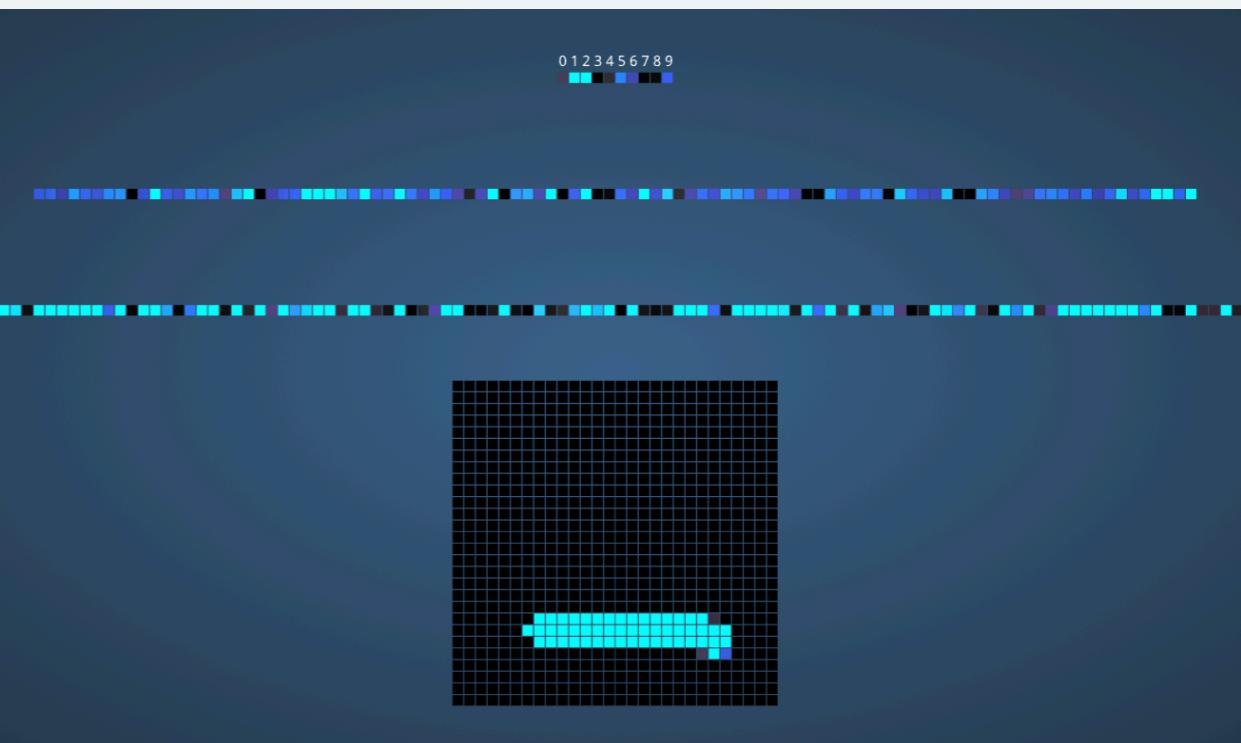
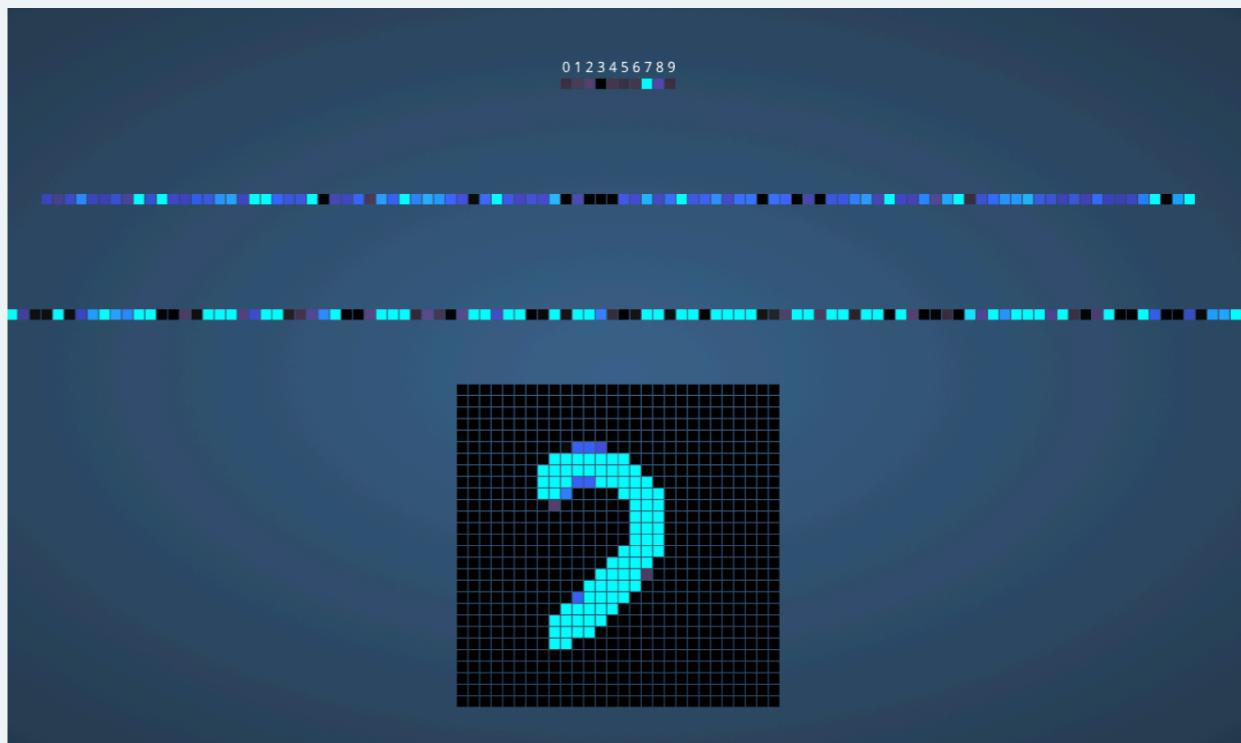
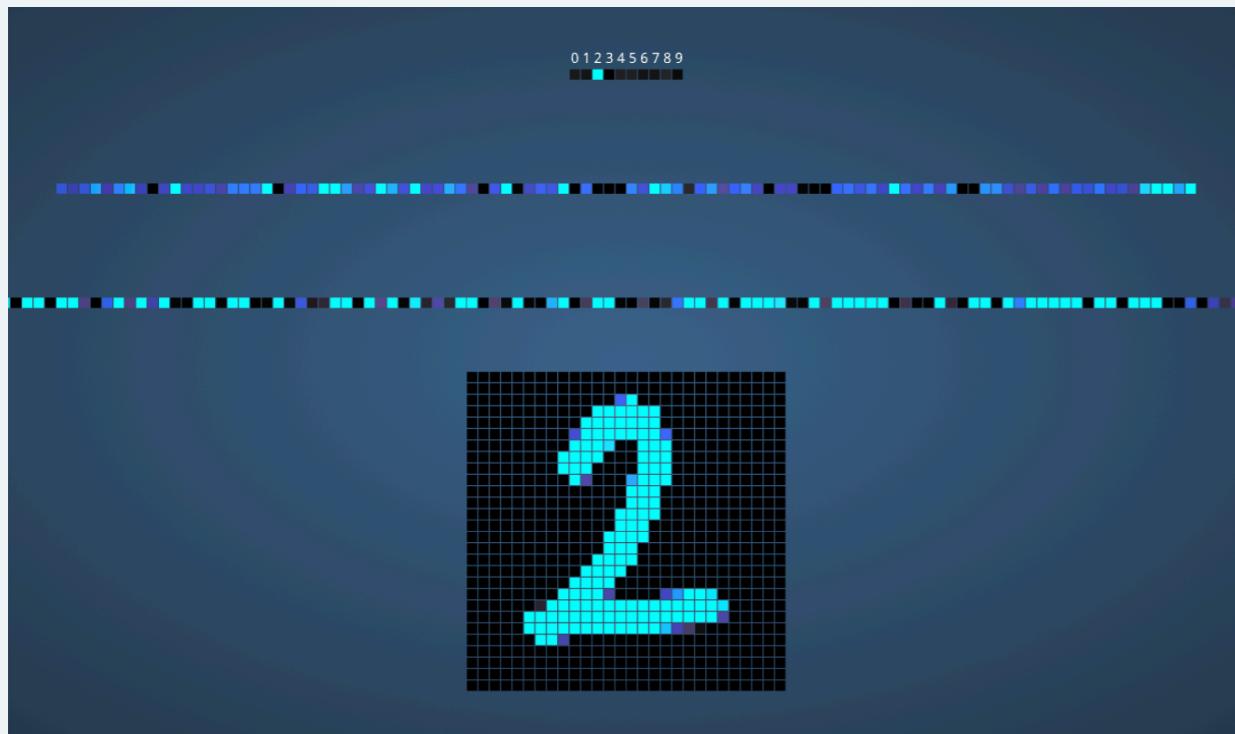
# Subsymbolic representations



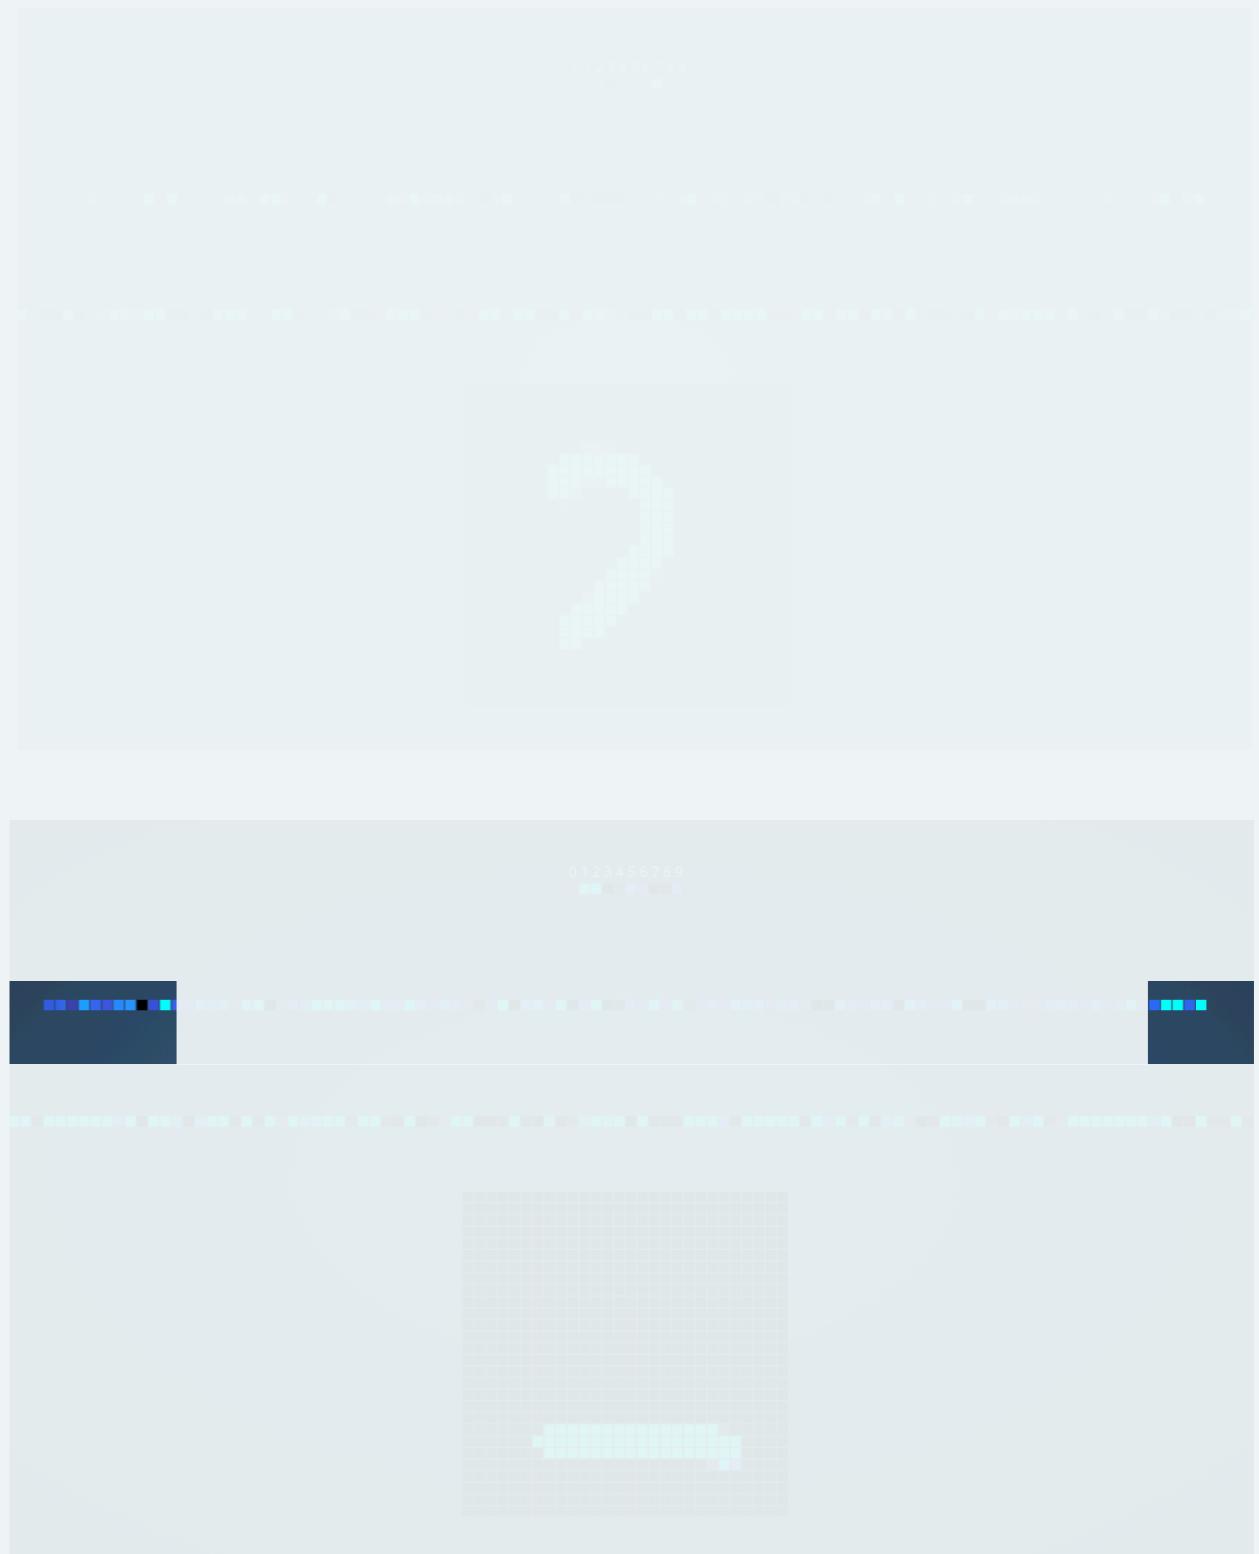
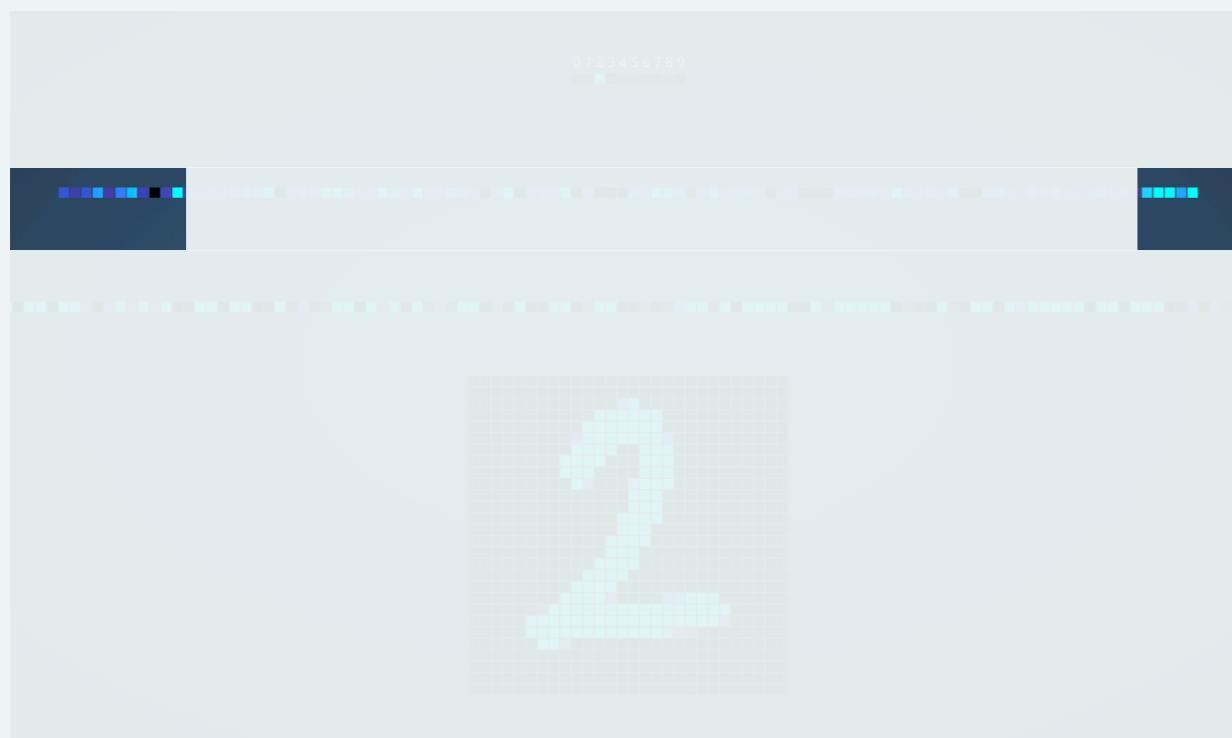
# Subsymbolic representations



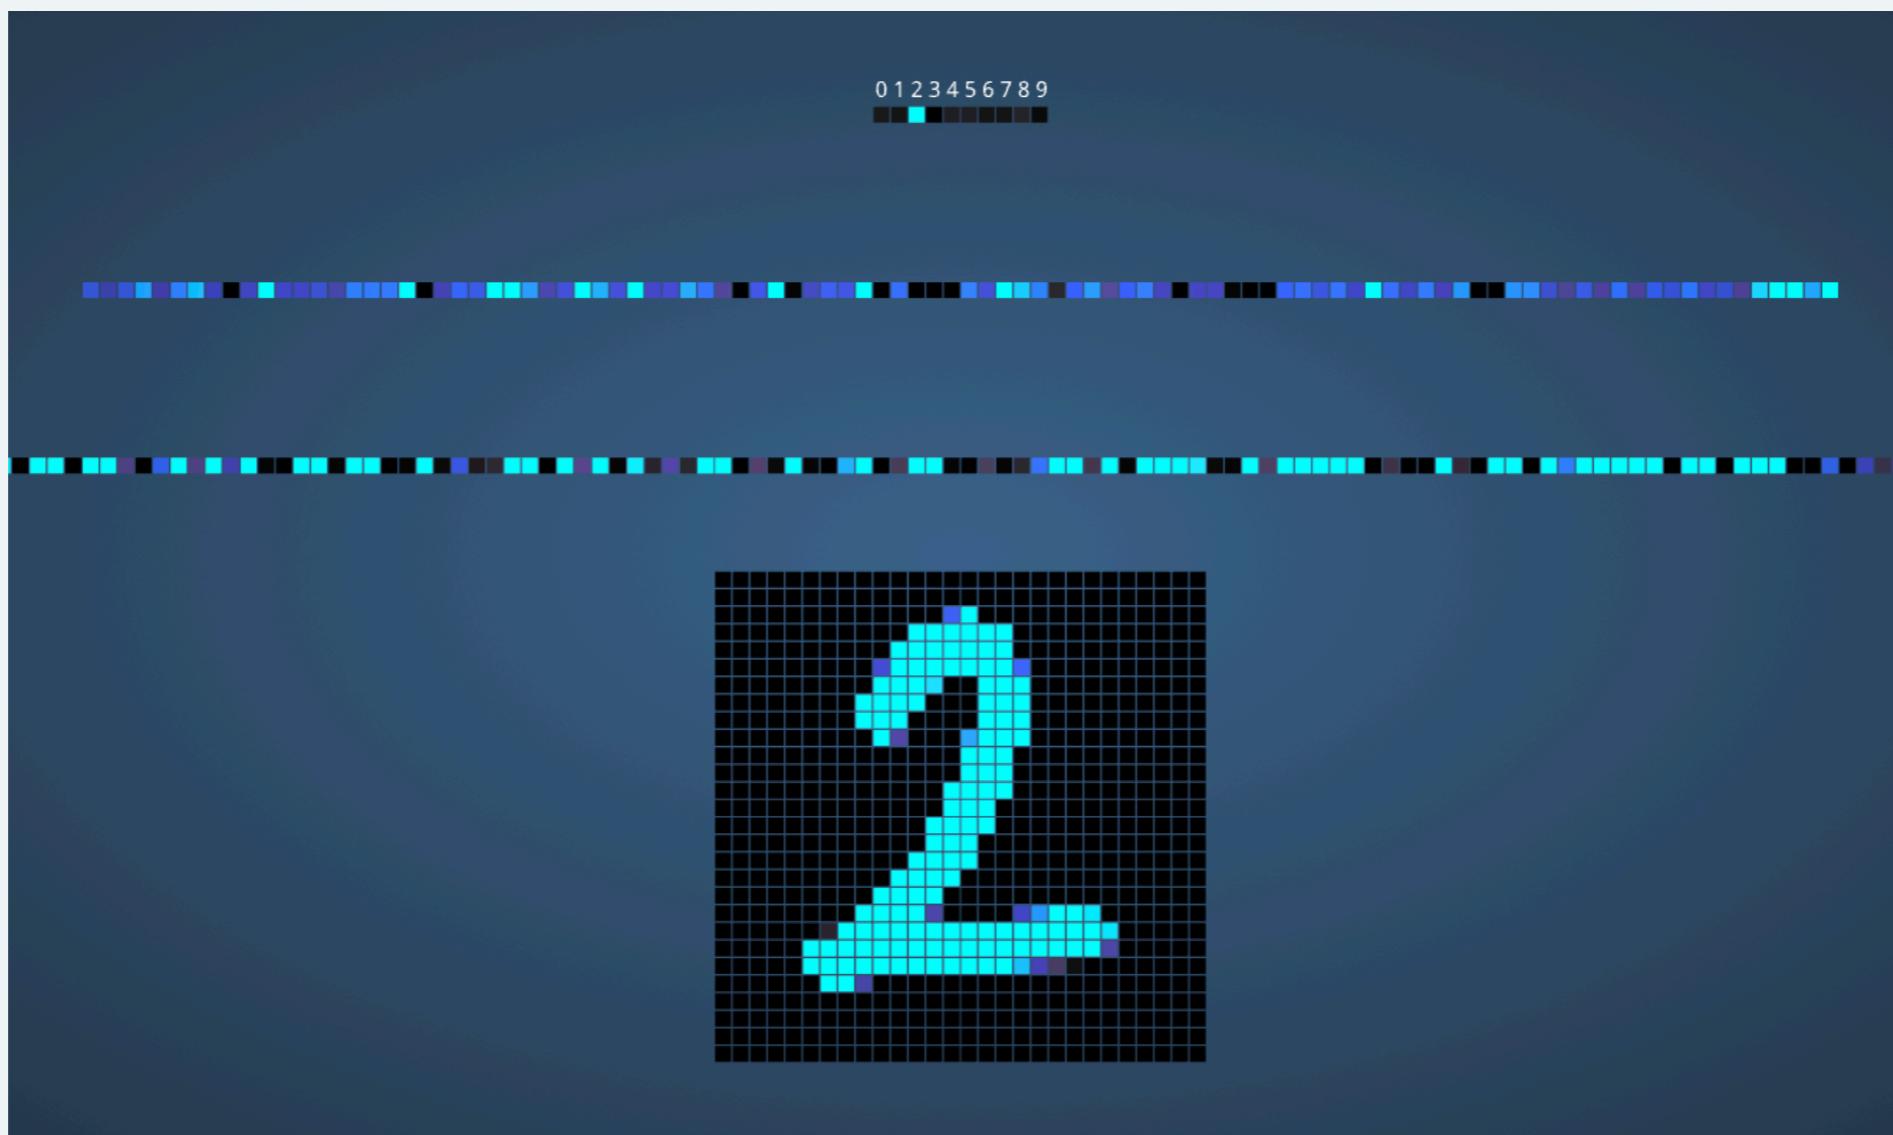
# Subsymbolic representations



# Subsymbolic representations



# Graceful degradation

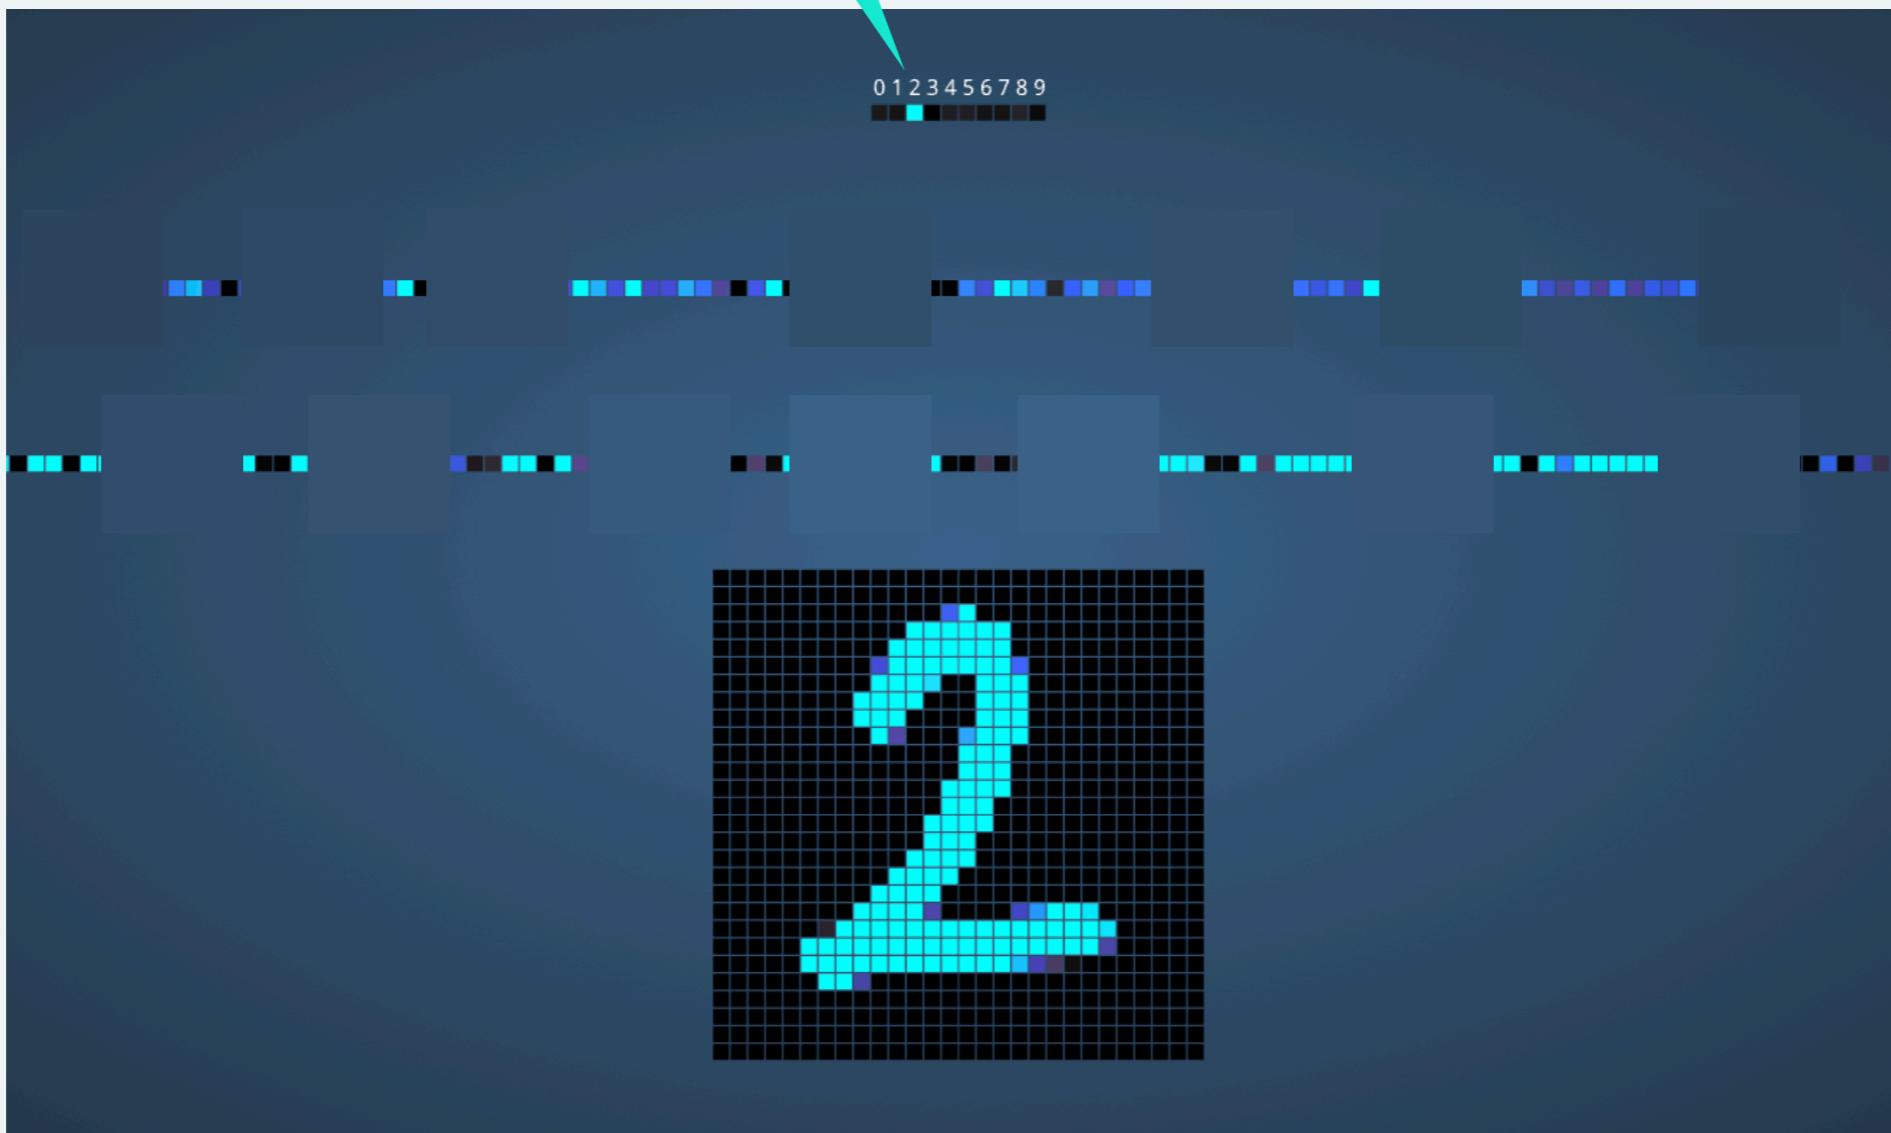


# Graceful degradation

I'm still OK!

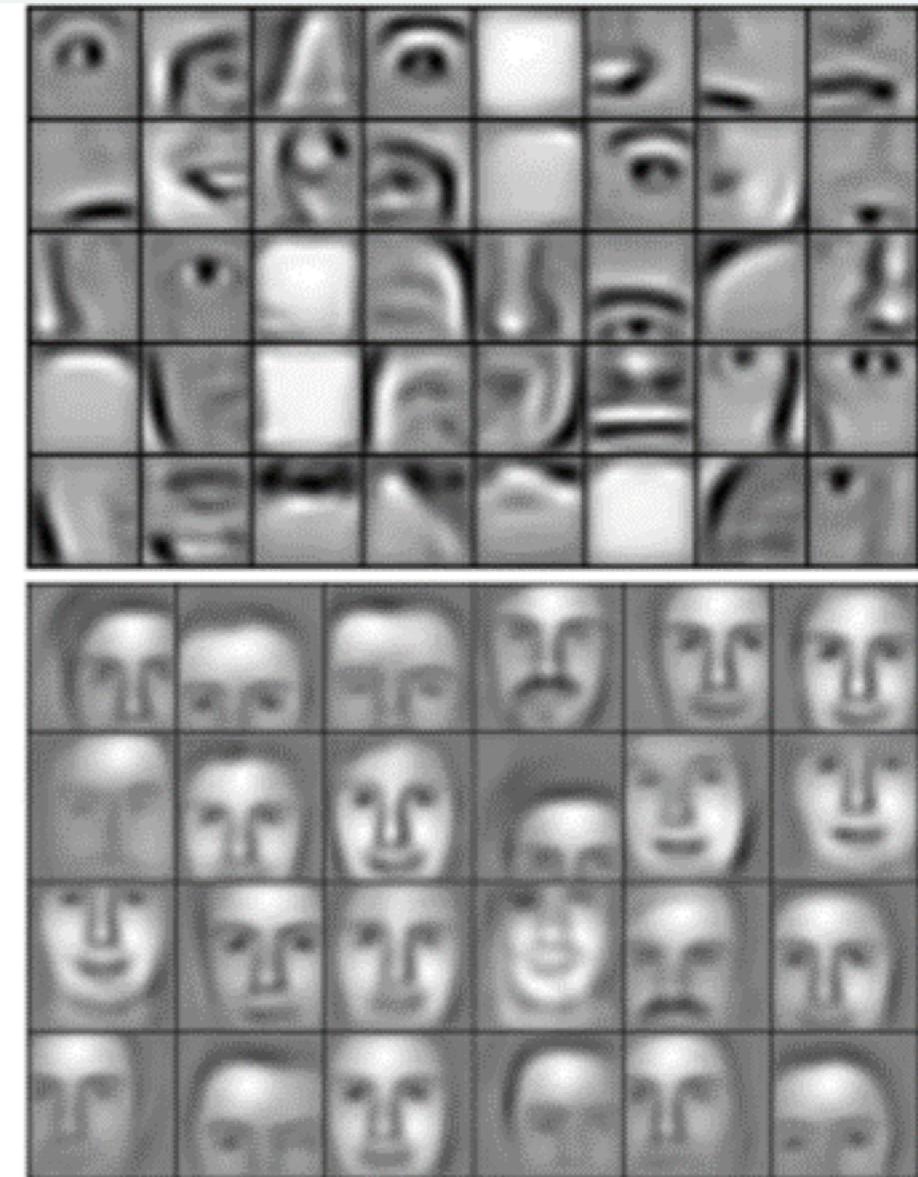
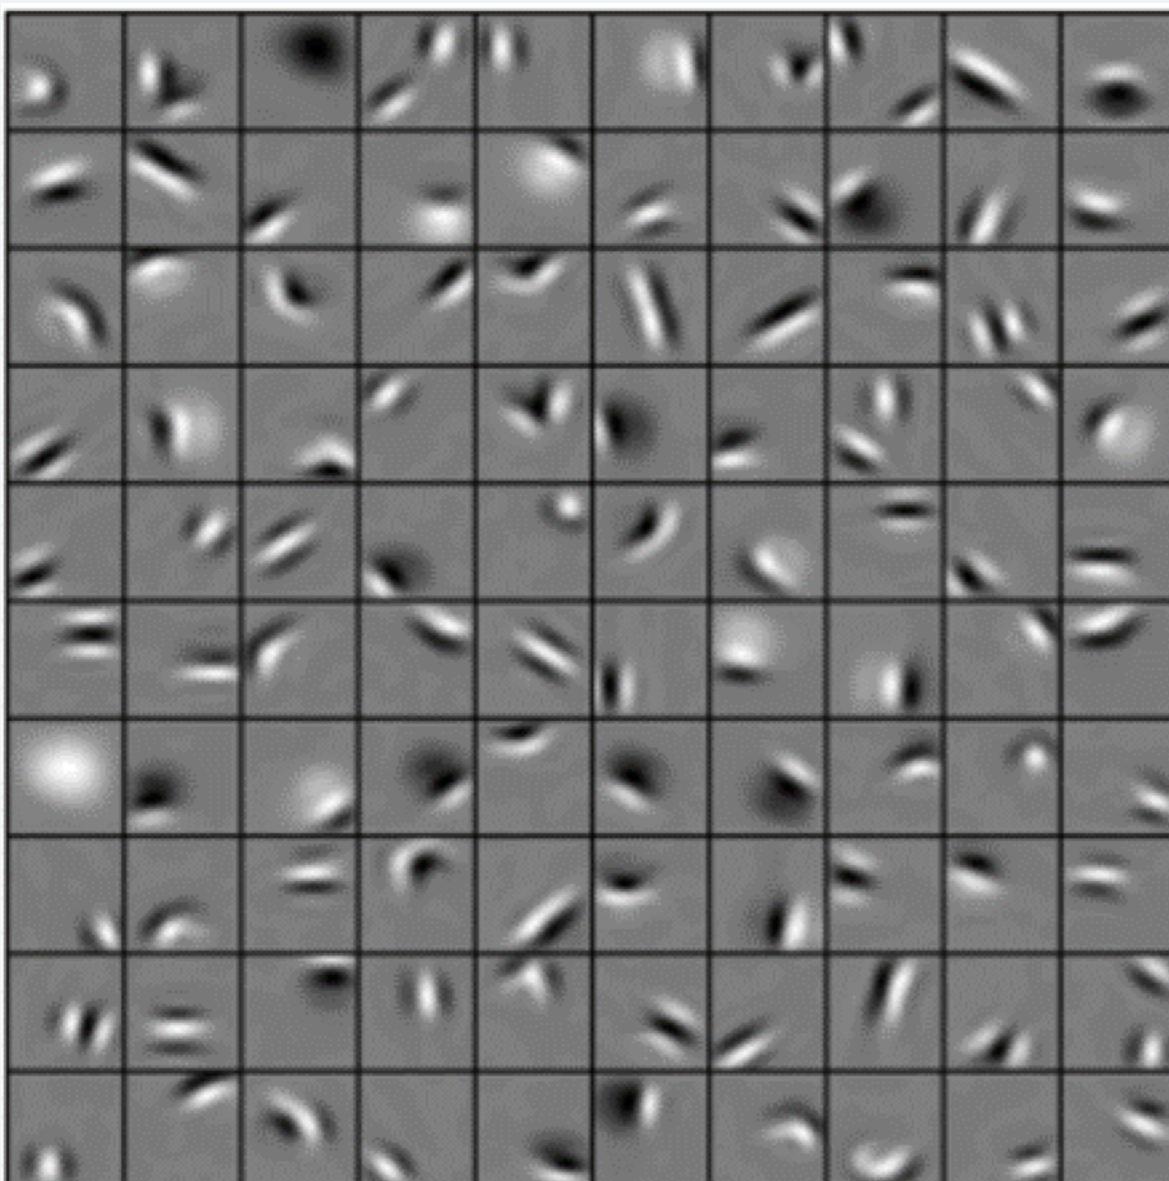
Performance doesn't plummet with removal of a few units

Not the case with PSS



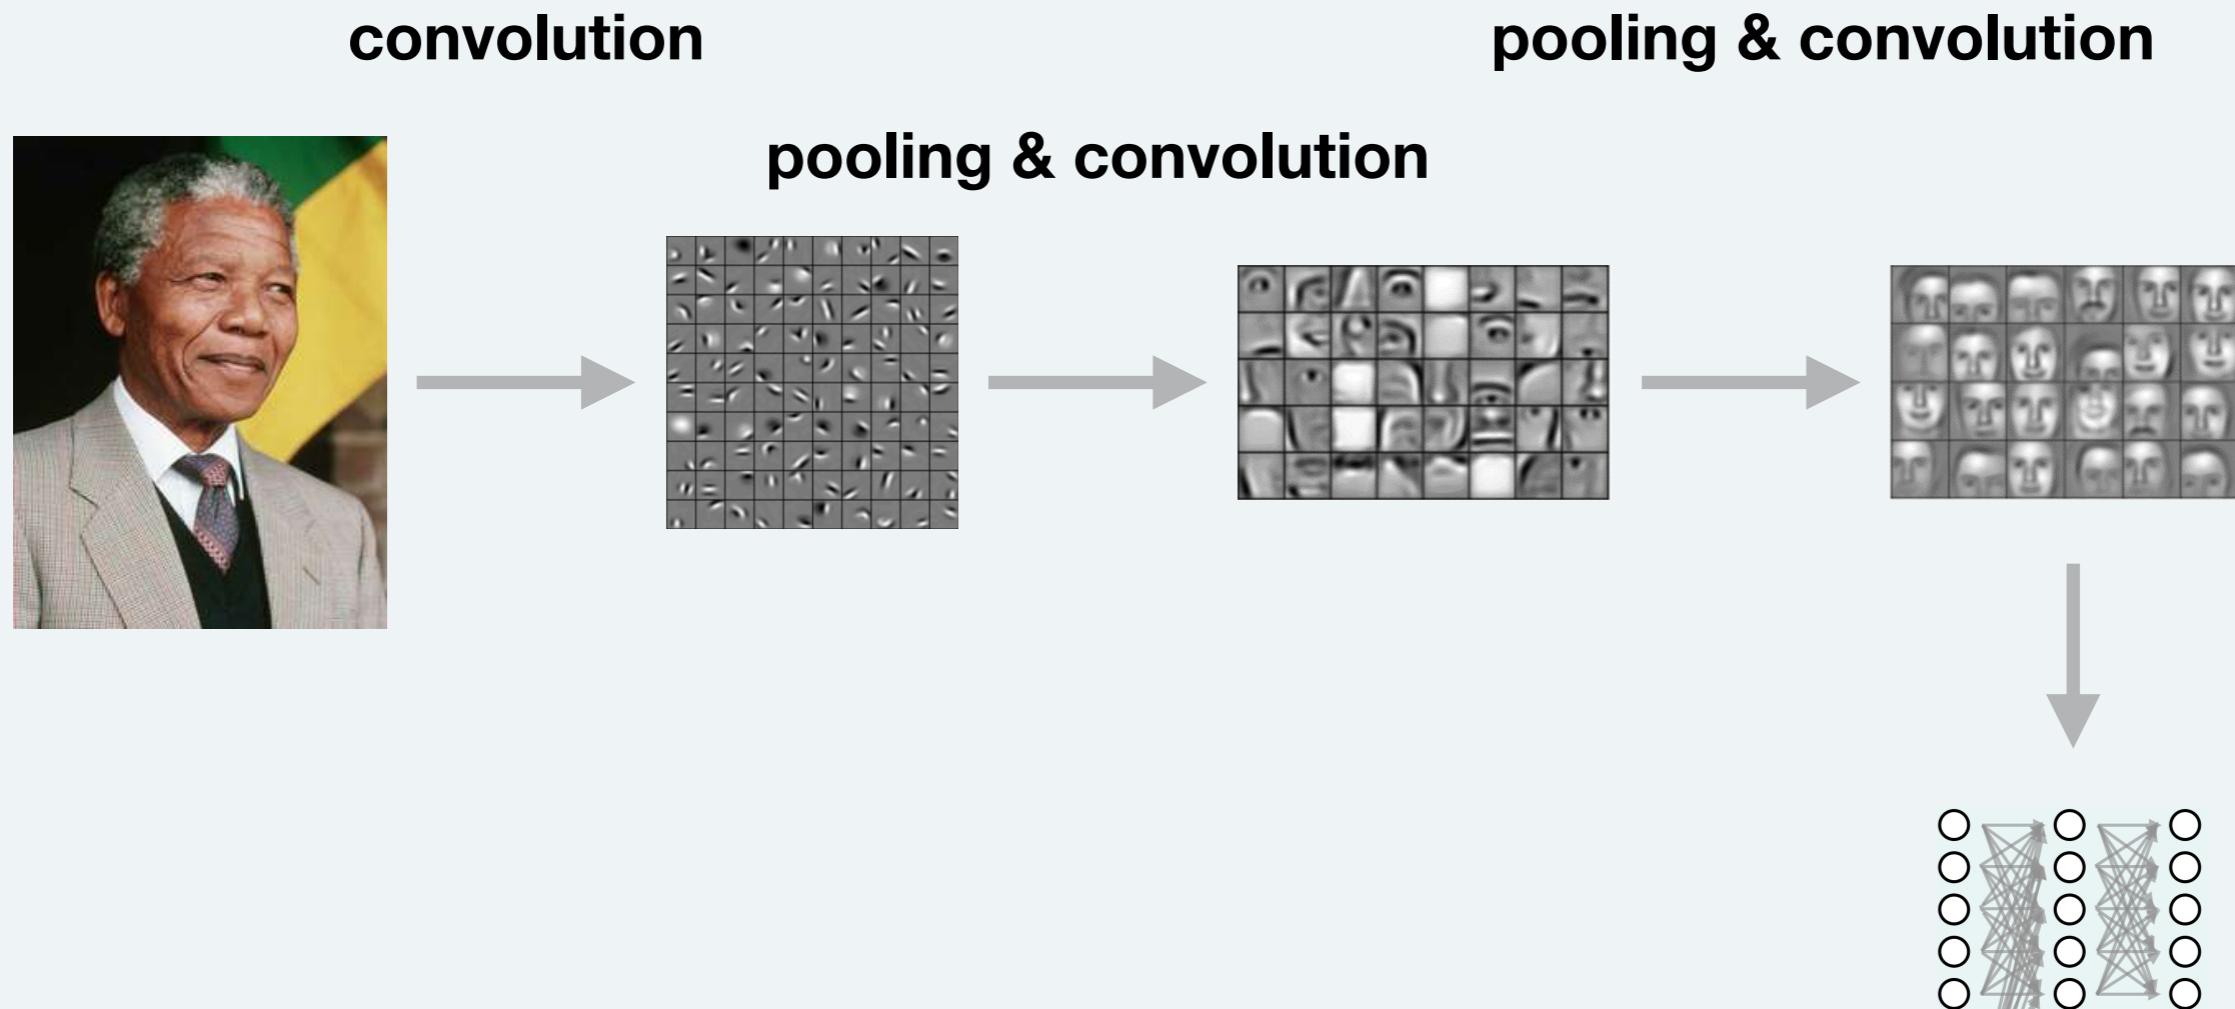
# Connectionist networks

## Convolutional Network



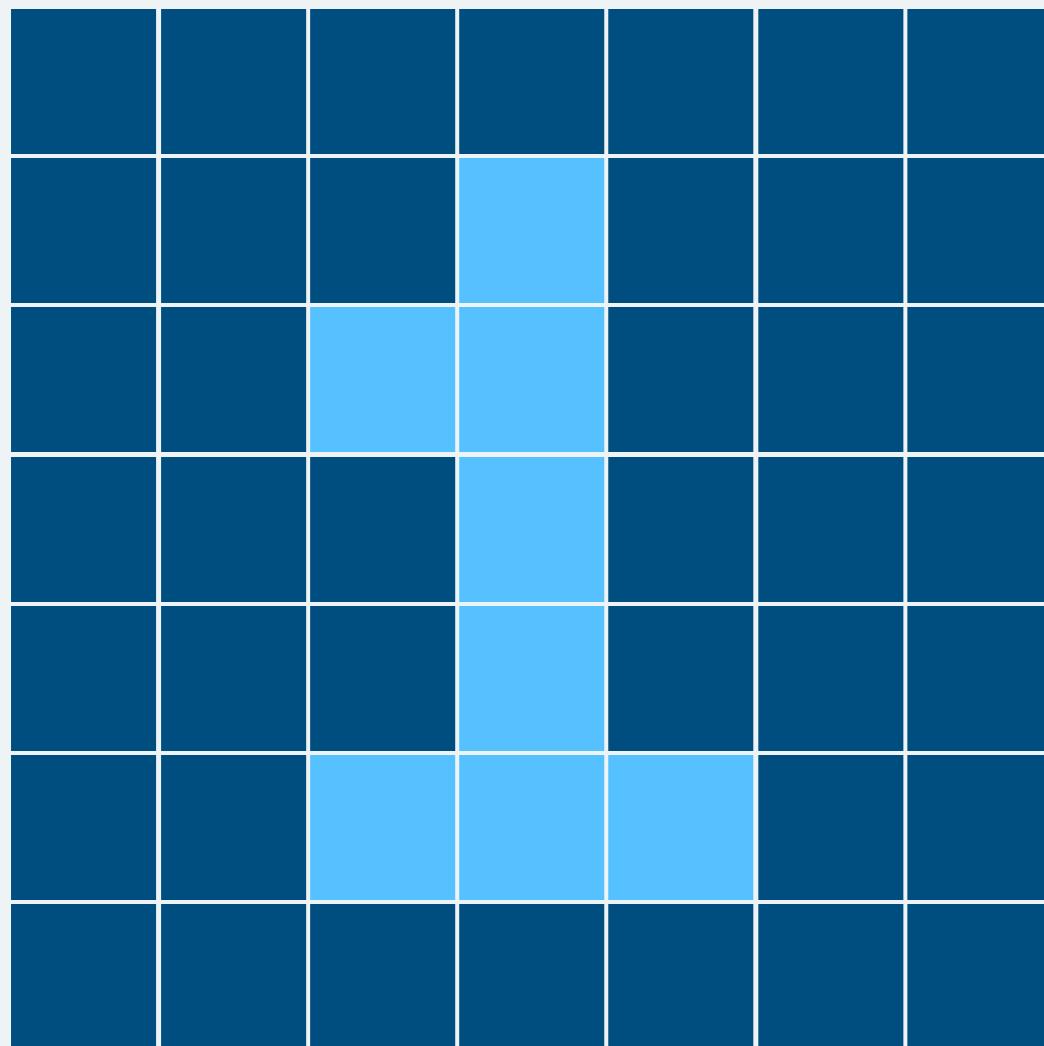
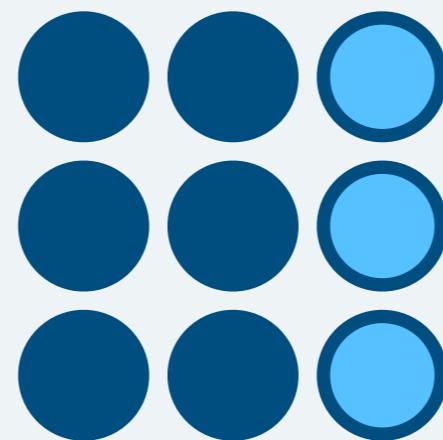
# Connectionist networks

## Convolutional Network

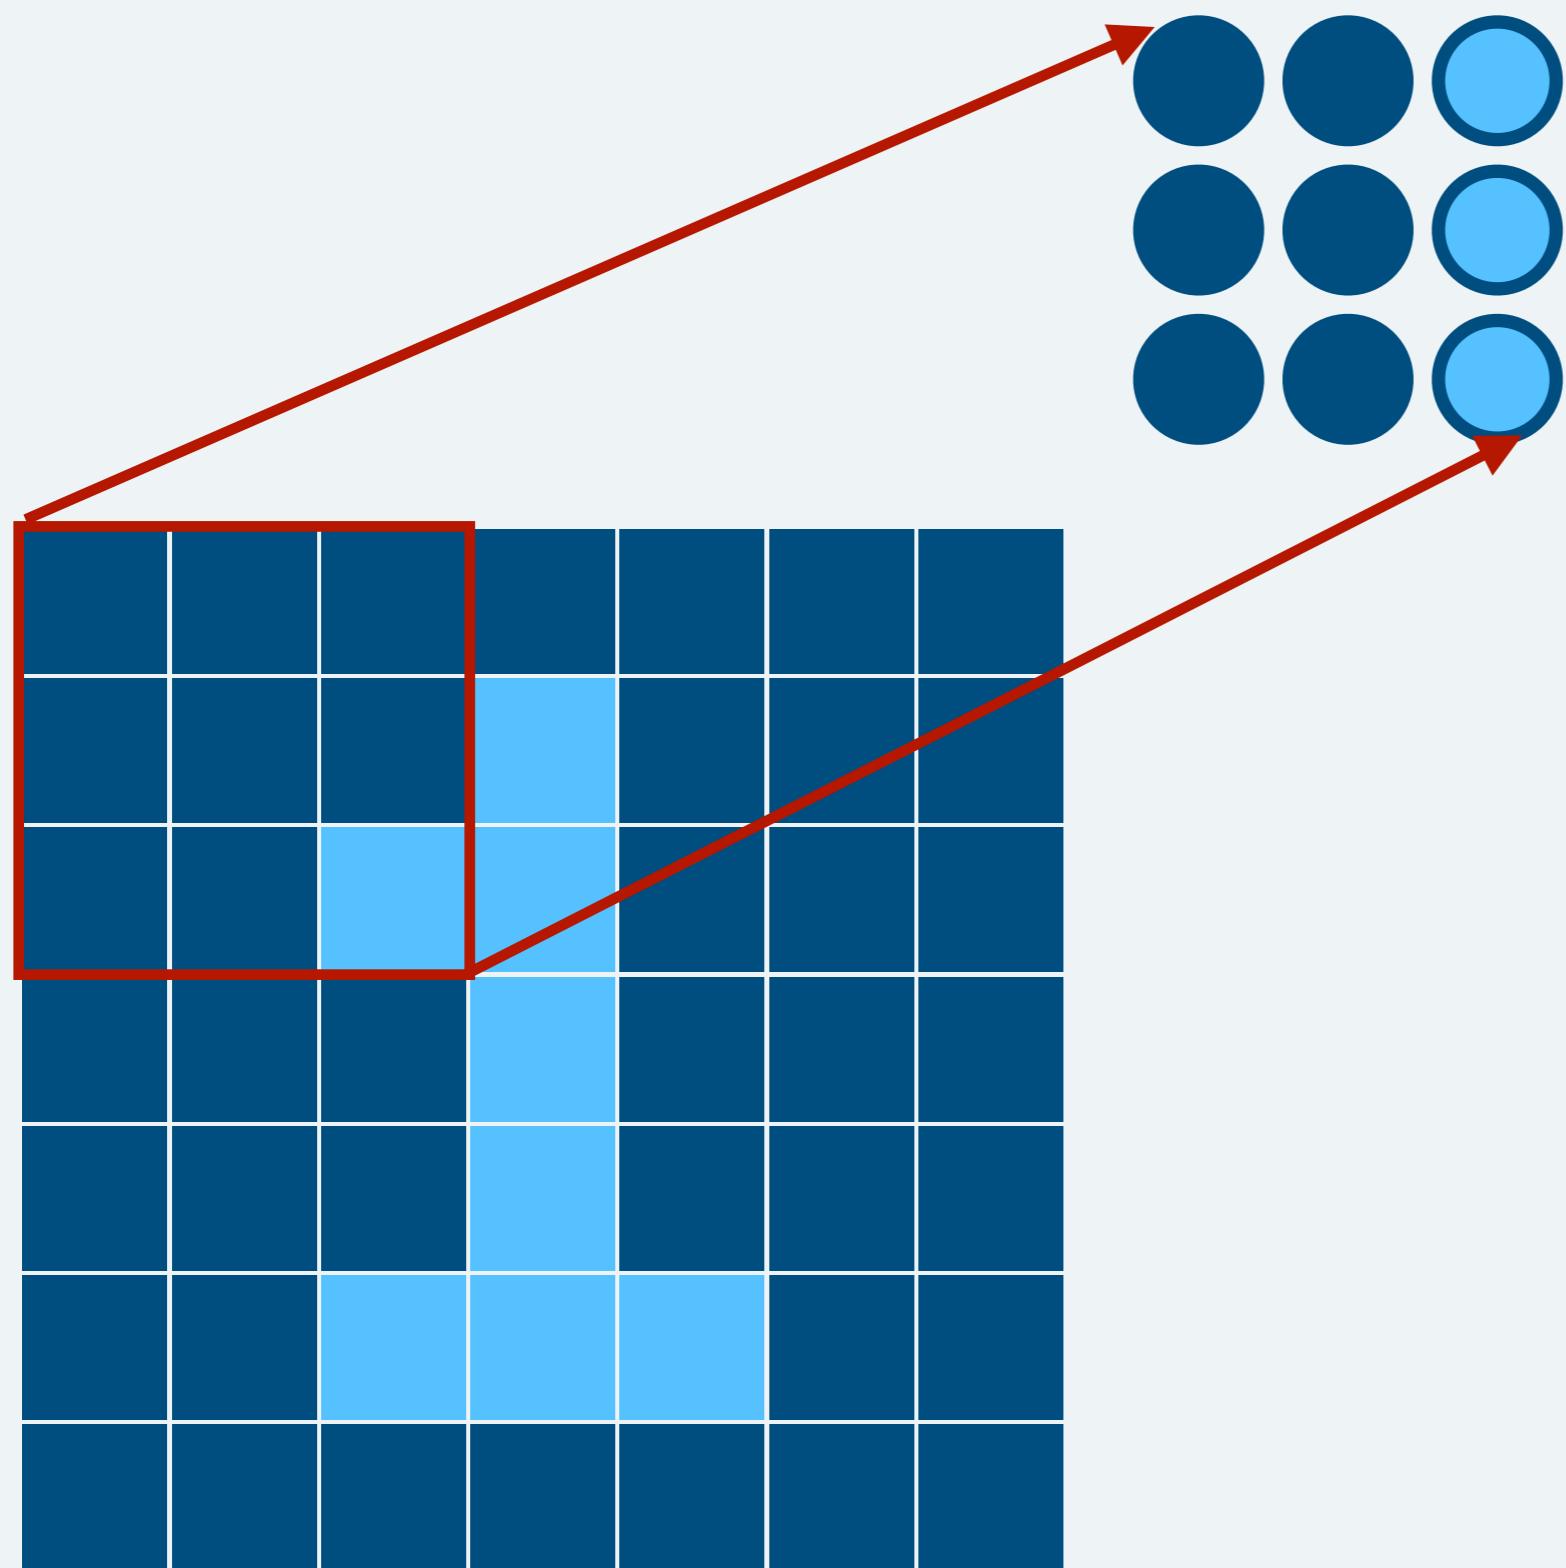


<https://www.youtube.com/watch?v=FmpDlaiMleA>

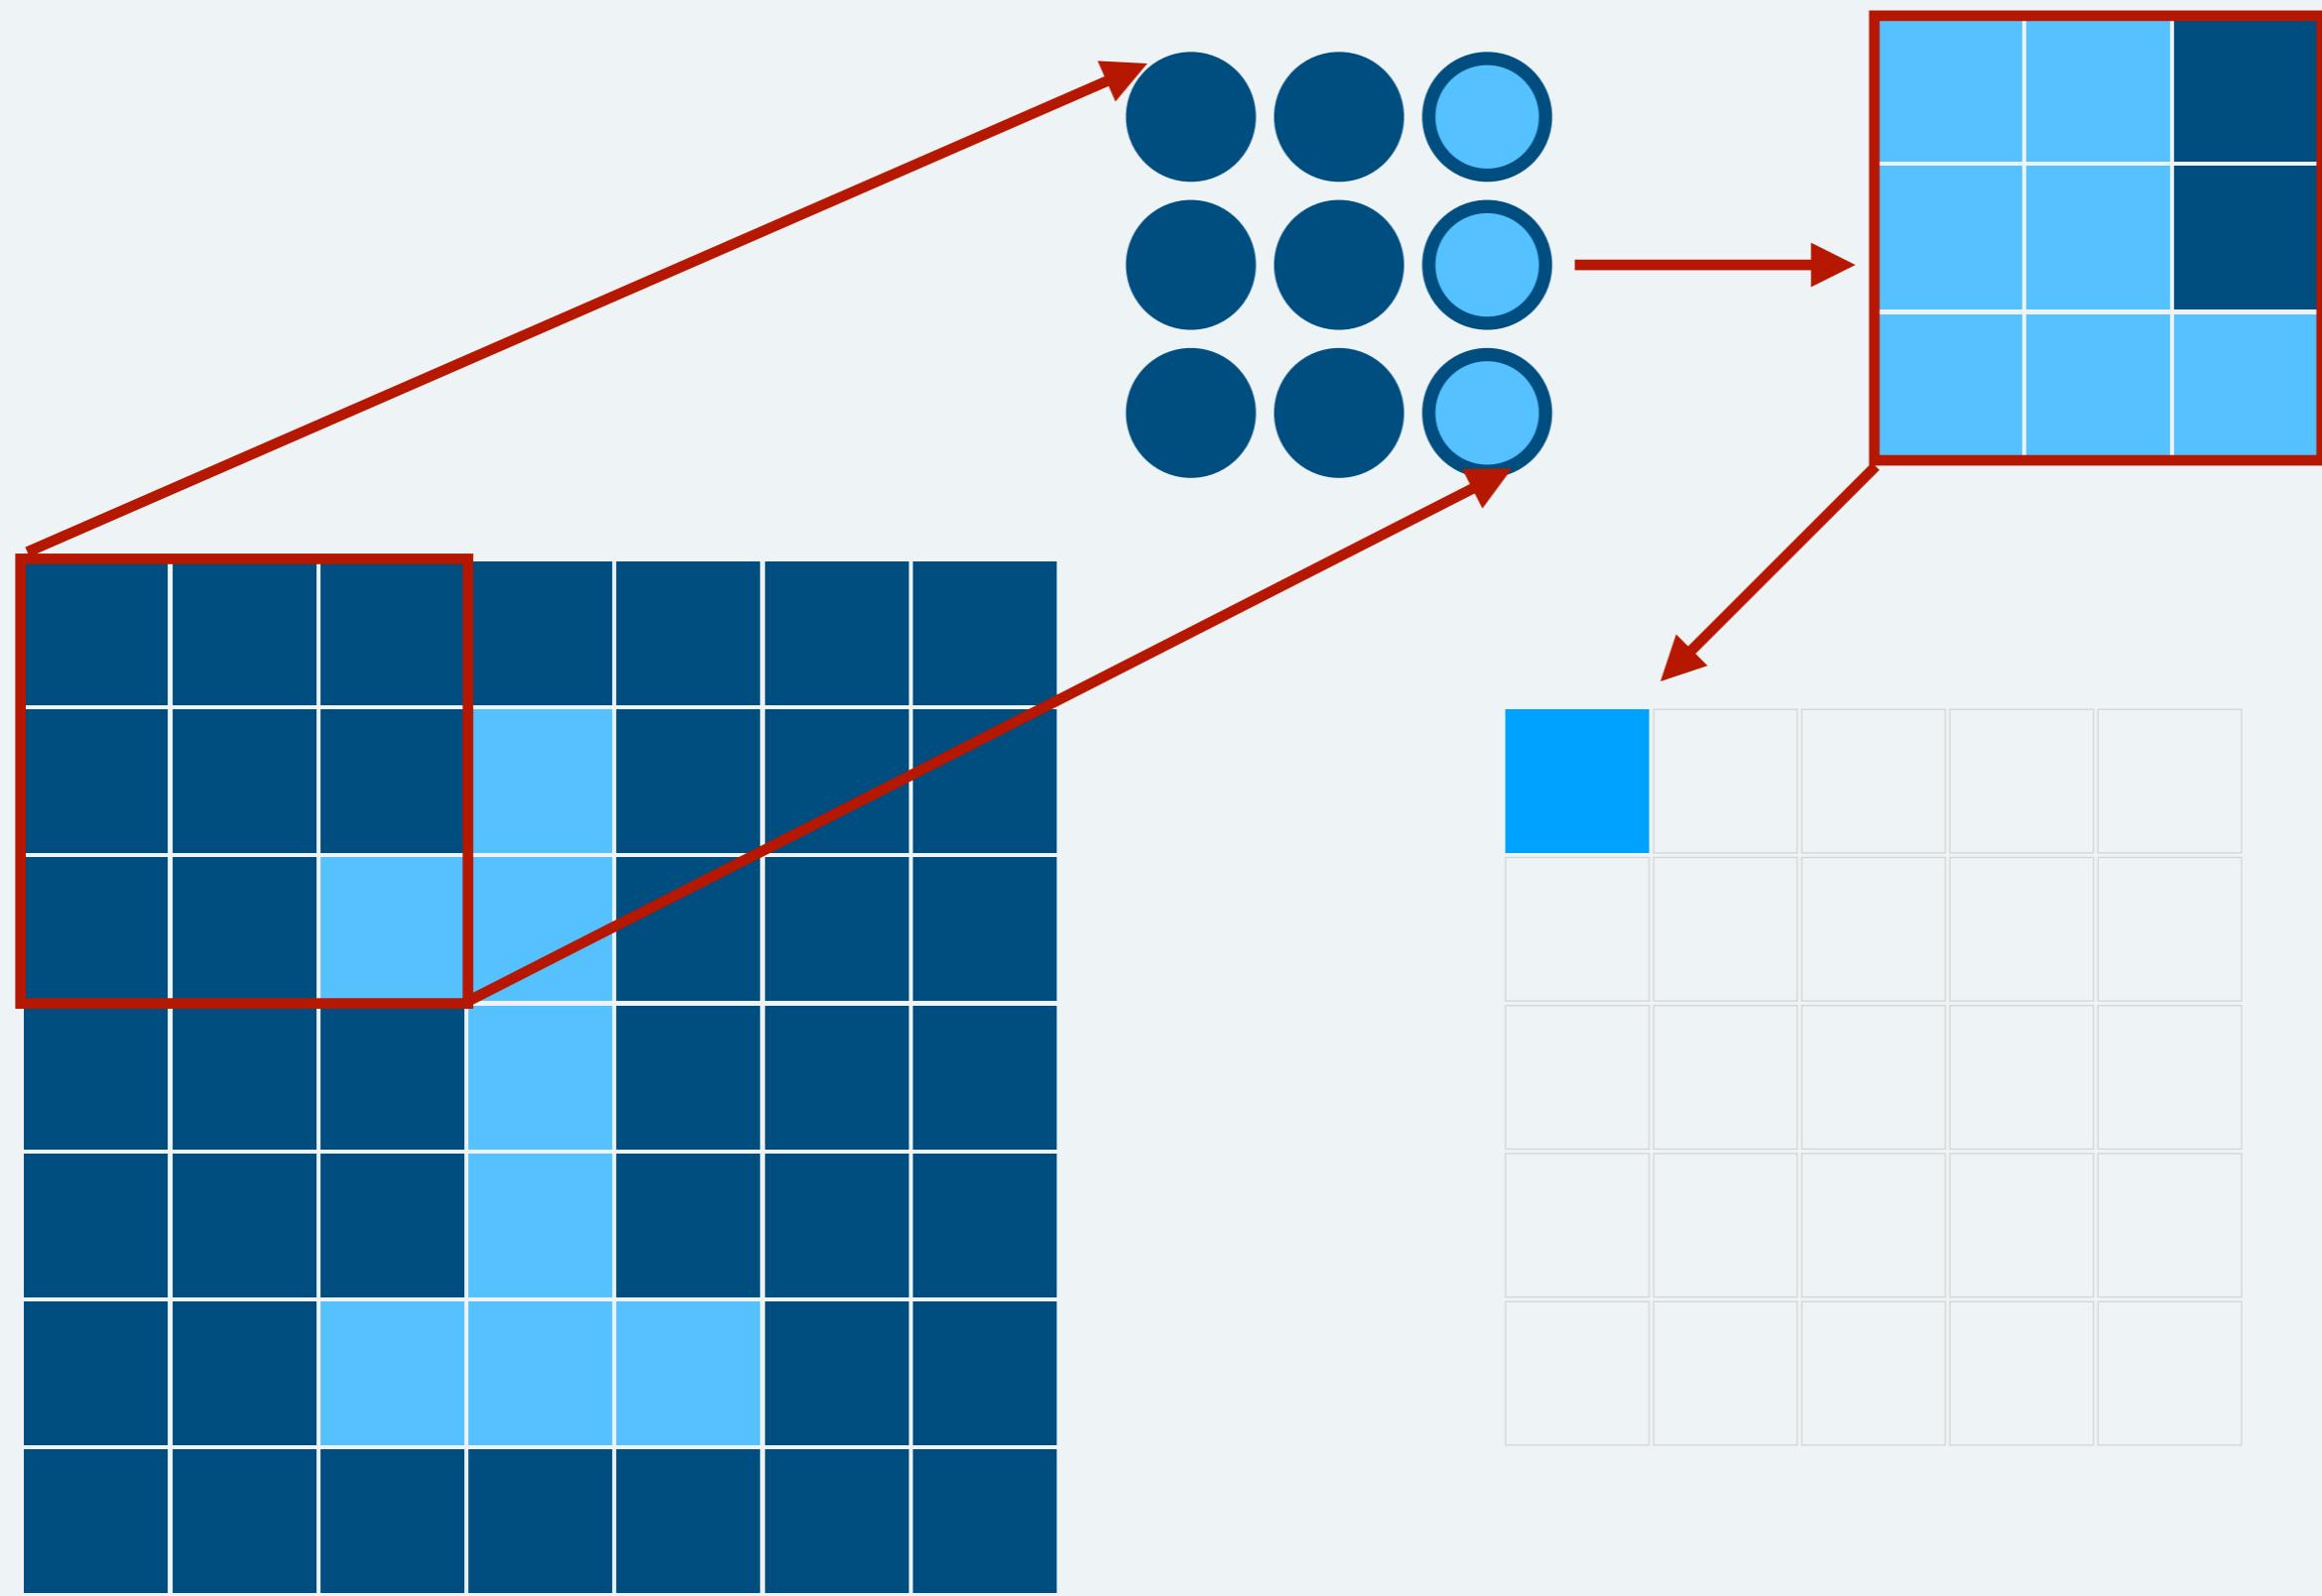
# Convolution



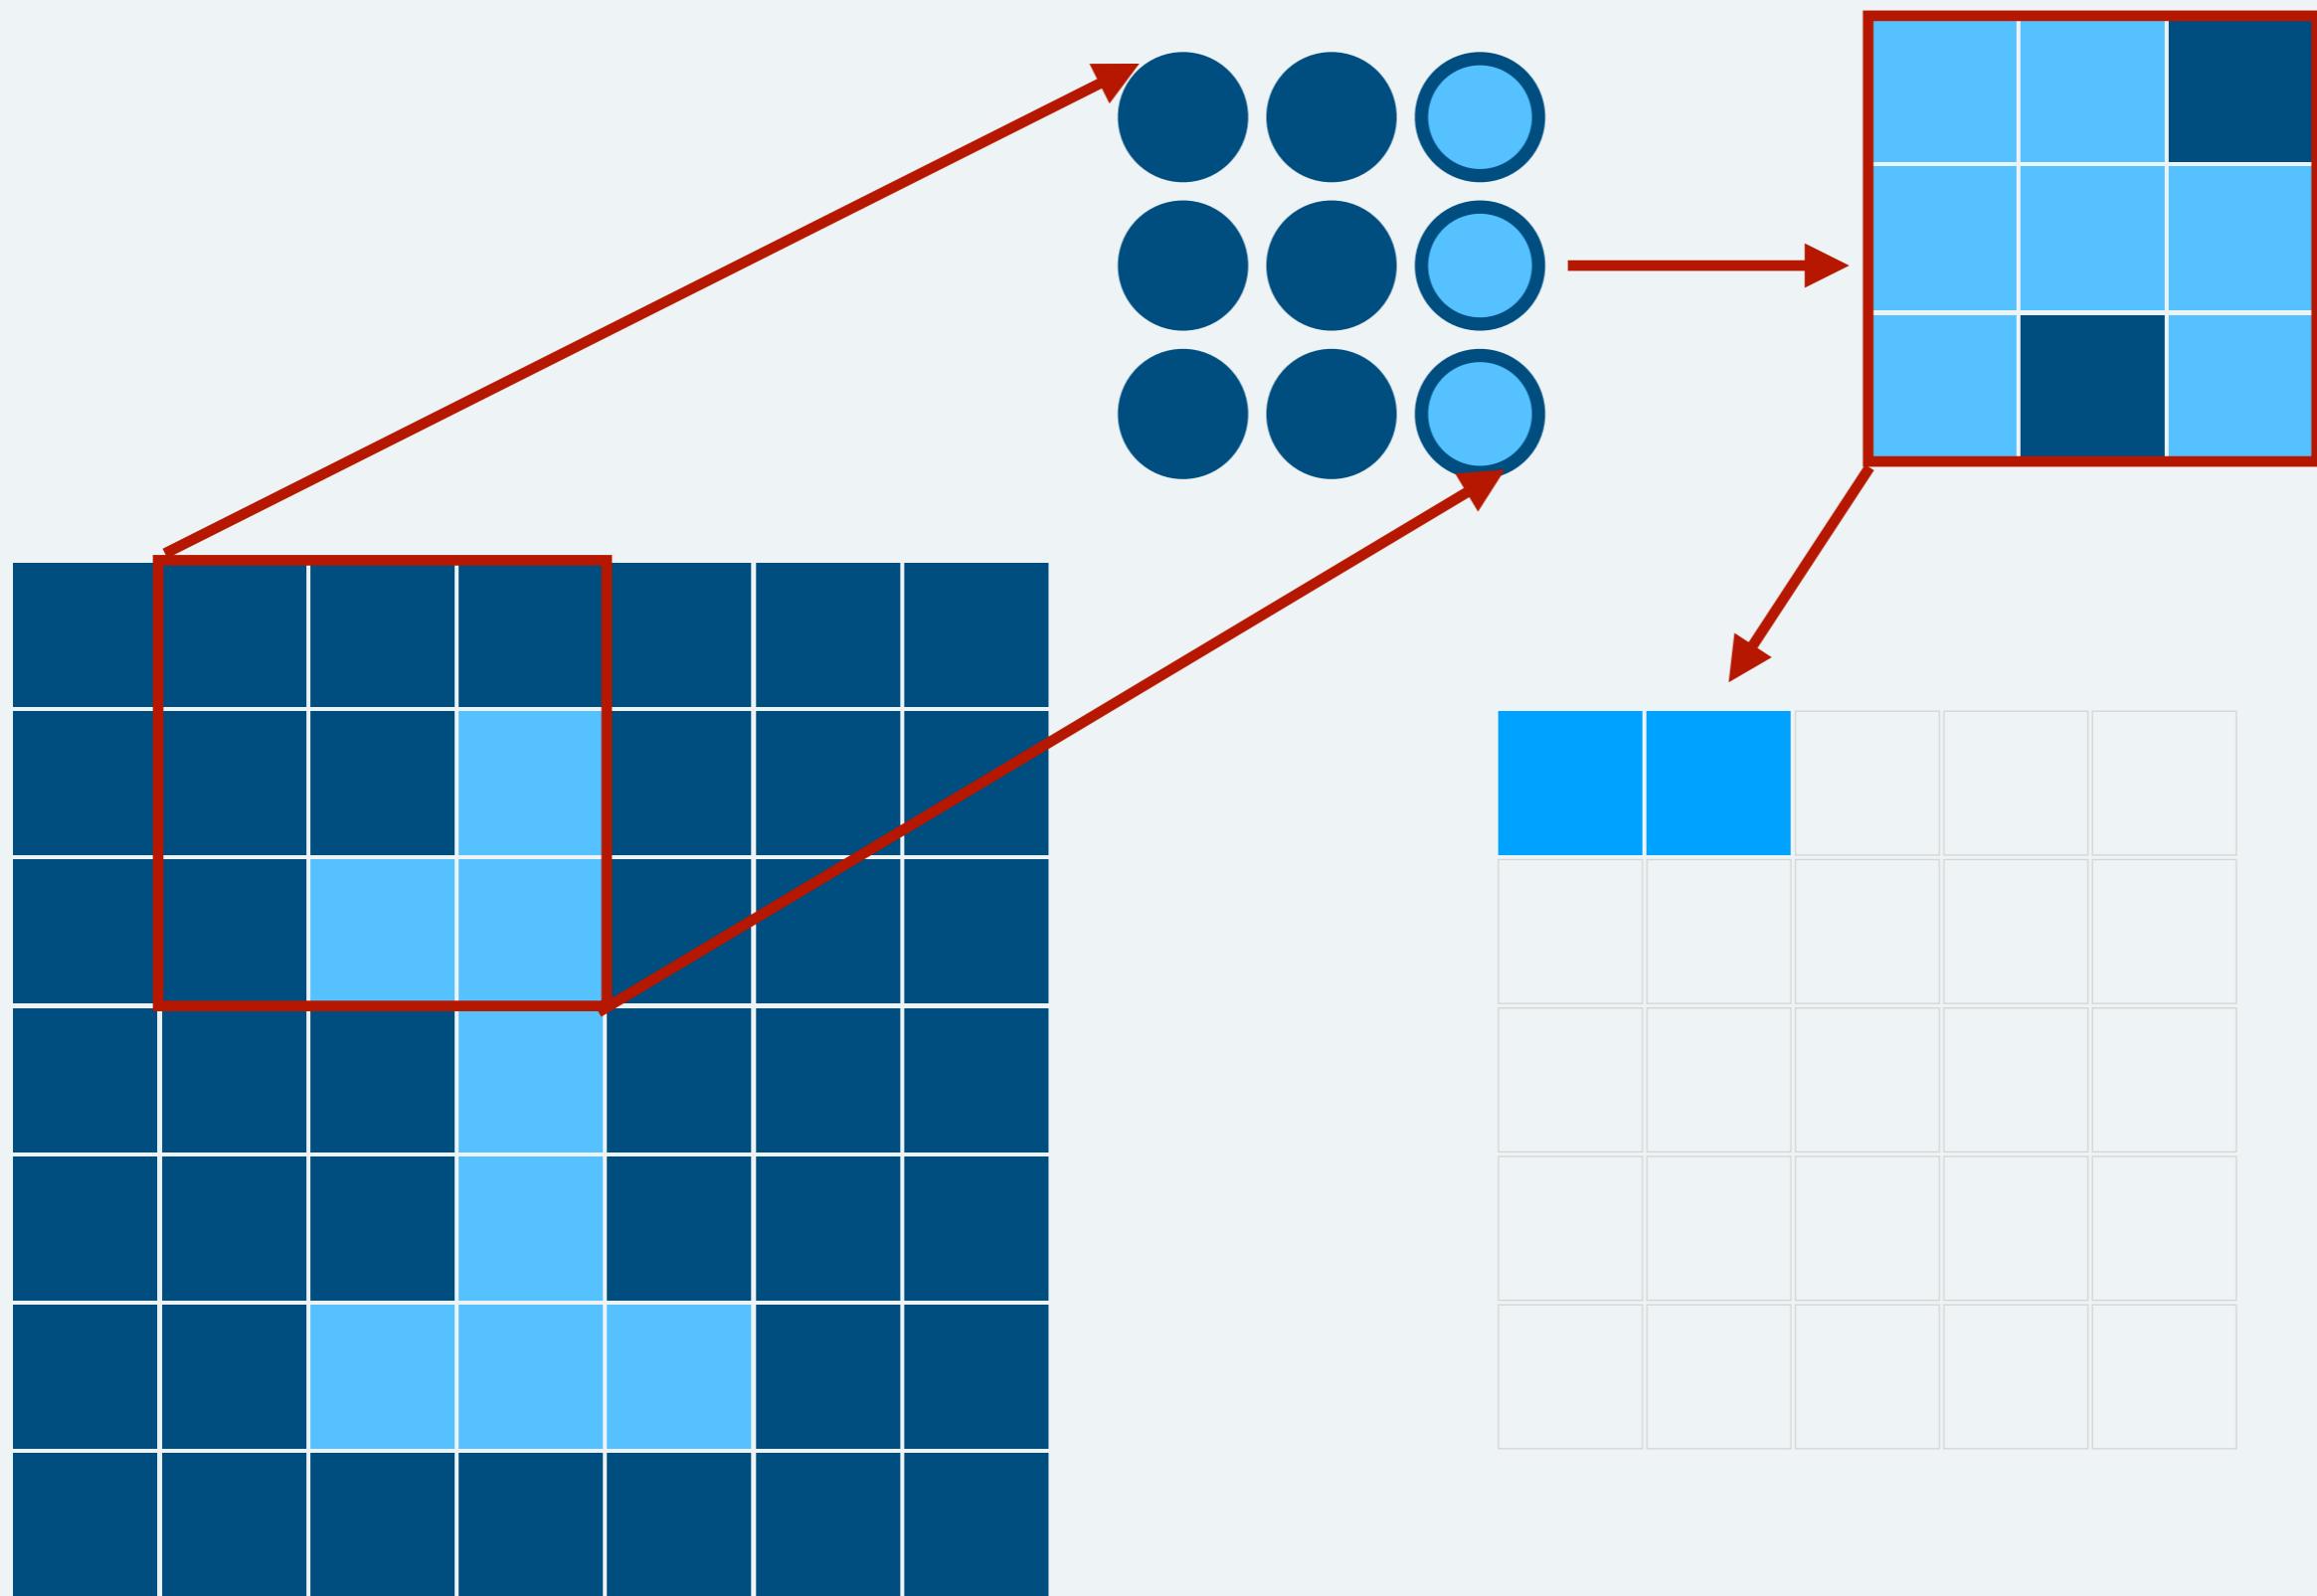
# Convolution



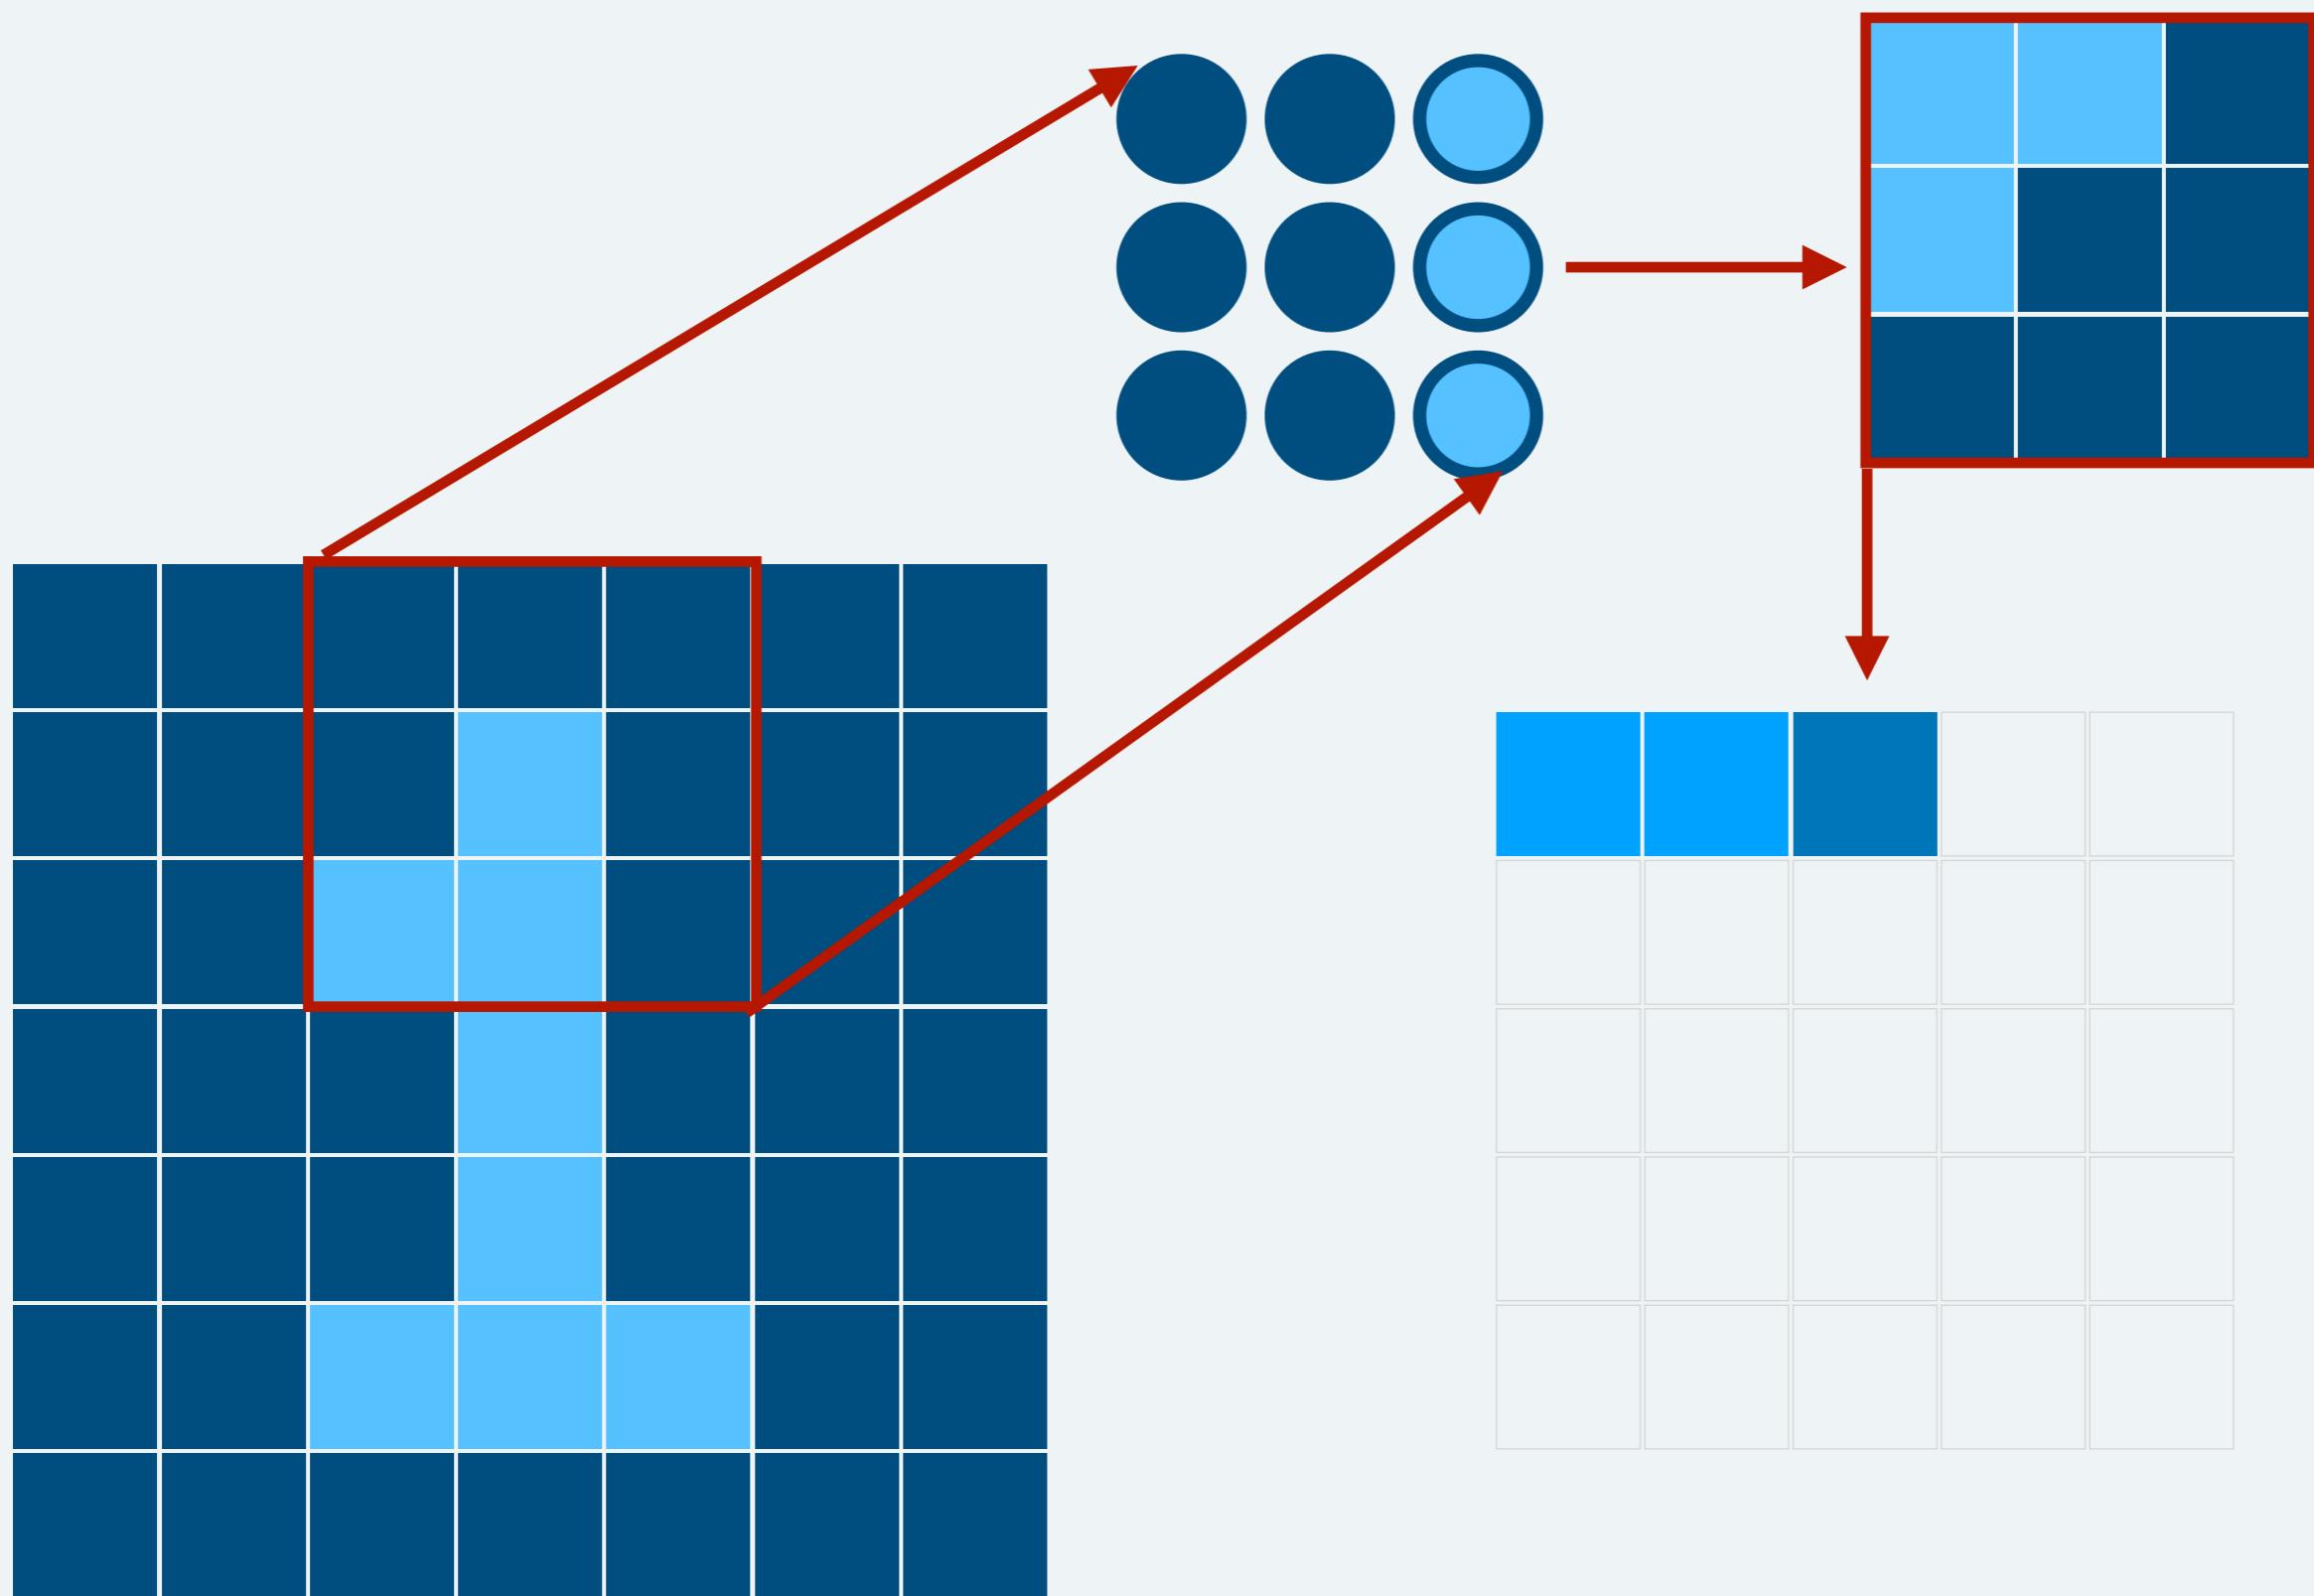
# Convolution



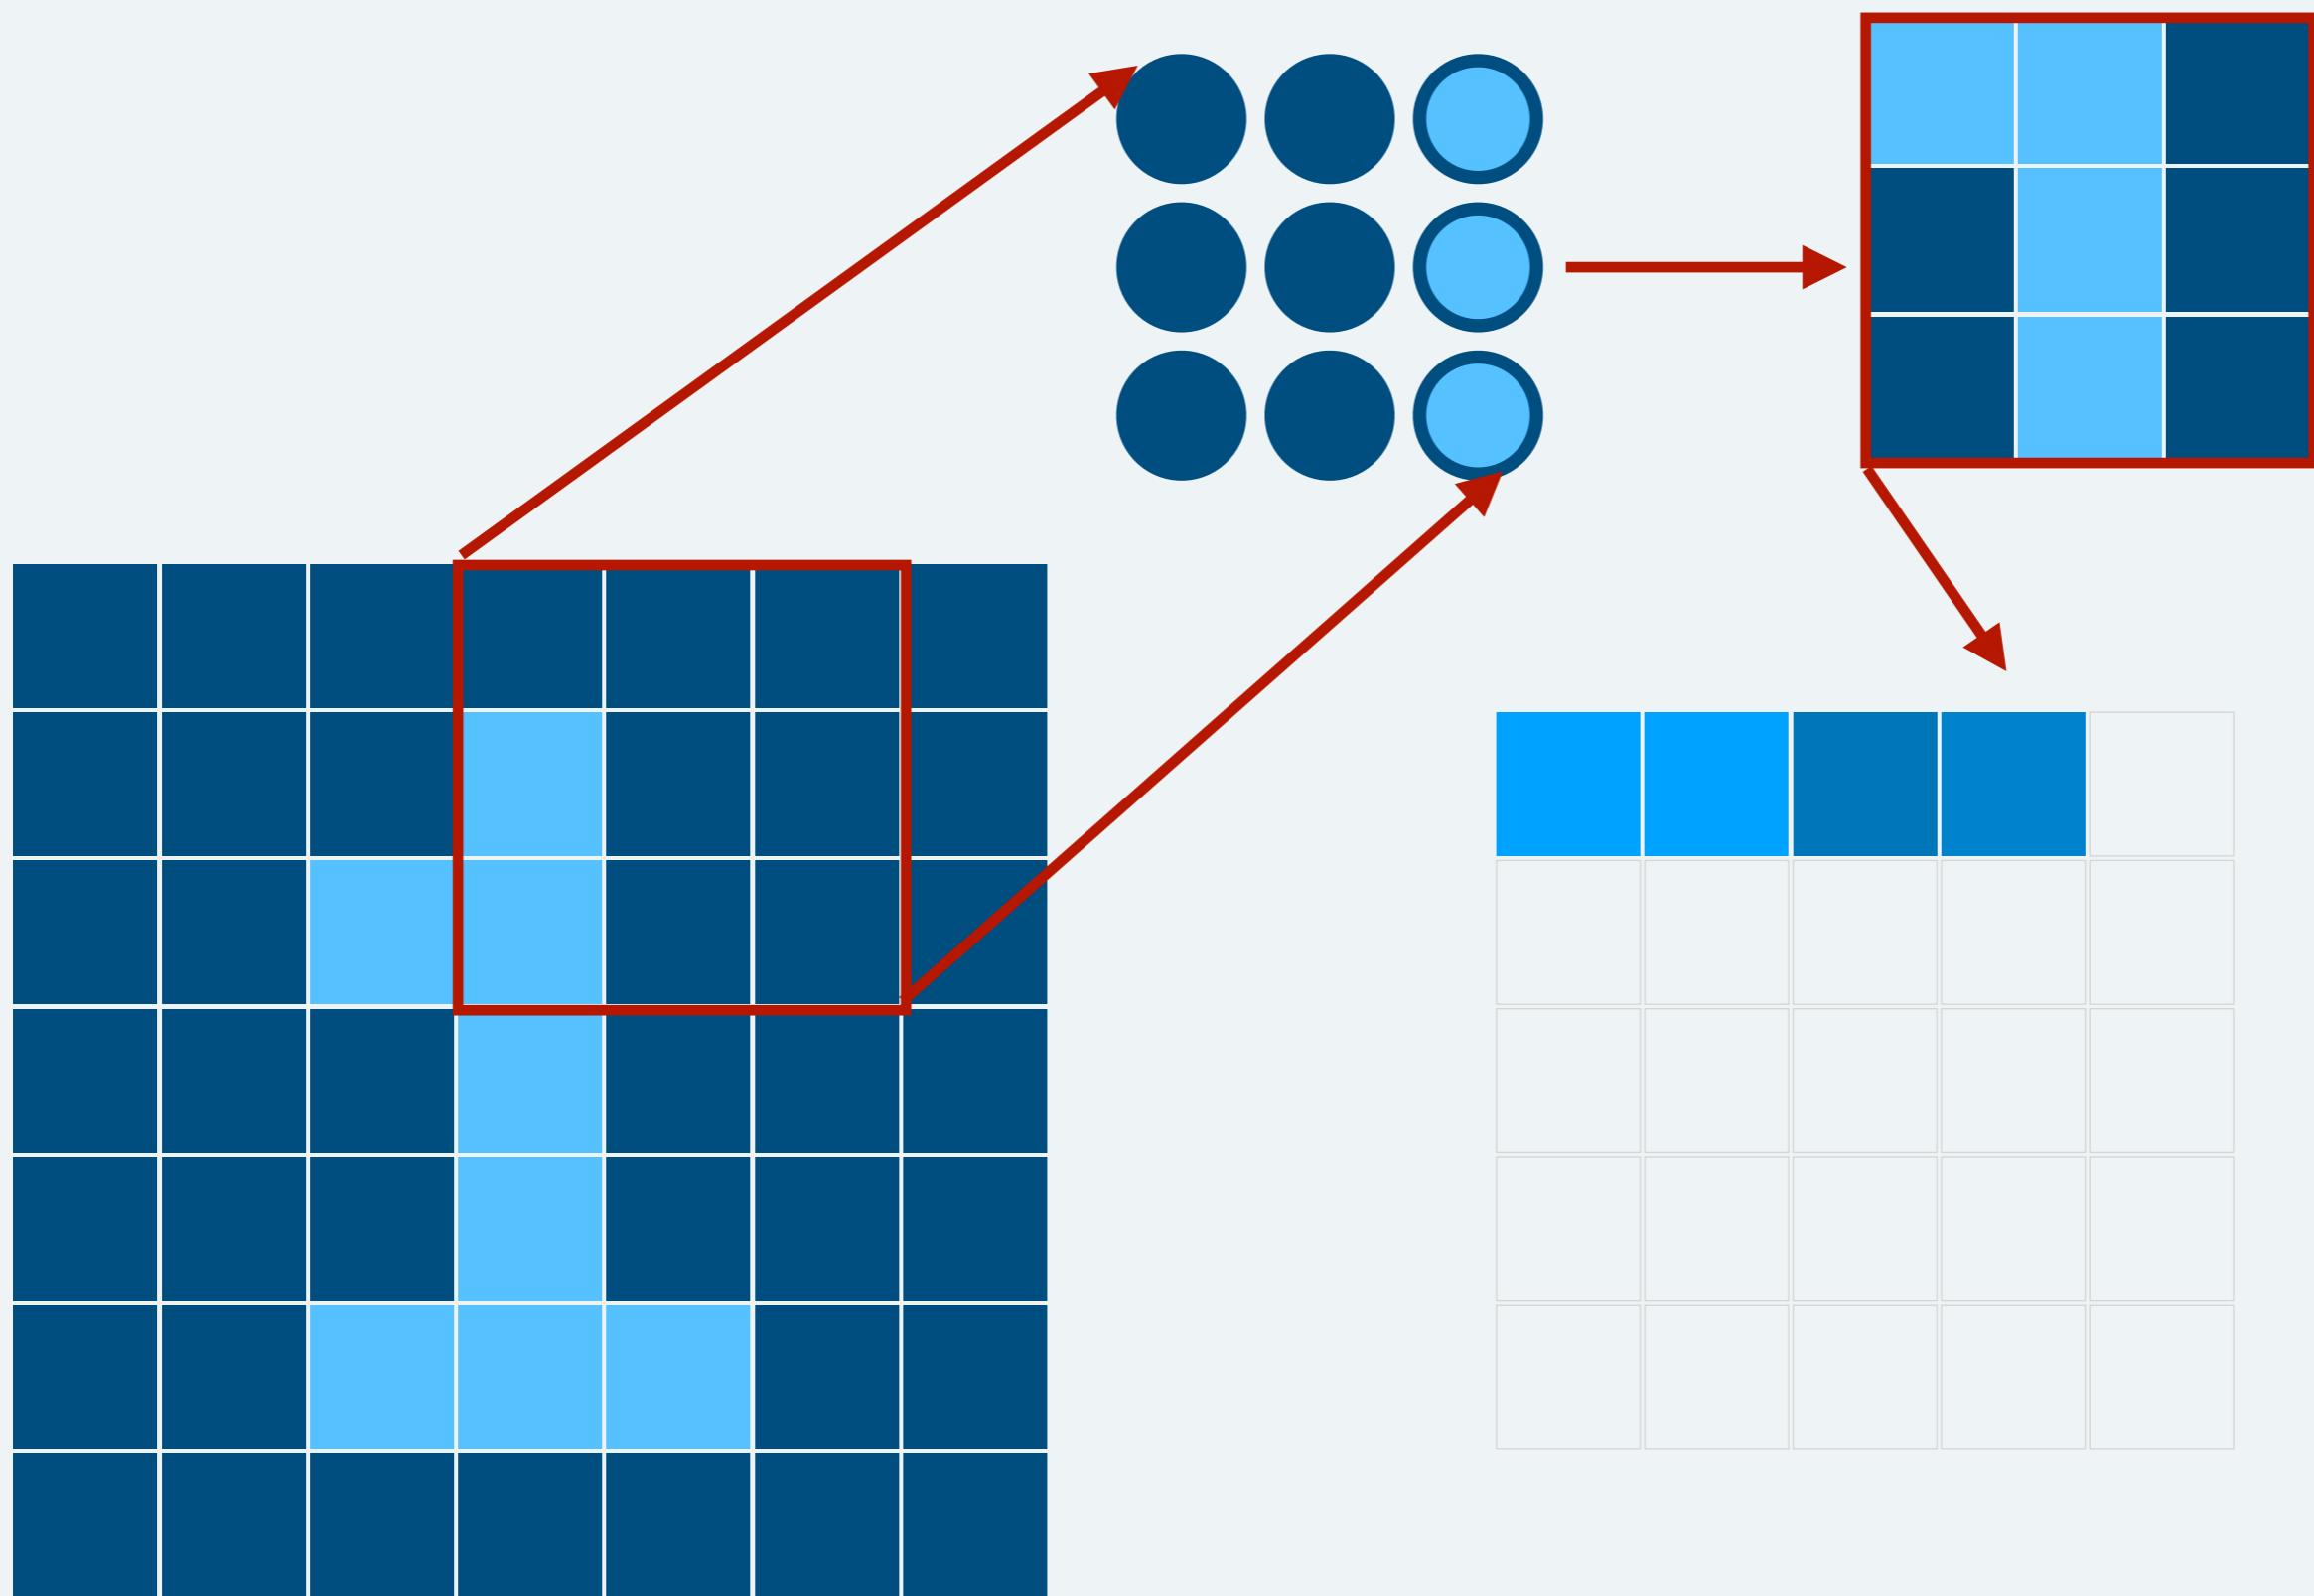
# Convolution



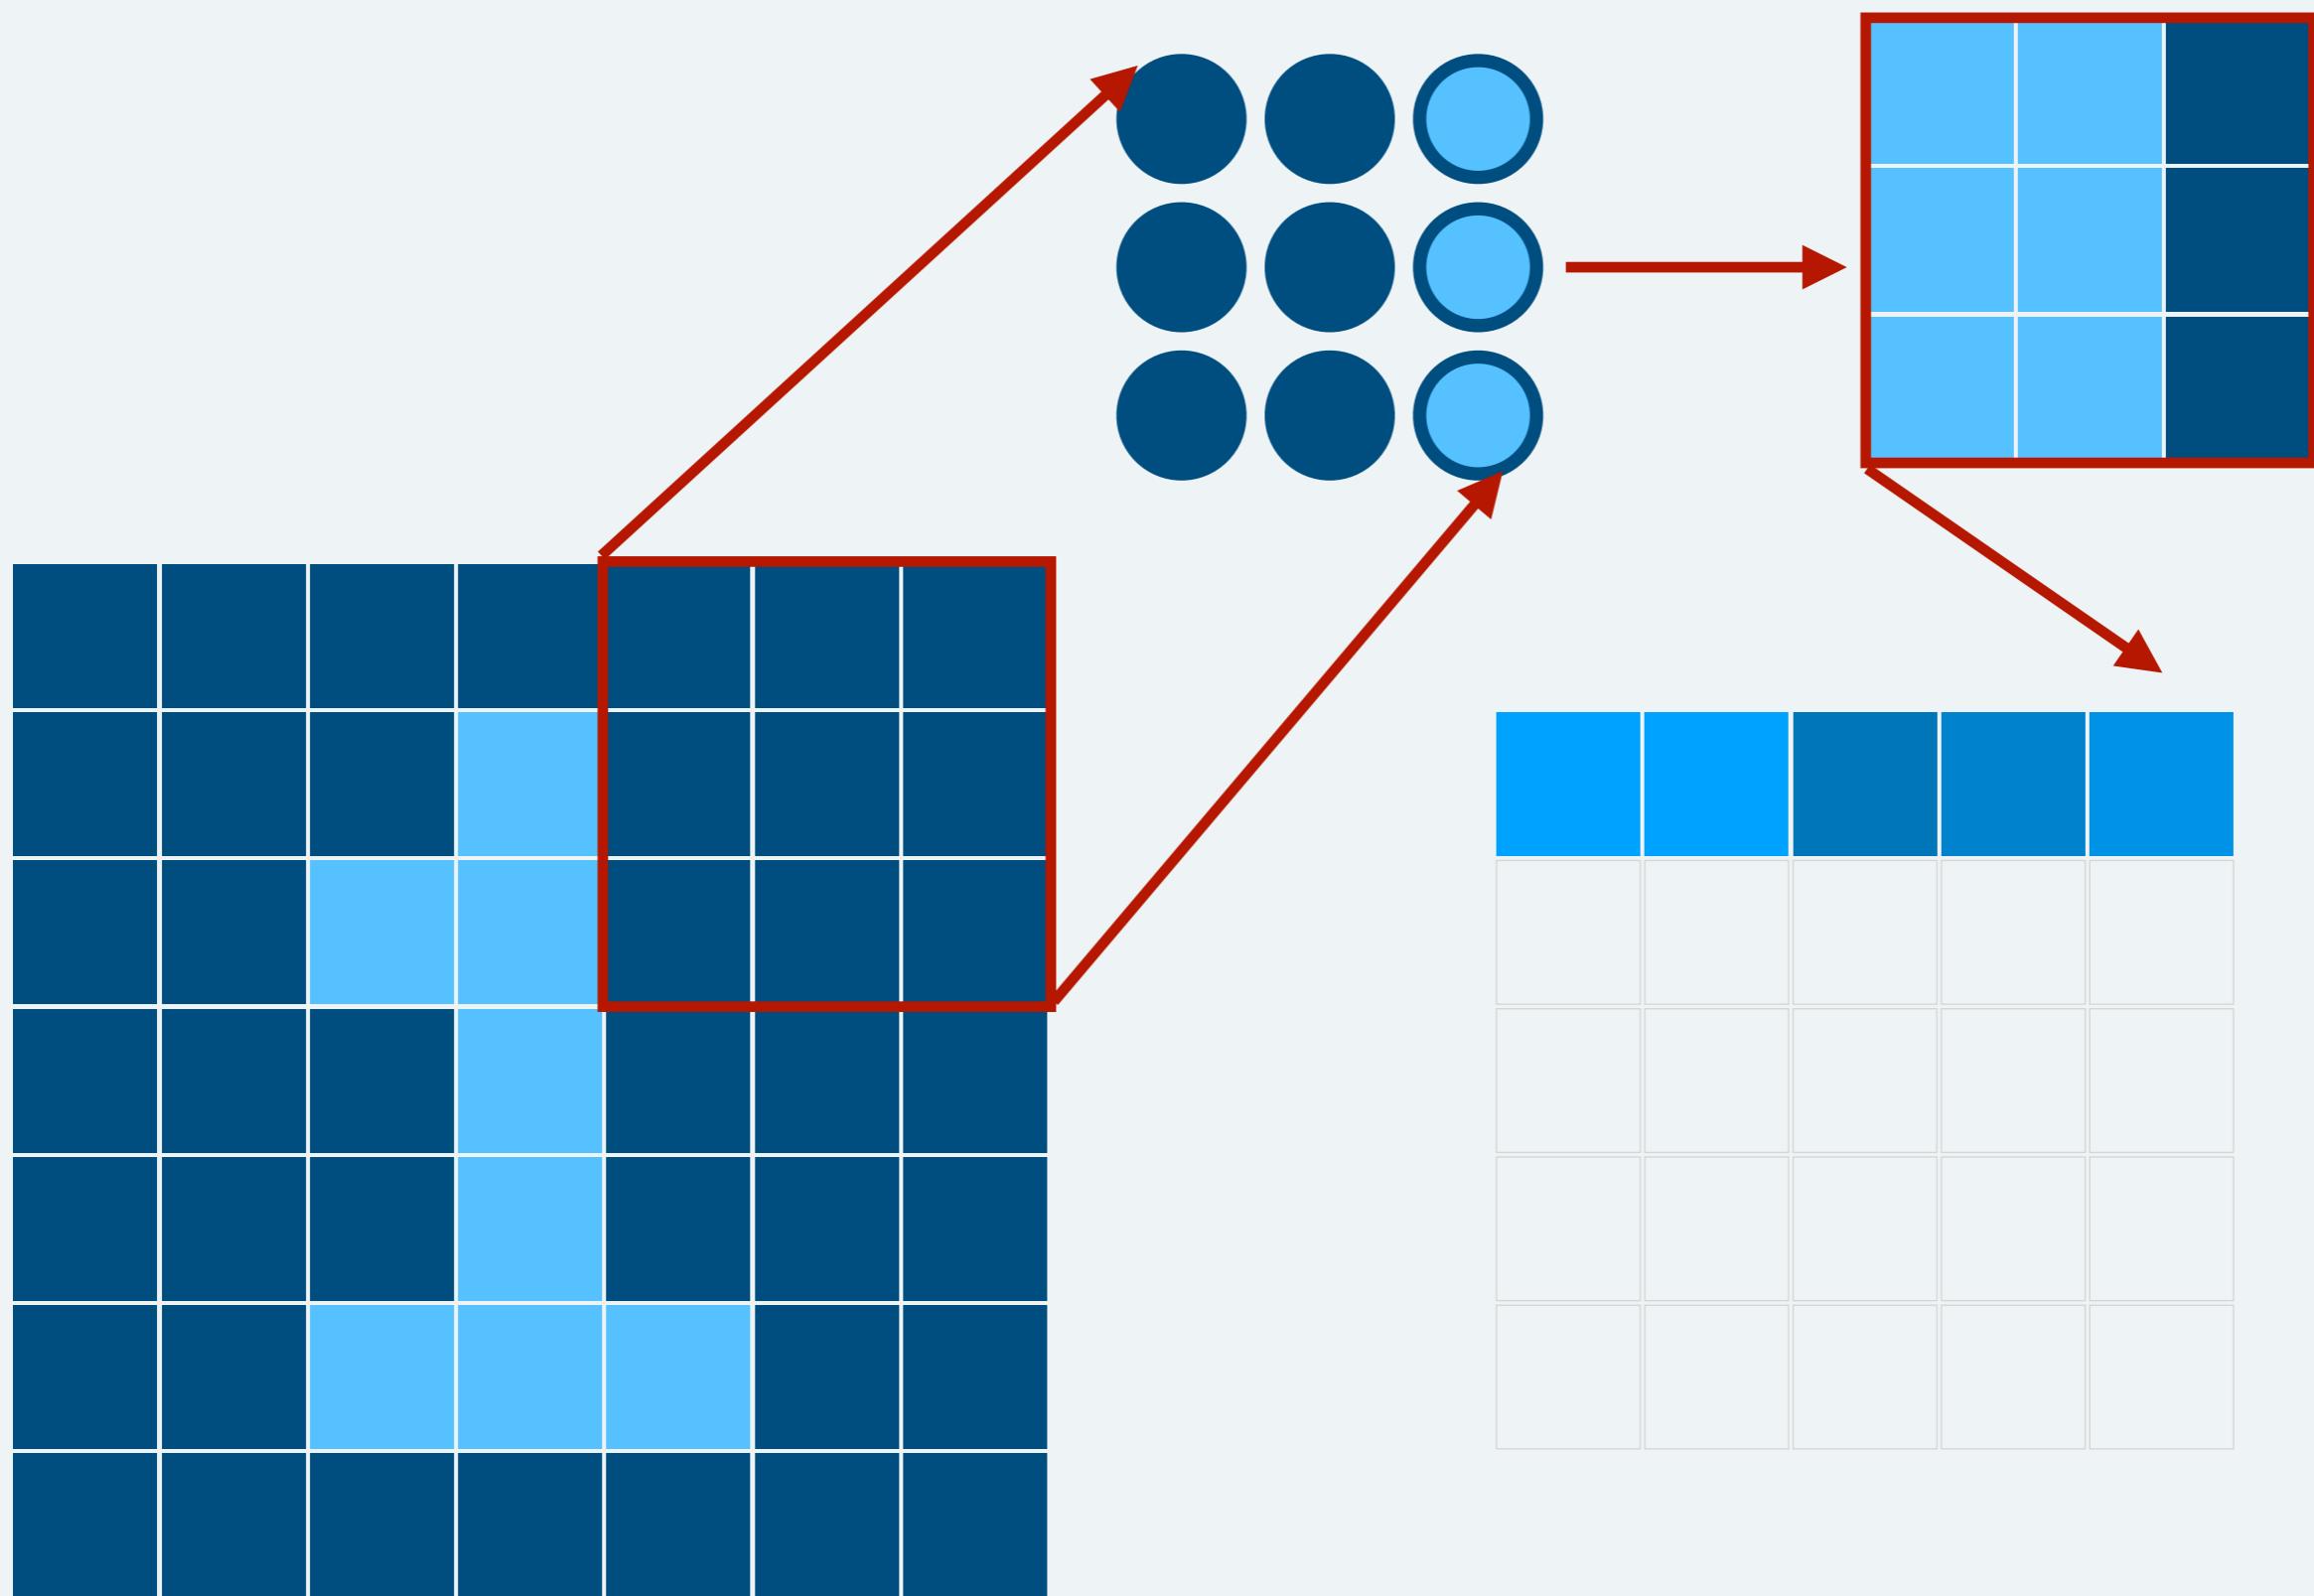
# Convolution



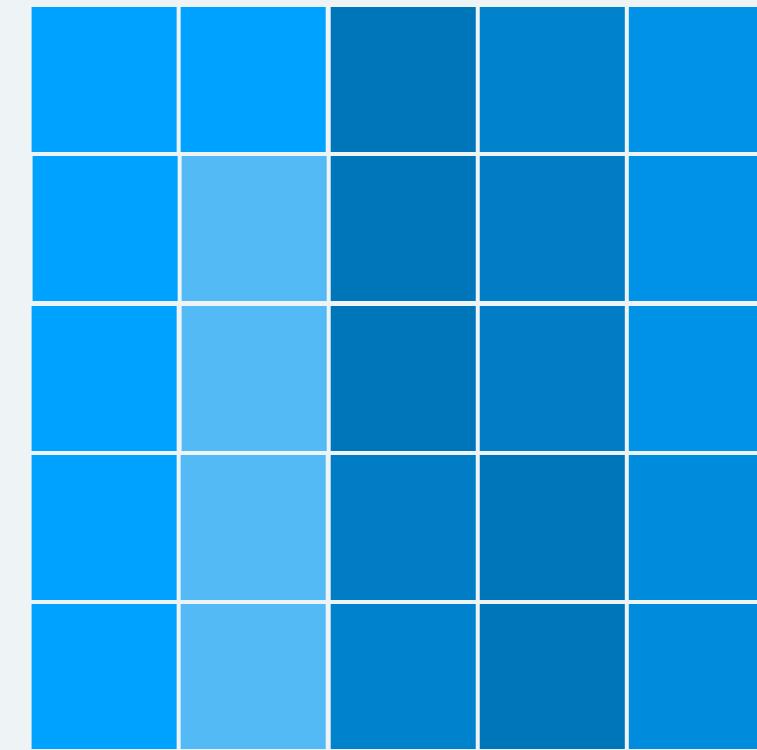
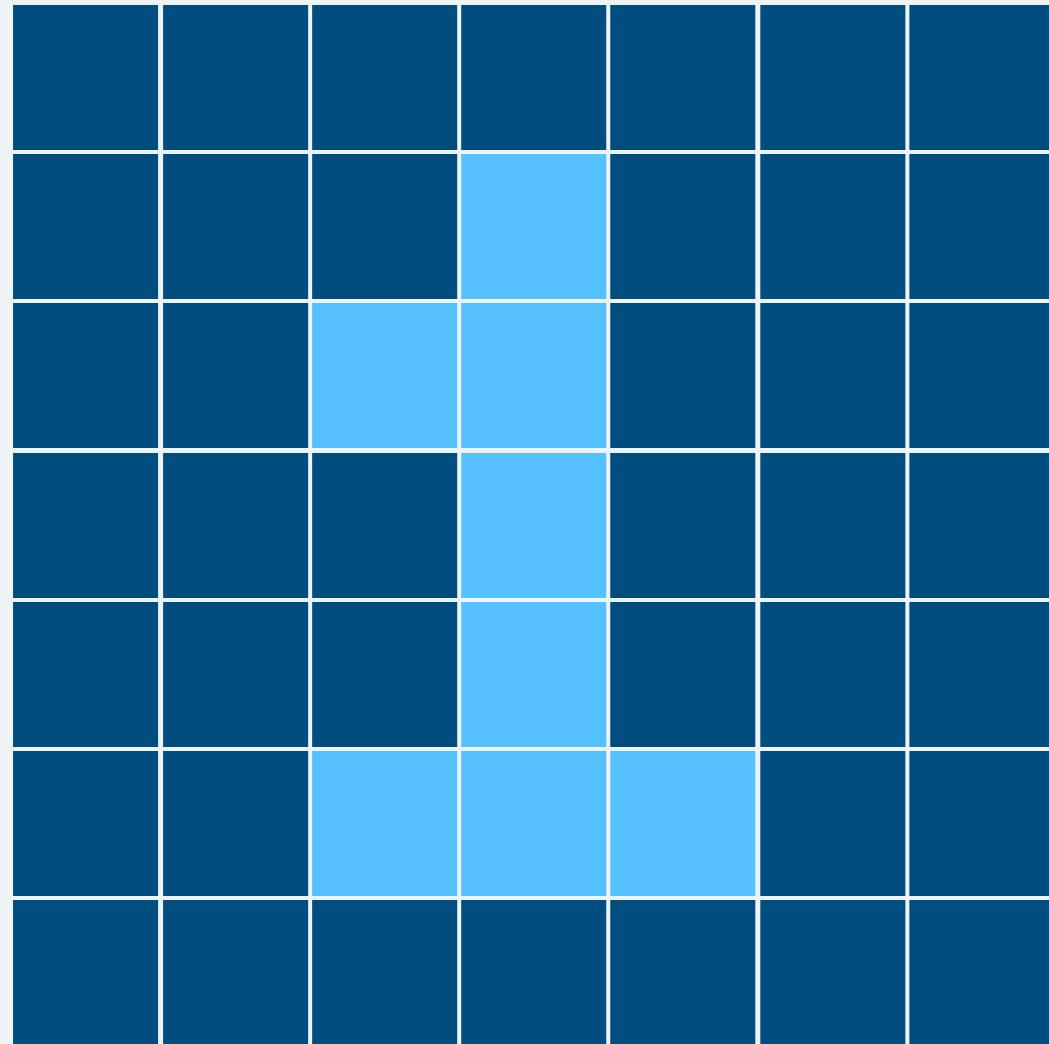
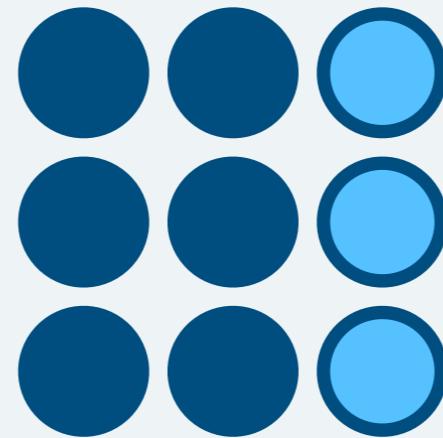
# Convolution



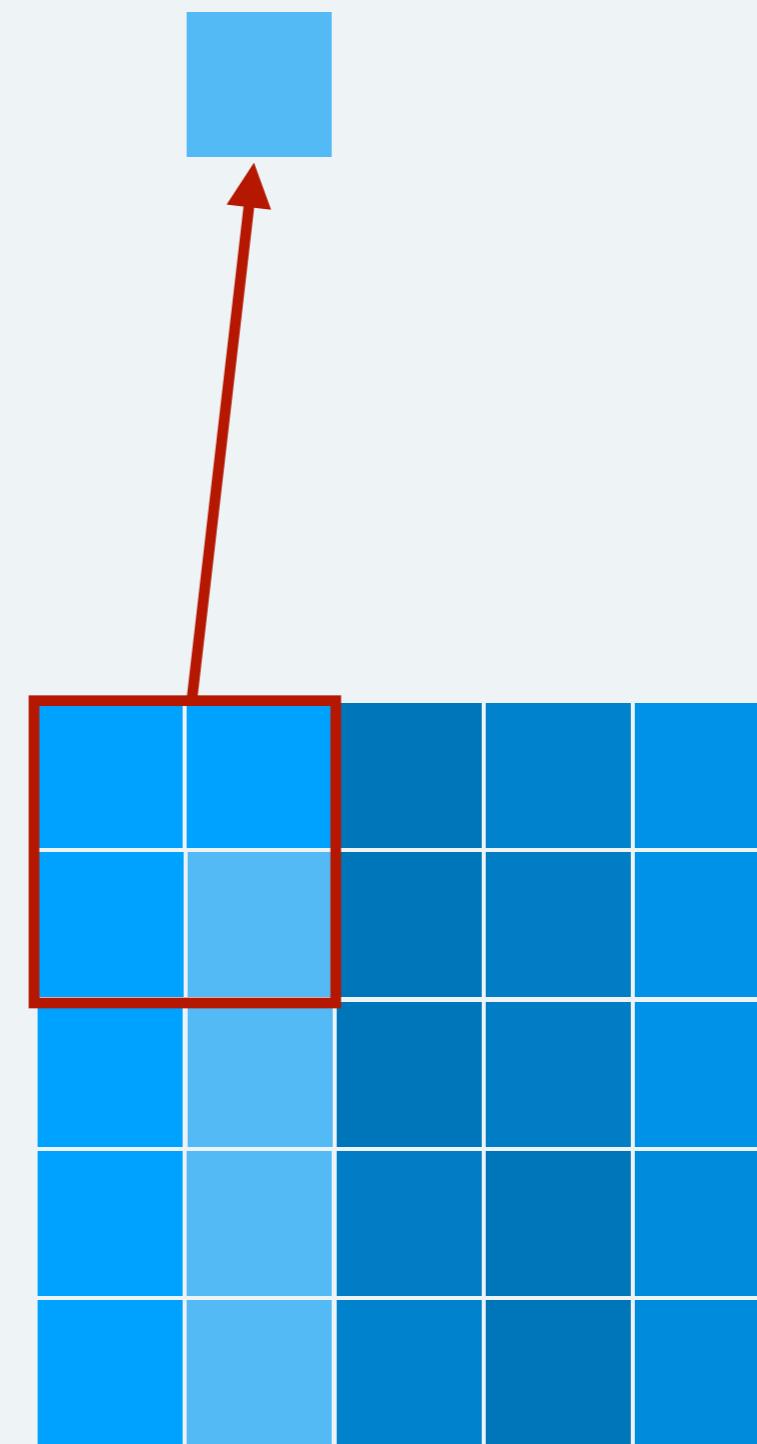
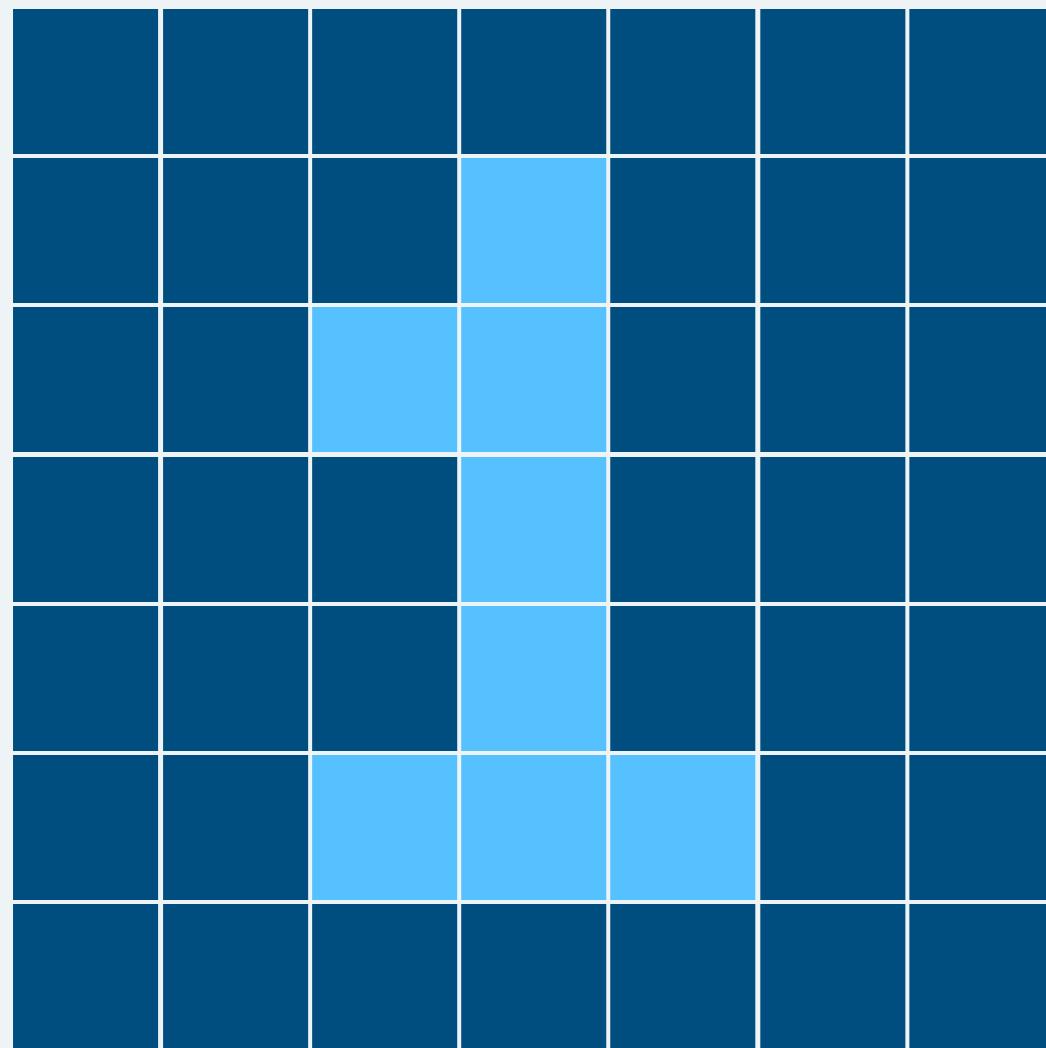
# Convolution



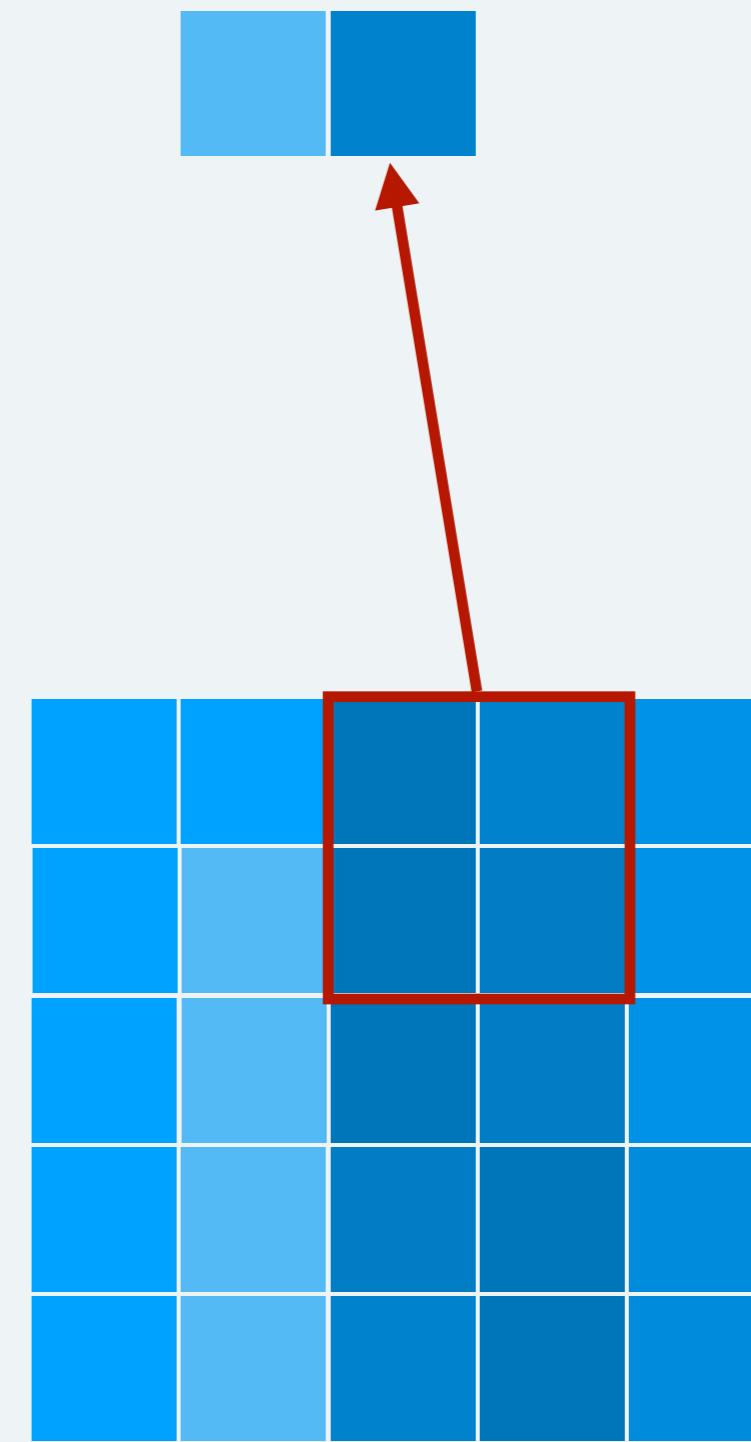
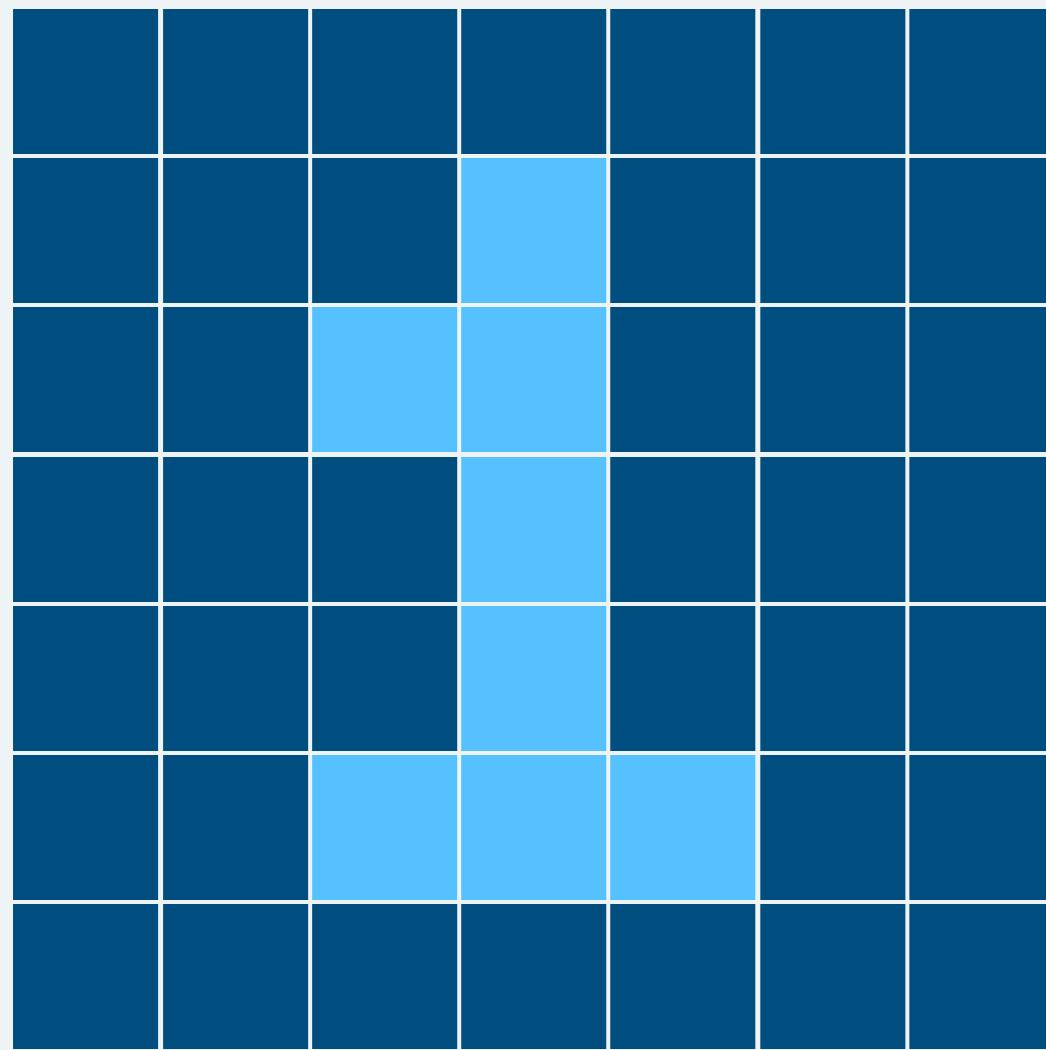
# Convolution



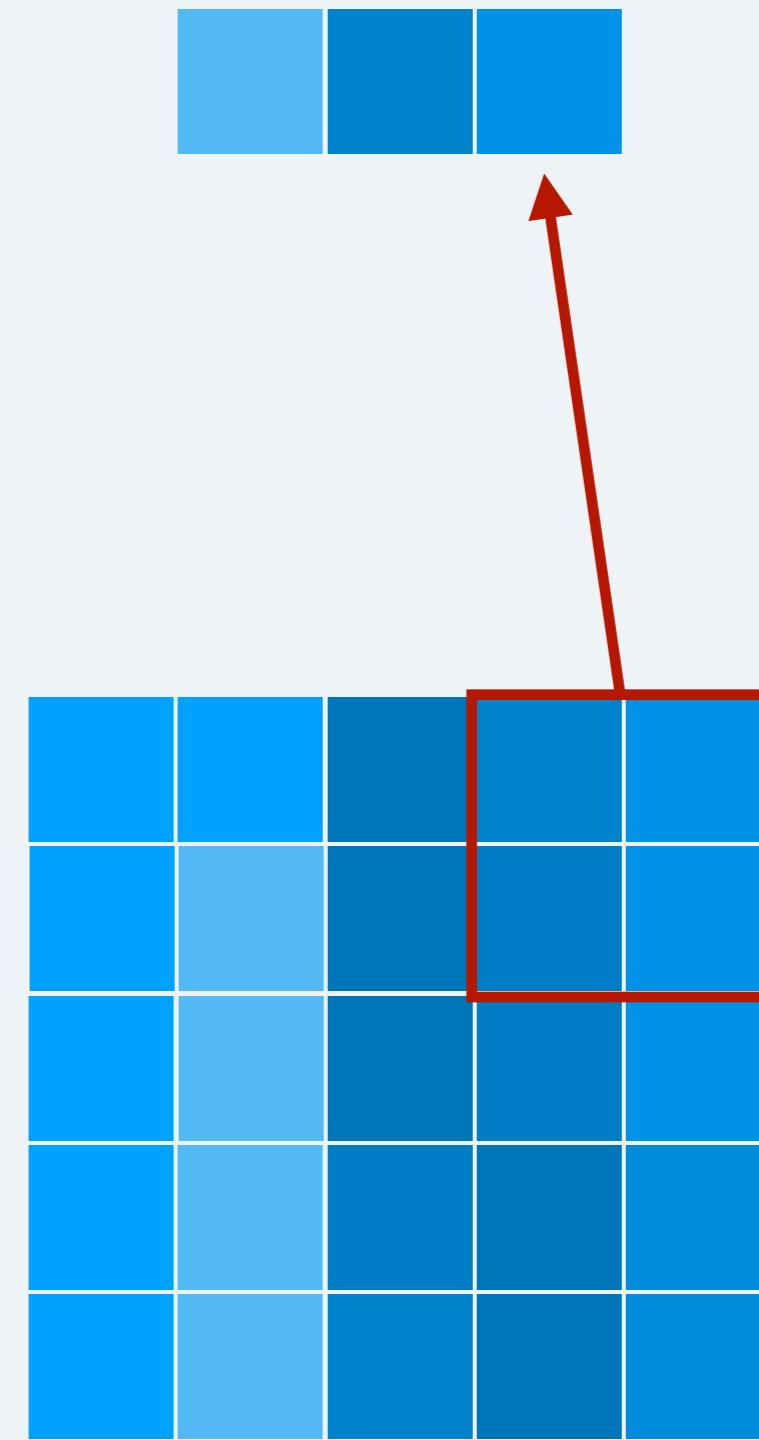
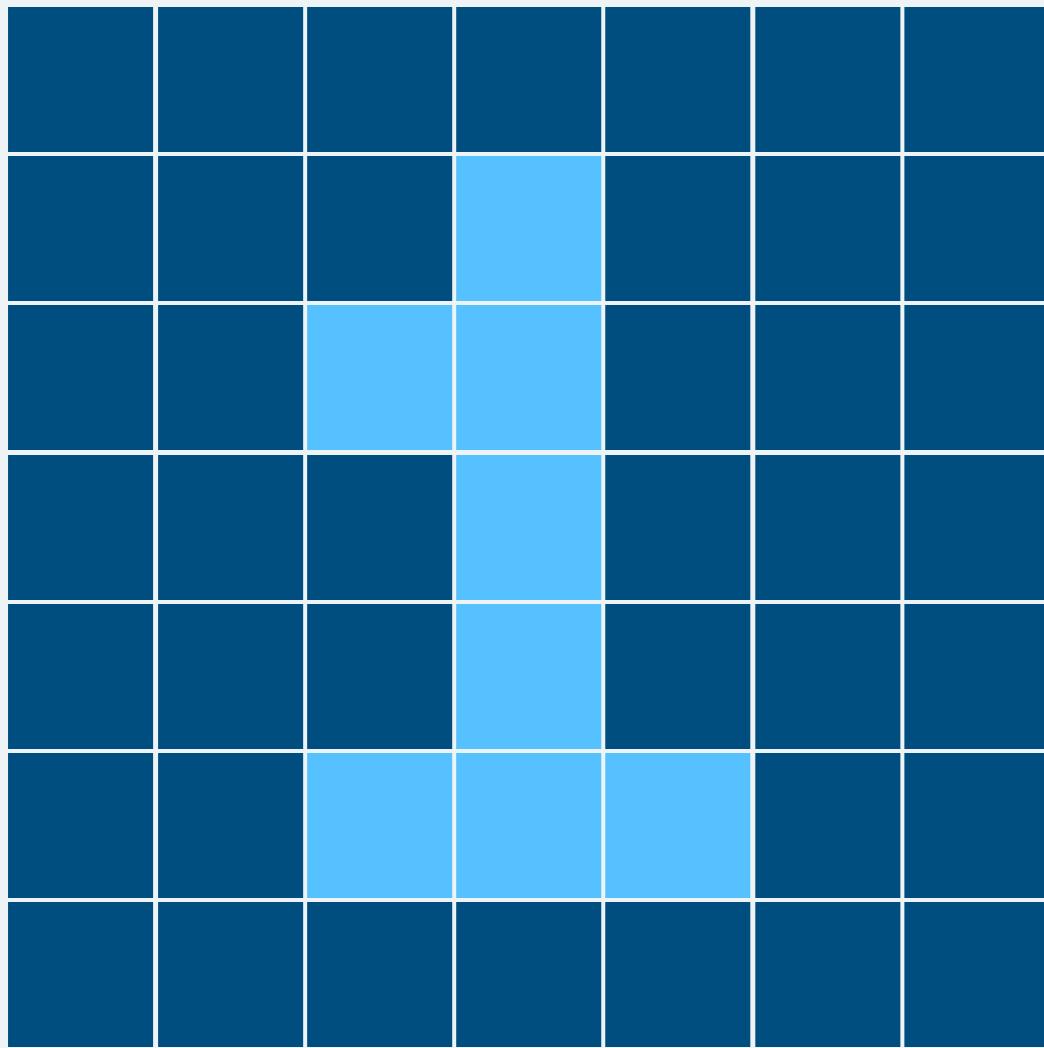
# Pooling



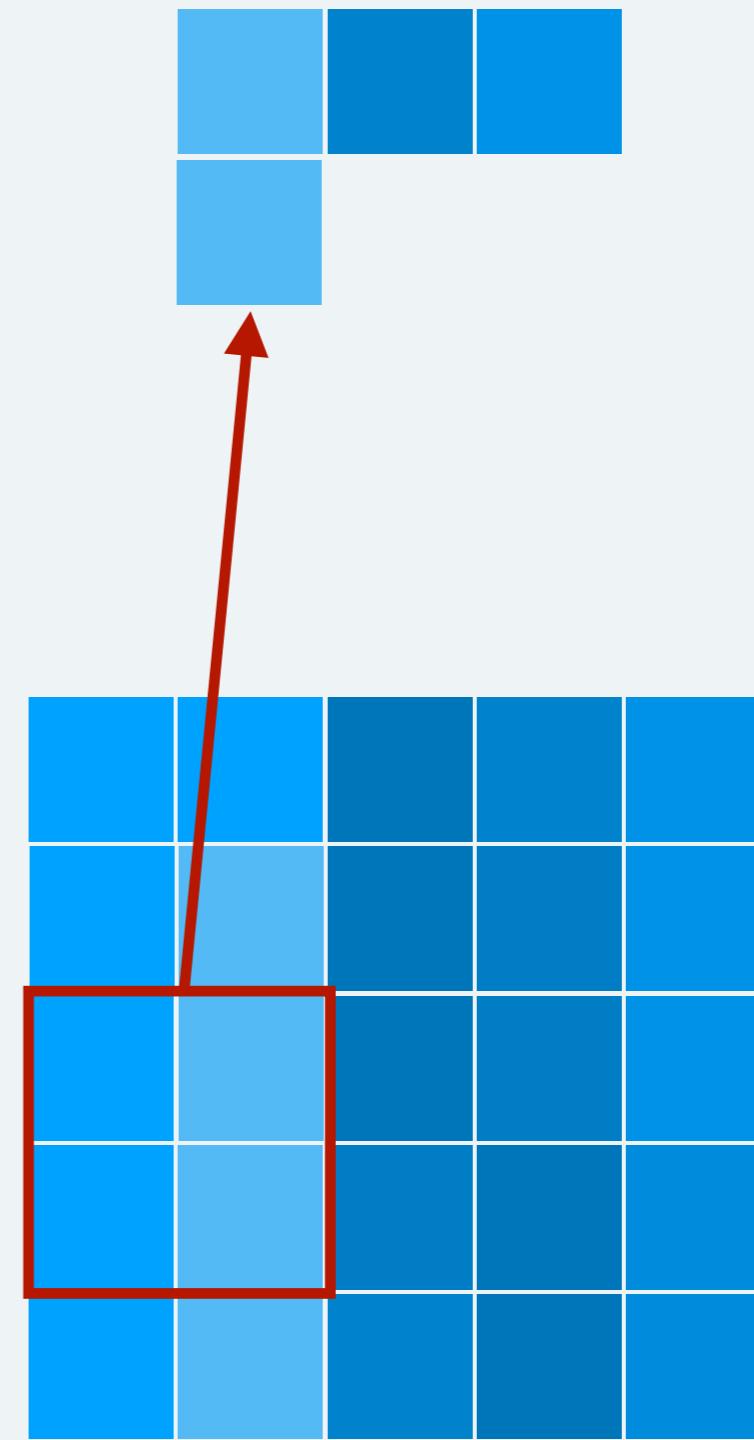
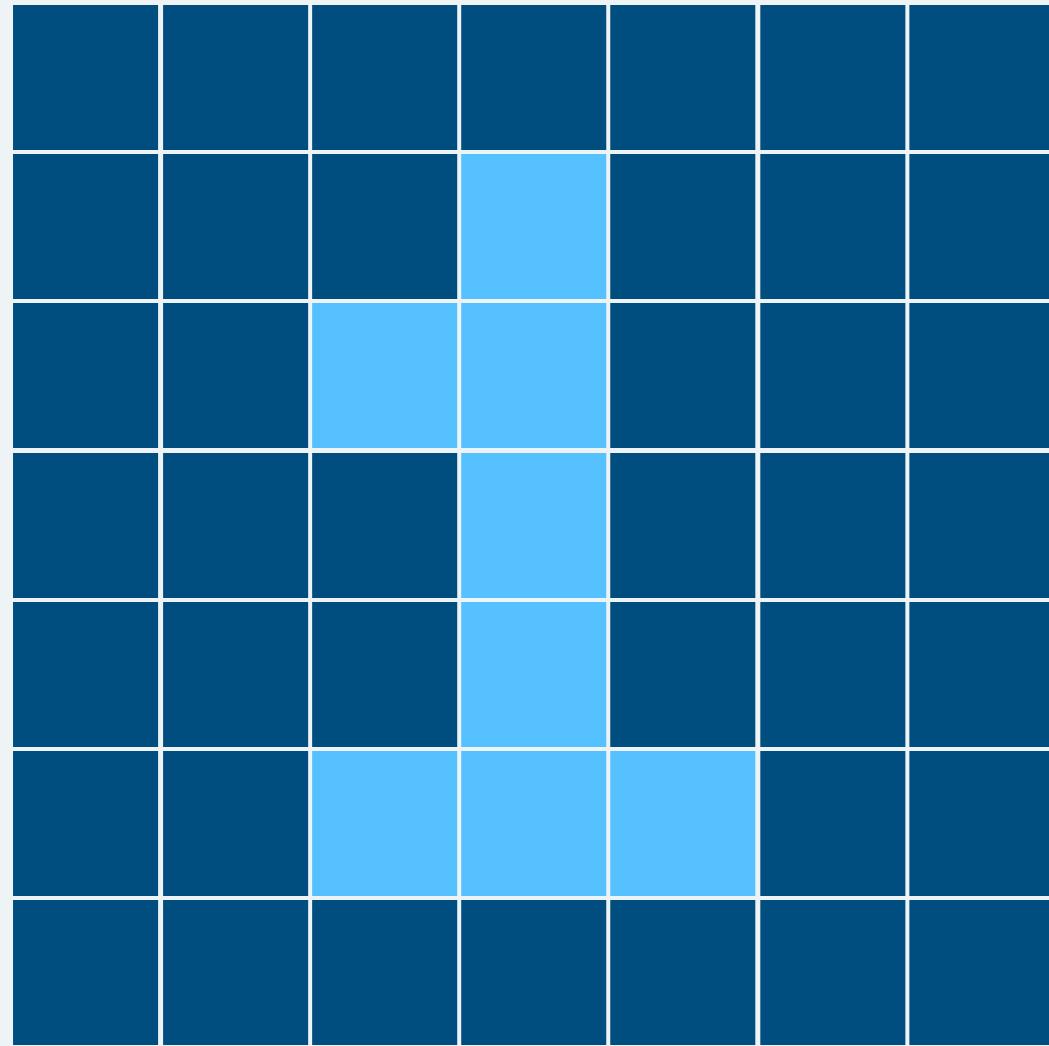
# Pooling



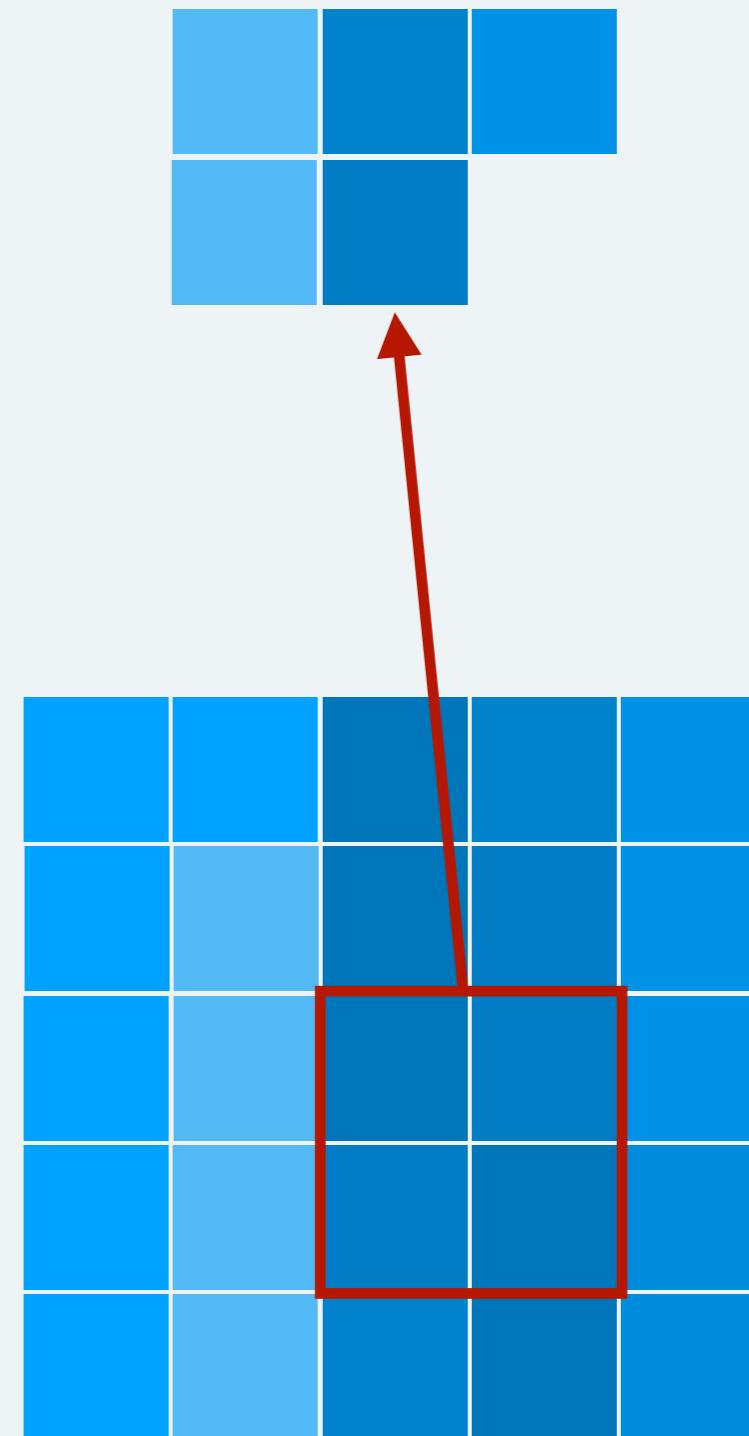
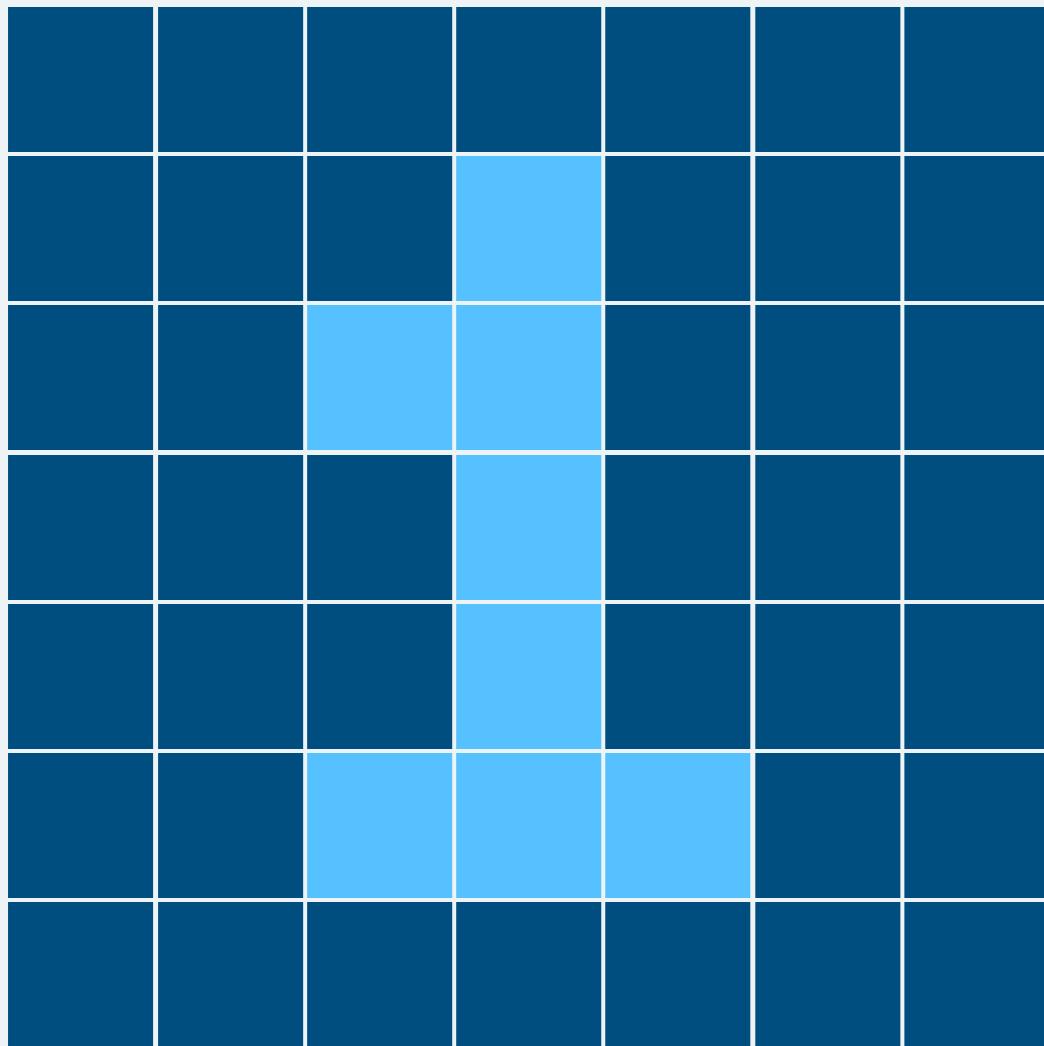
# Pooling



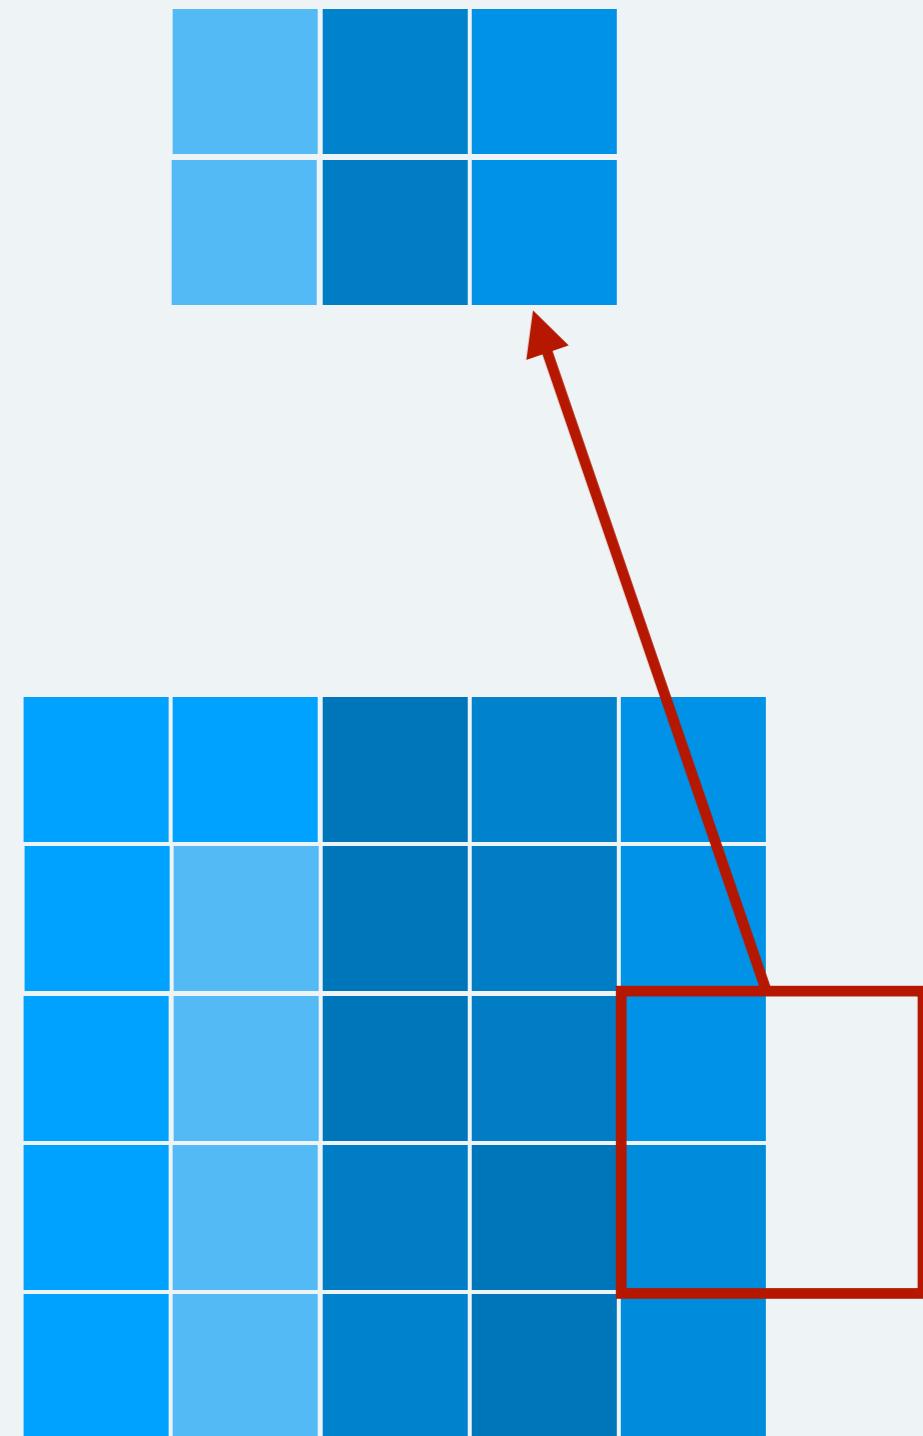
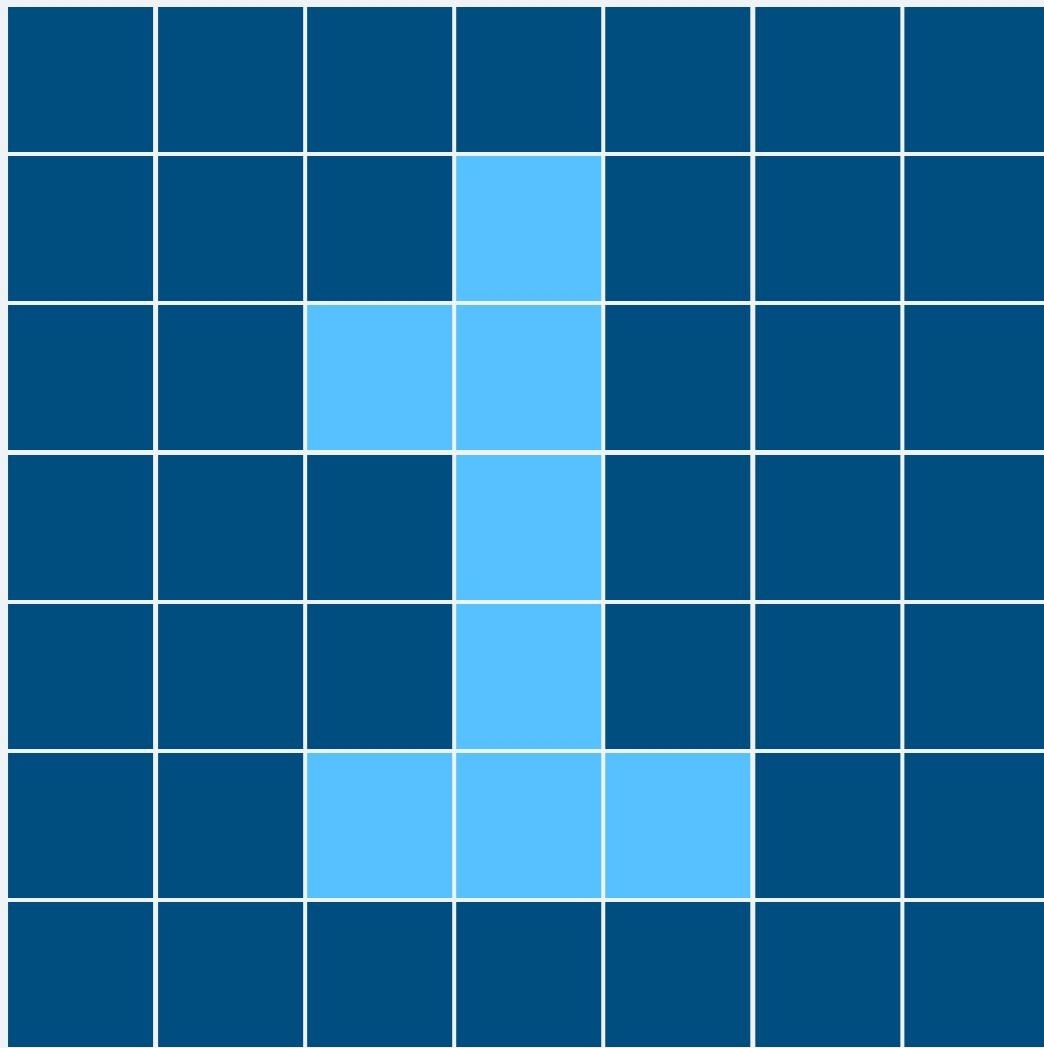
# Pooling

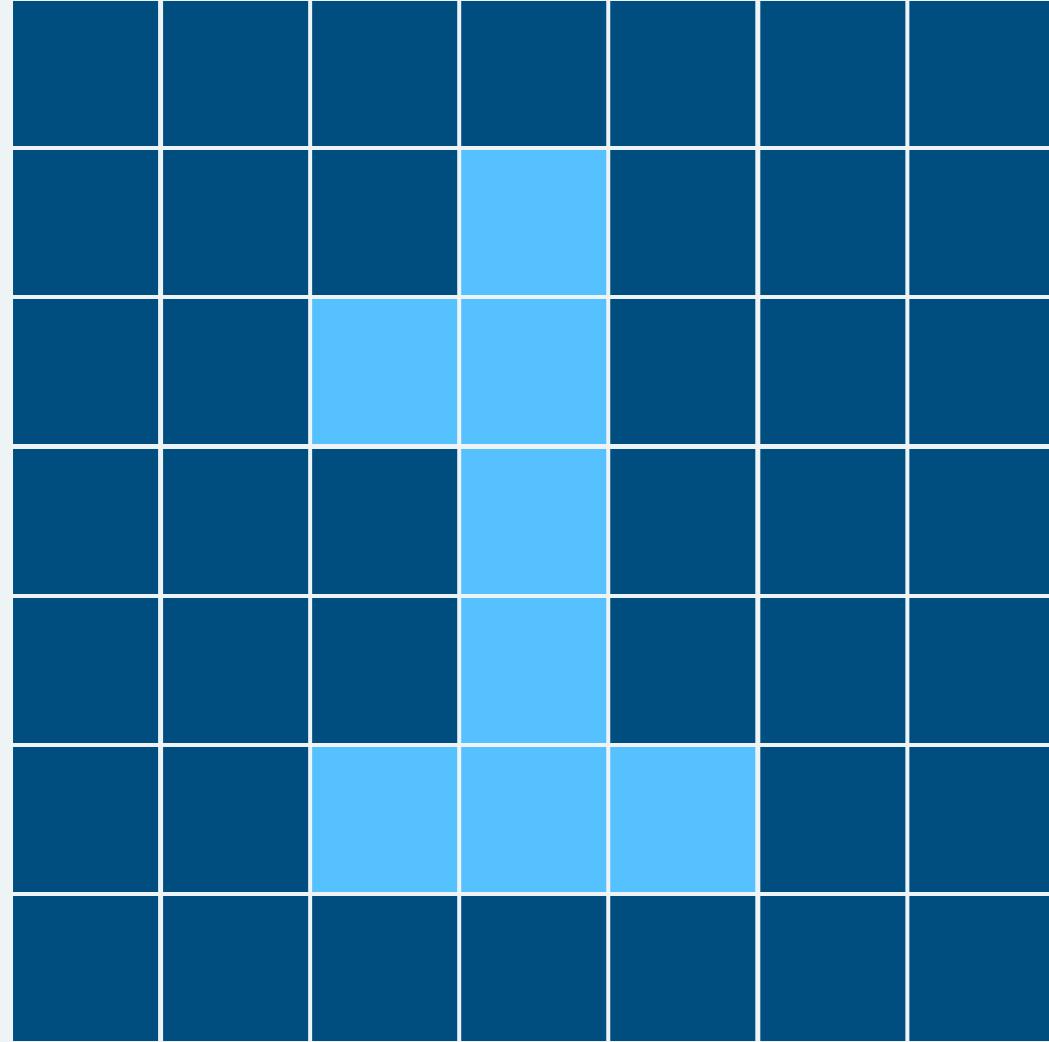


# Pooling

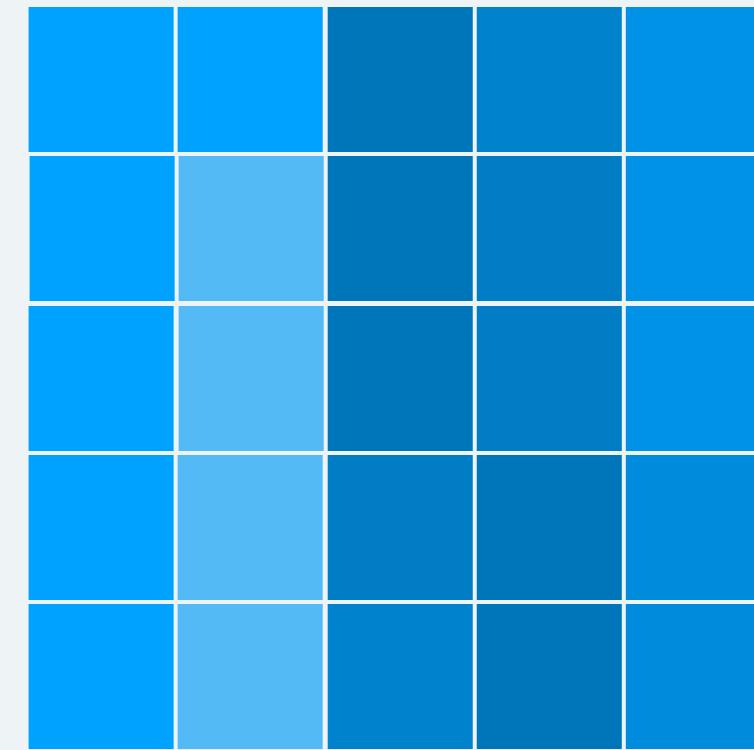


# Pooling

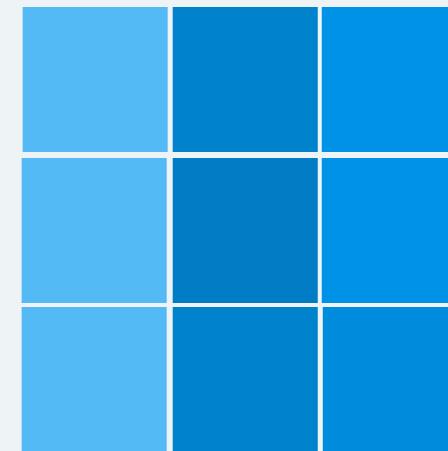


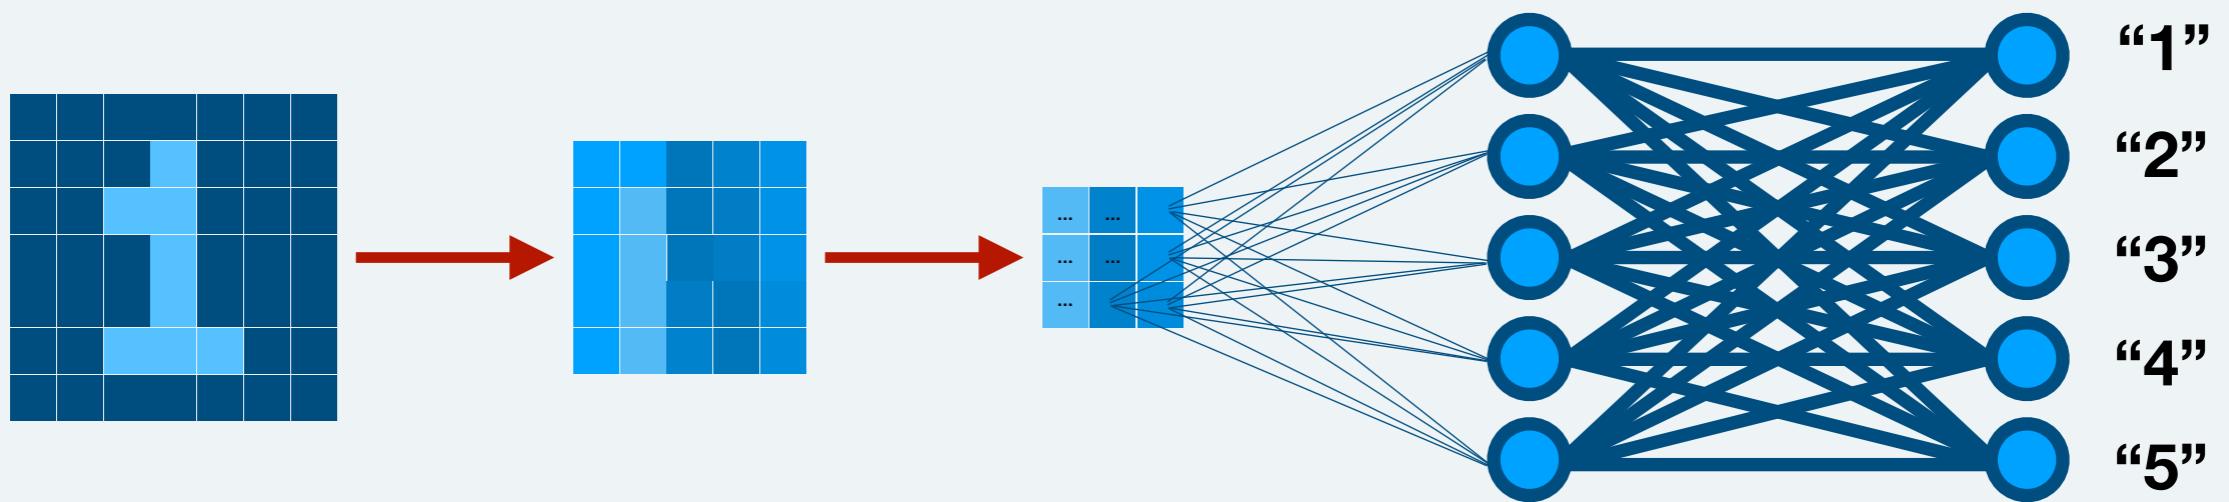


**Convolution**

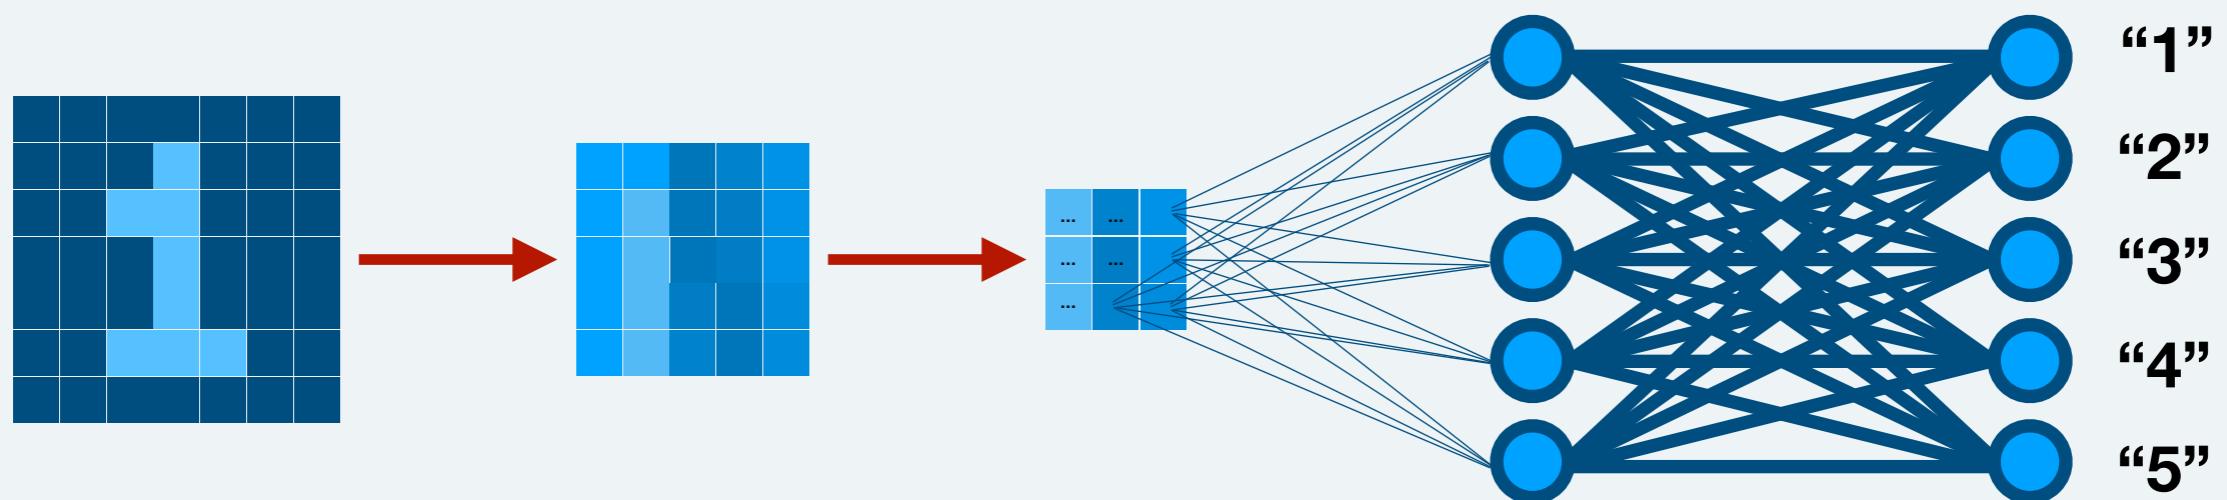


**Pooling**



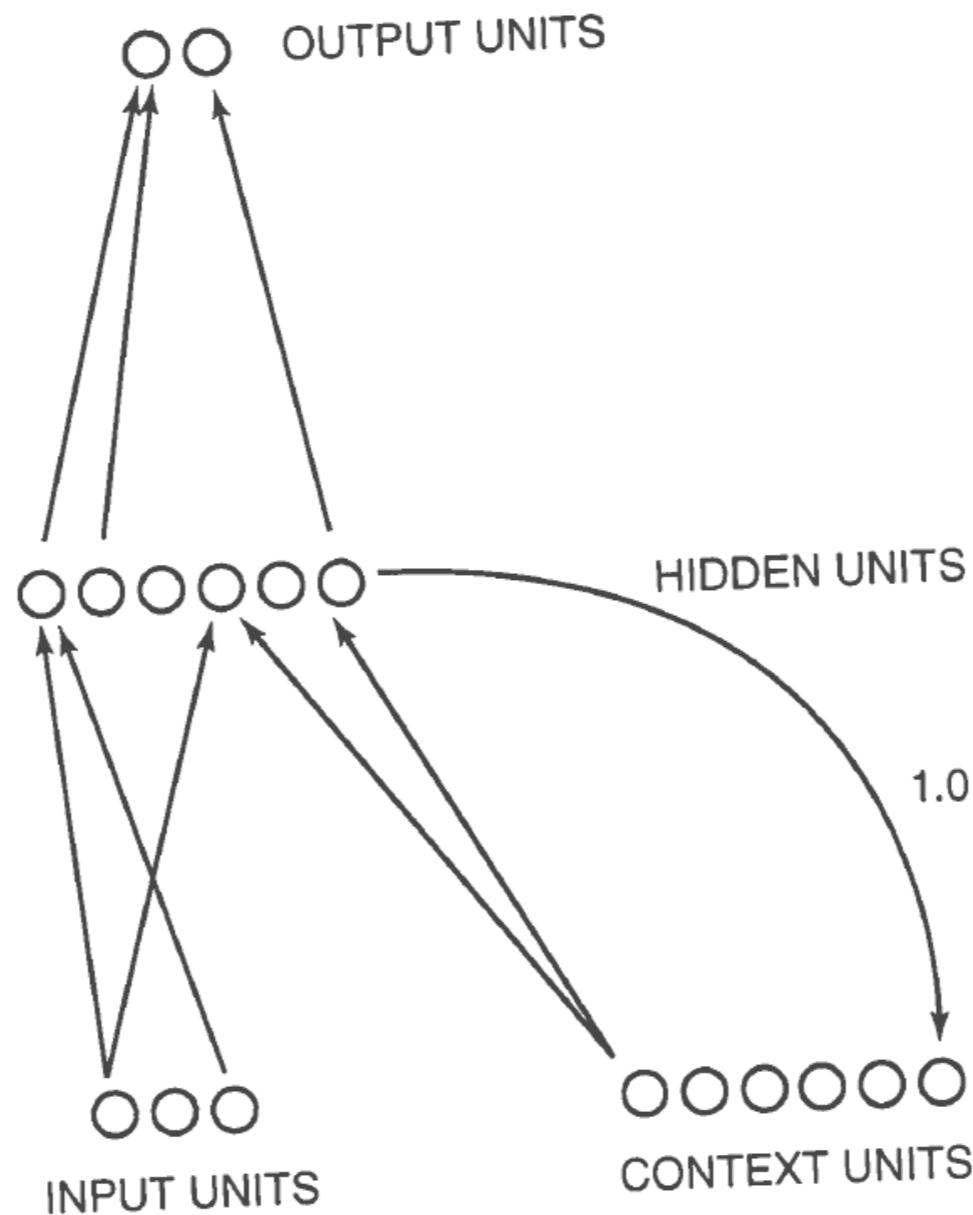


[https://adamharley.com/nv\\_vis/](https://adamharley.com/nv_vis/)



# Connectionist networks

## Recurrent network



# Revision

Let's start by taking stock of what we've learnt about connectionism so far...

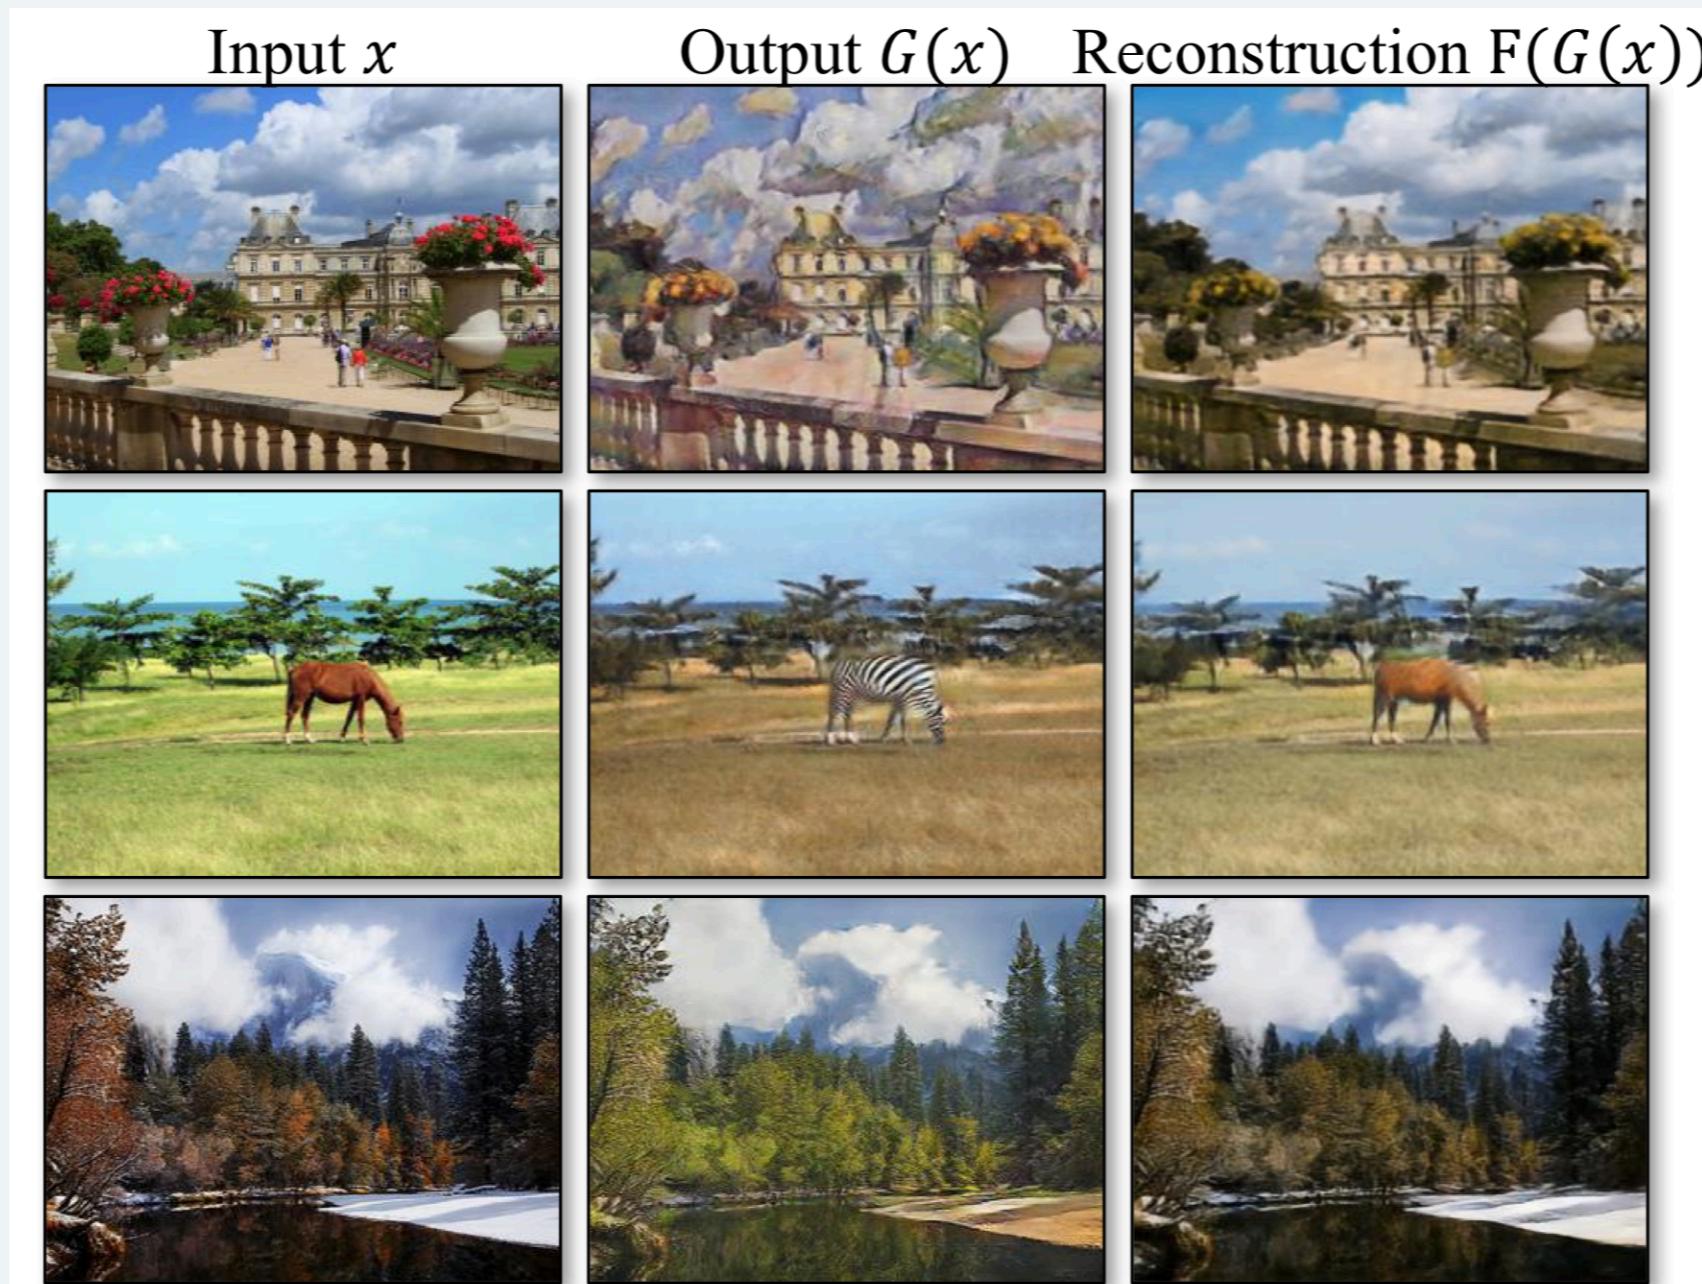
- the perceptron
- the perceptron learning rule
- linear separability
- limitations of single-layer perceptrons
- multi-layer neural networks
- gradient descent
- distributed representations
- fully connected, convolutional, recurrent network

# Bad robots!



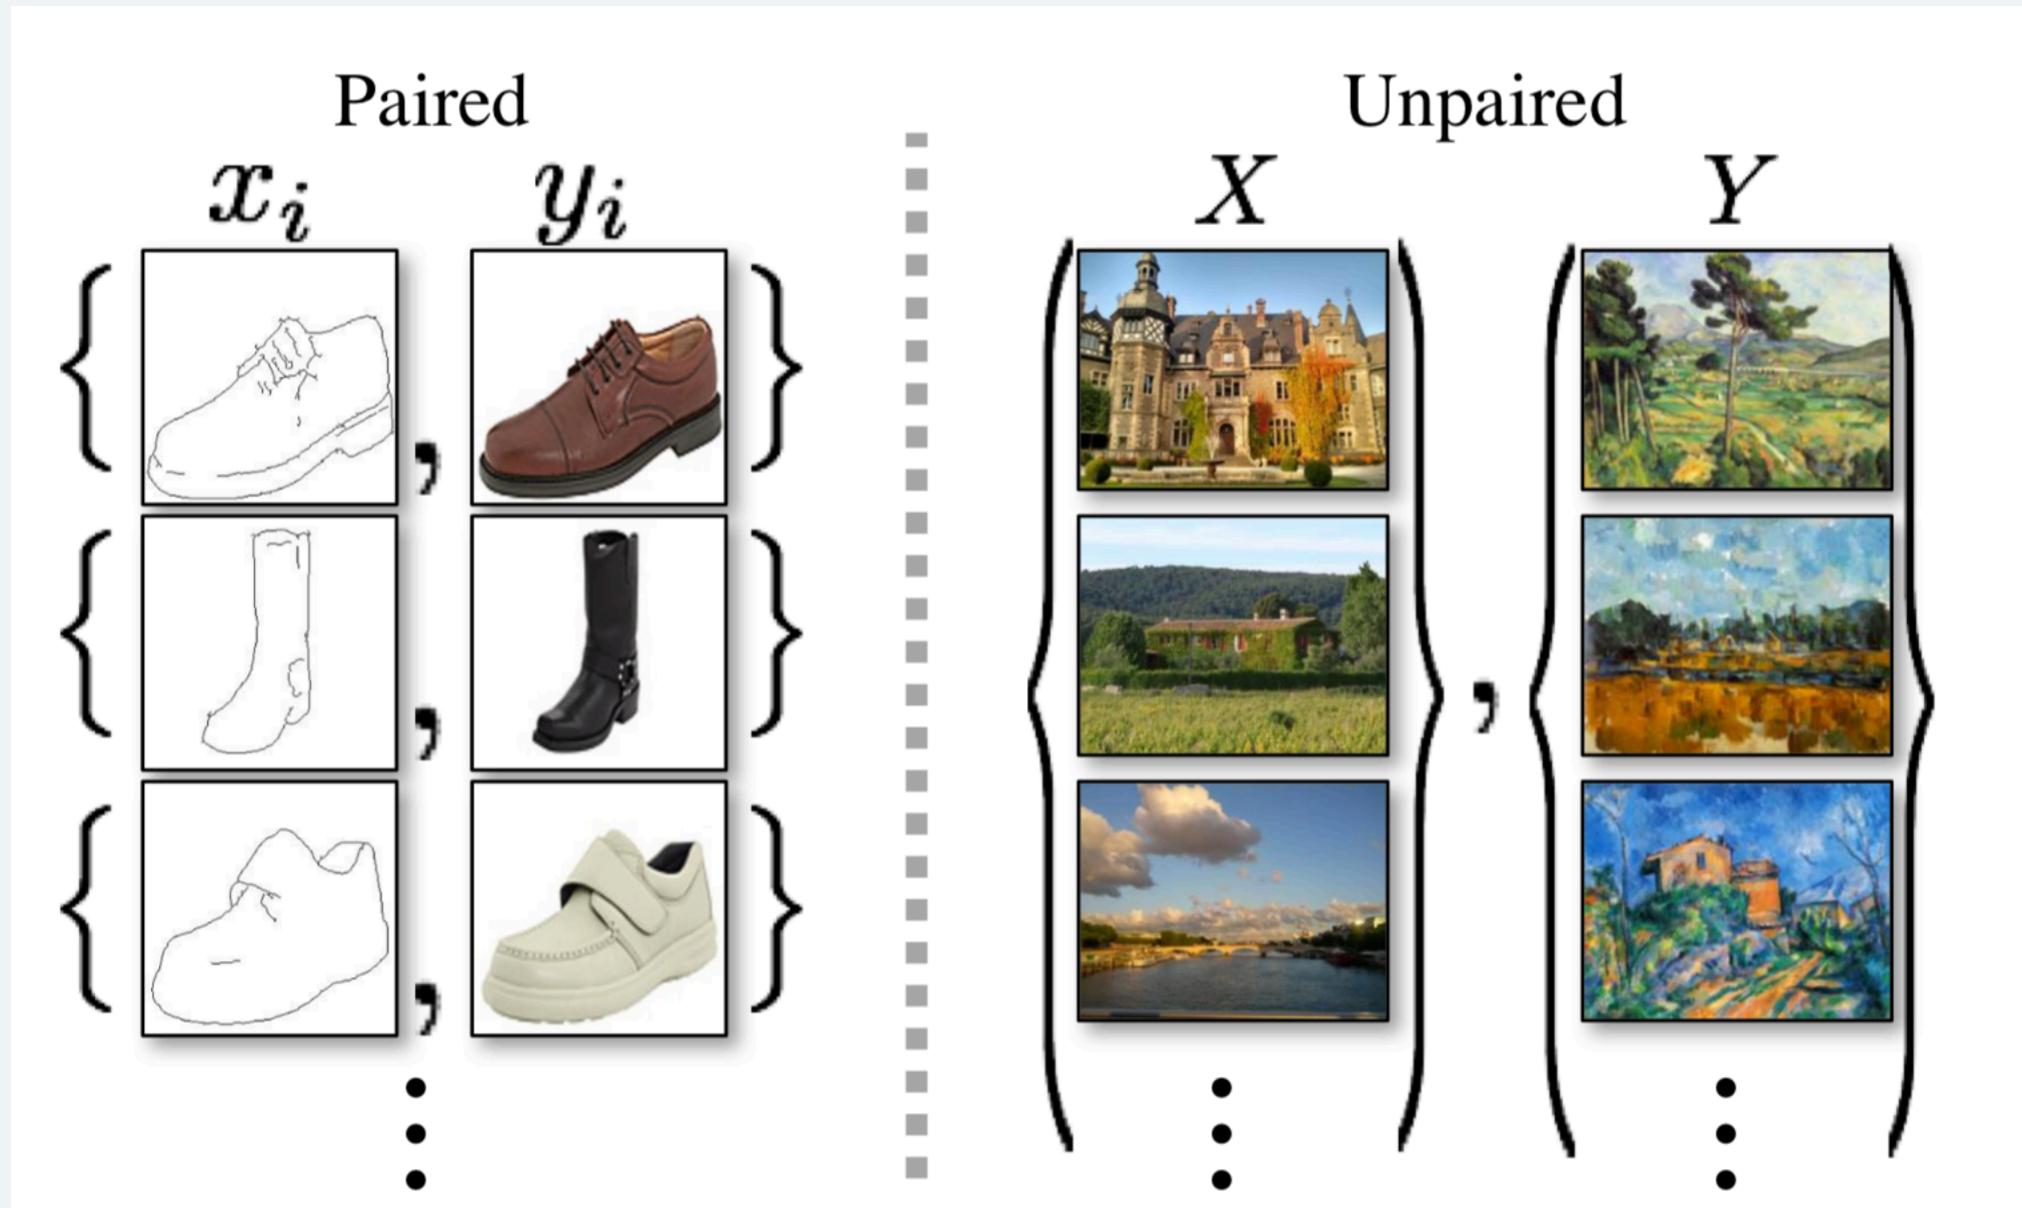
# CycleGAN

**CycleGAN: a deep neural network for converting images from one domain into another**



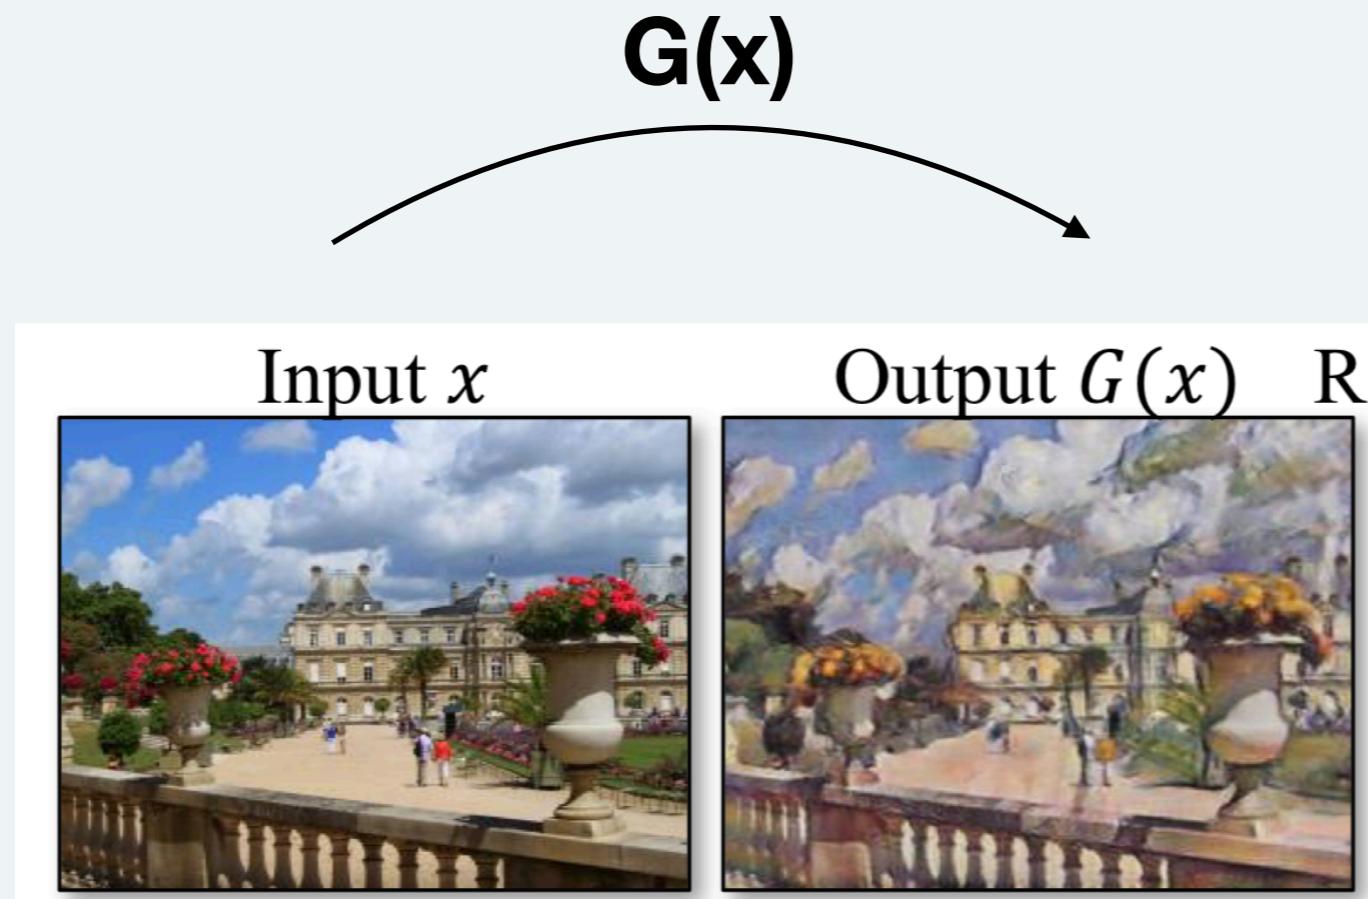
# CycleGAN

**CycleGAN does this without relying on paired images**



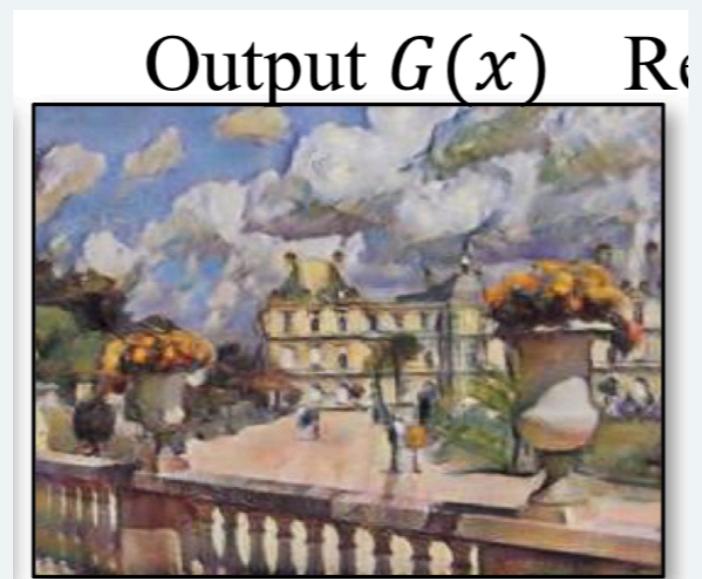
# CycleGAN

**first network: learns conversion  $G(x)$  from domain A to domain B**



# CycleGAN

**second network: tries to tell if  $G(x)$  is a ‘fake’ or a real image from  $Y$**

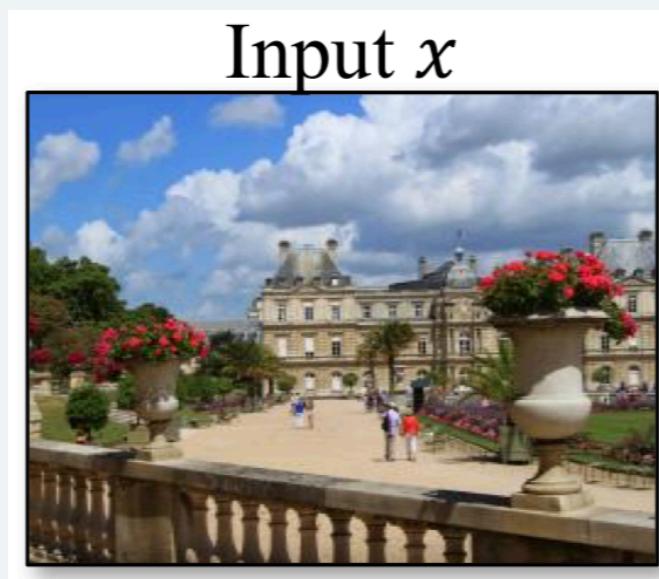


**vs.**



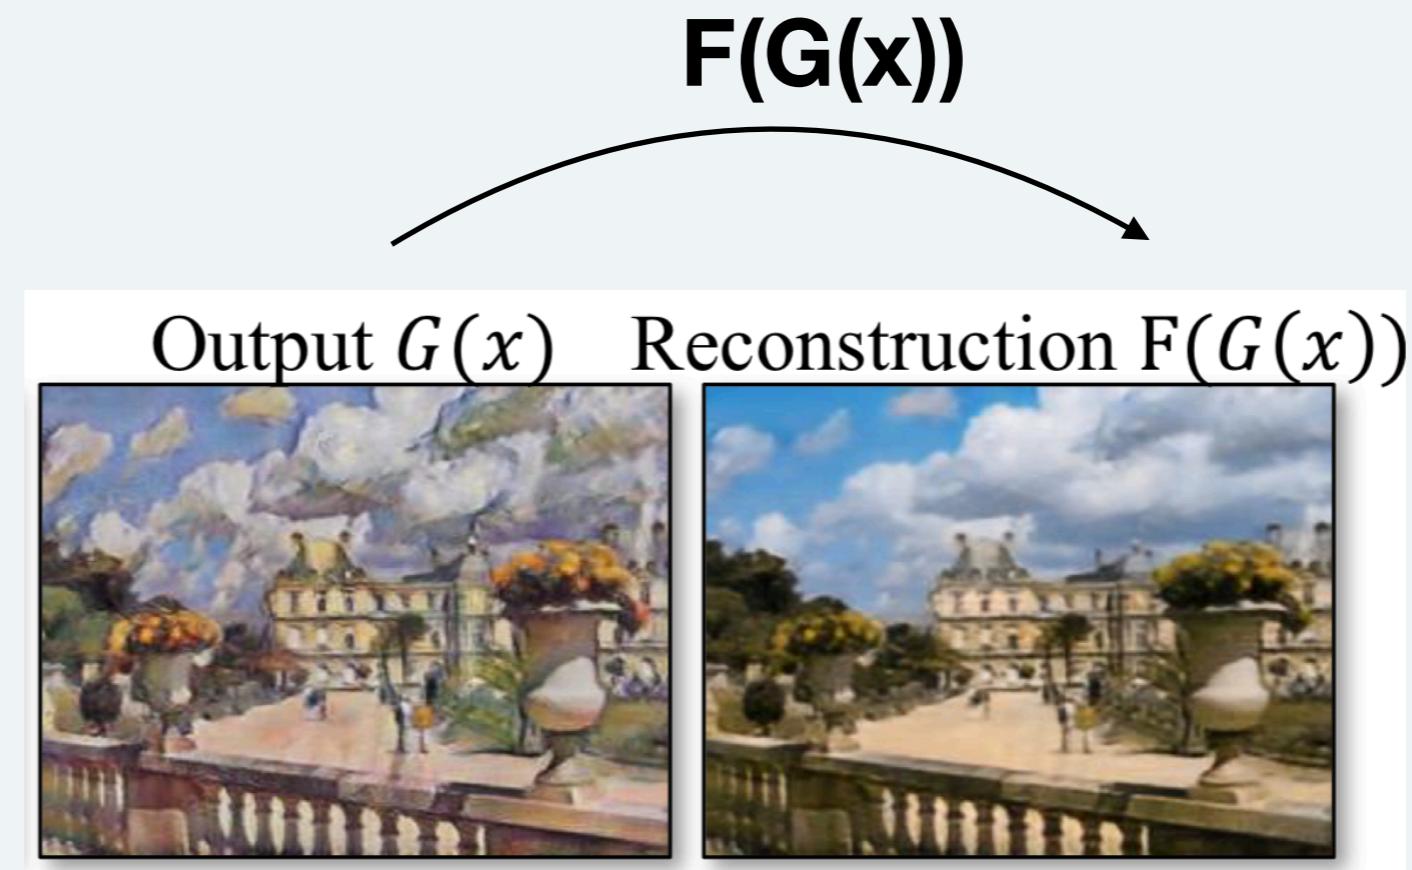
# CycleGAN

**but, at this point,  $x$  could in principle be translated into  
\*any\* image from  $Y$ :**



# CycleGAN

**third network: learns conversion  $F(y)$  from domain Y to domain X**



# CycleGAN

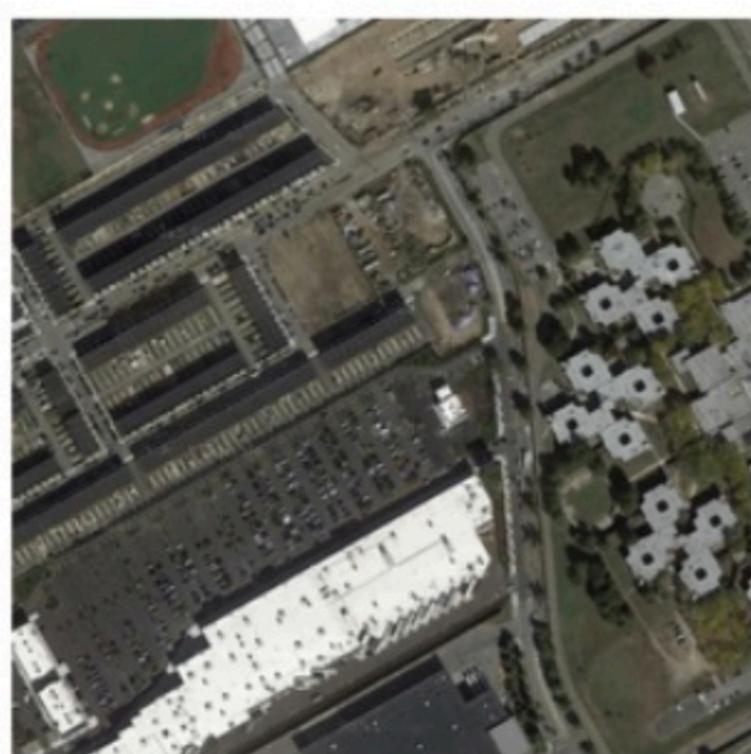
**the whole AI is constrained to minimise the difference between  $x$  and  $F(G(x))$**

$$\text{diff}( x, F(G(x)) )$$

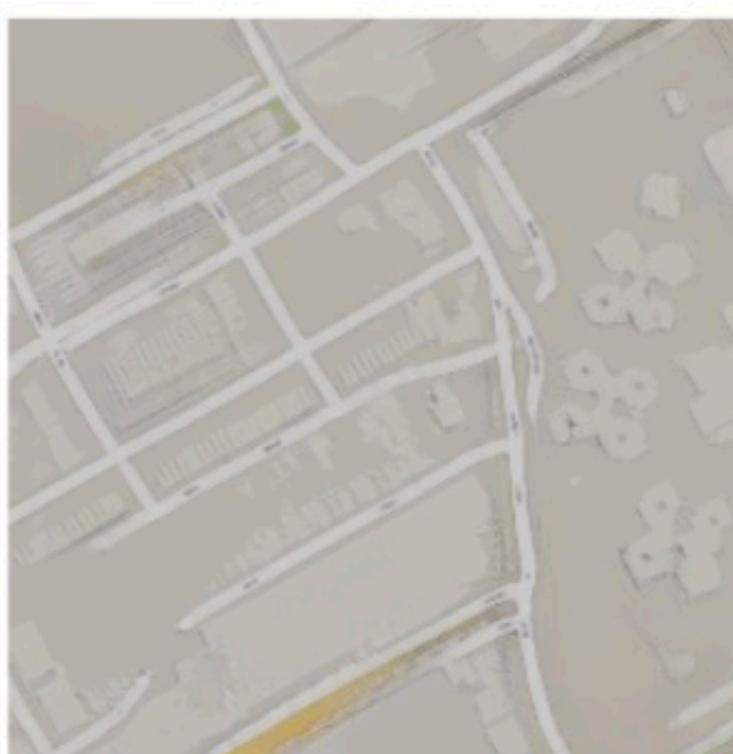


# A connectionist network for creating maps

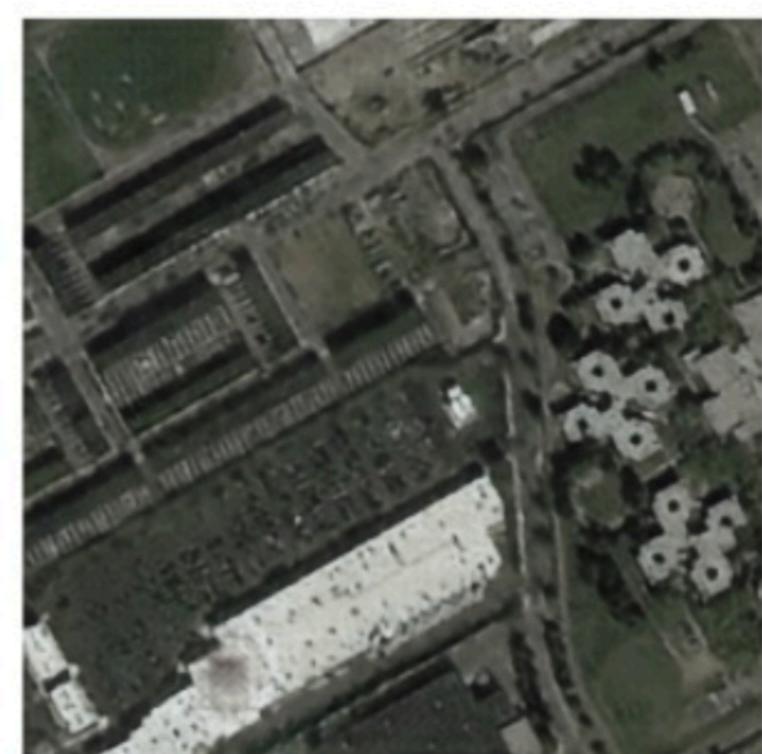
## CycleGAN



(a) Aerial photograph:  $x$ .



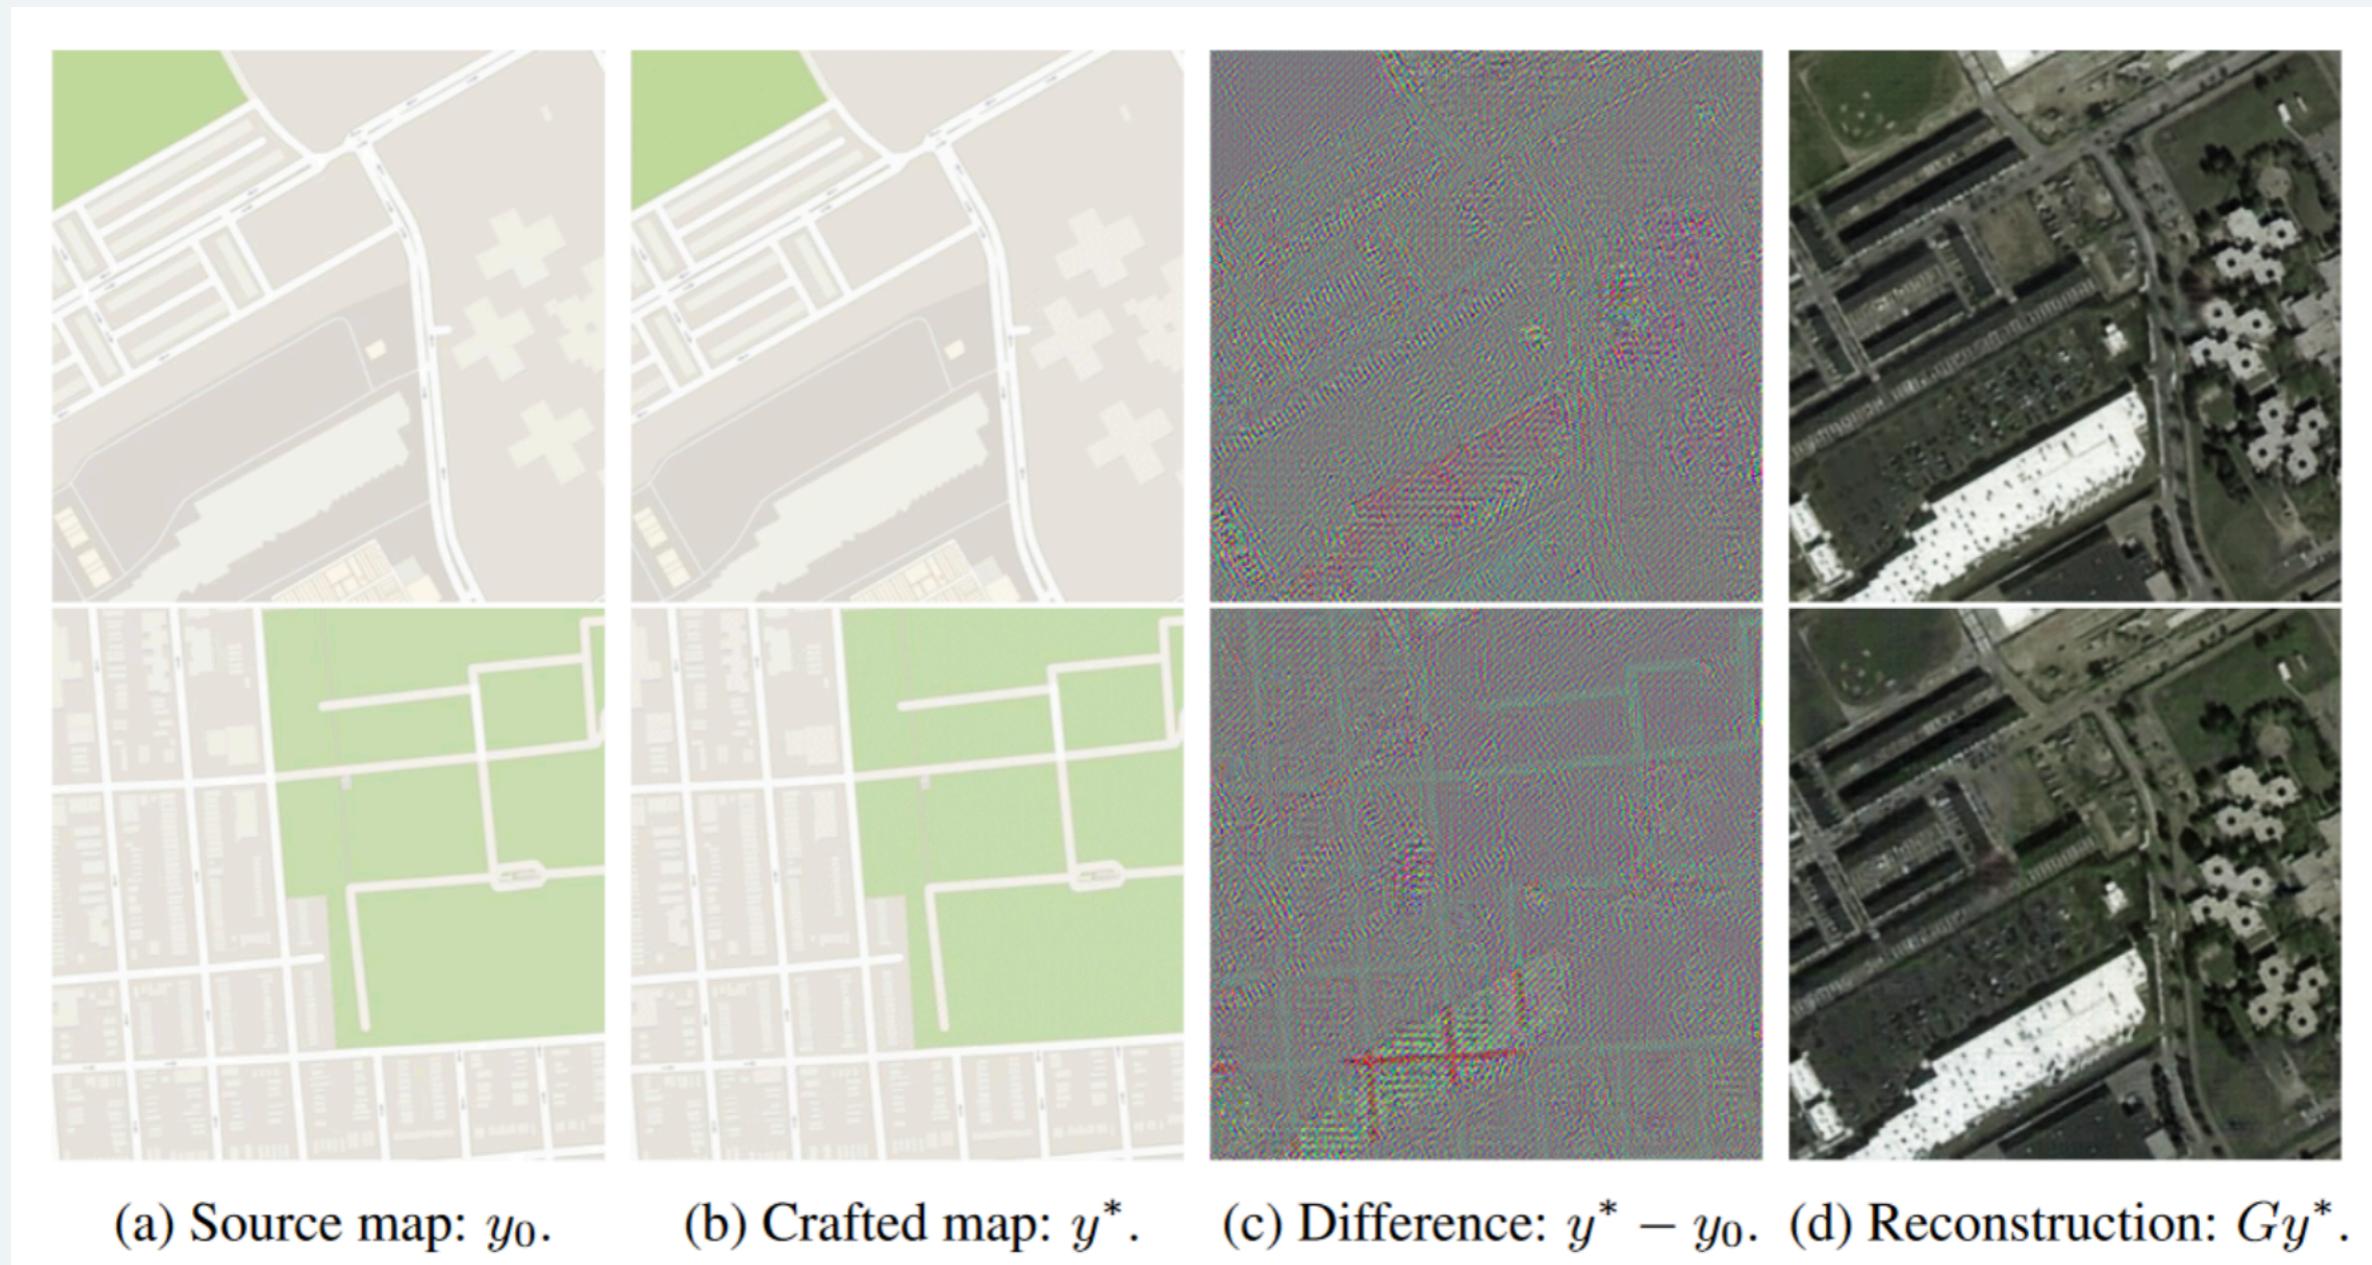
(b) Generated map:  $Fx$ .



(c) Aerial reconstruction:  $GFx$ .

# A connectionist network for creating maps

CycleGAN



**“But in fact this occurrence, far from illustrating some kind of malign intelligence inherent to AI, simply reveals a problem with computers that has existed since they were invented: they do exactly what you tell them to do.”**

**source: TechCrunch,  
[link](#)**